

6809 Forth — Open Items Checklist (Part 2 of 2: Unit Test Framework)

This is part 2 of 2, picking up exactly where part 1 leaves off: the assembly-level unit test framework's own introduction, and every glossary-section test group (3.1 through 3.13 so far) built on it since, including every bug found and fixed via real MAME testing along the way. See part 1 for the initial 6809 port and its core bug fixes that preceded this work.

Known bugs — resolved since the original checklist (continued)

- **New capability, not a bug fix: an assembly-level unit test framework added, with its first test (`TSTDUP`) written.** Gated by a new conditional-assembly flag, `UNITTESTS` , matching this file's established `IFEQ` convention (`0` = included, `1` = excluded entirely). Two insertion points: a `JSR TSTRUNNER` call at the true end of `COLDSTRT` (after `INITSERIAL` , before `JMP COLD` - at that point `U / S` are already valid, since `COLDSTRT` sets them at its very start, but nothing else - `APPVARS / DPHERE / CODEHERE / LATEST / BASE` - has been initialized yet), and the test framework's own body, living in what was previously unused ROM space right after `INOUT` 's shadow (`USROMSTRT` , `$C100`) - pure `FILL` padding before this existed. `UNITTESTS=1` reverts this block to exactly that padding, computed automatically via the `ROM` label rather than a fixed byte count (`FILL $FF,BASEDICT-ROM` , was `BASEDICT- USROMSTRT`), so it stays correct either way without needing hand-adjustment when the test code's size changes. Each test is independent by construction, per the stated design principle: it saves the data stack pointer (`U`) before touching anything and unconditionally restores it at the end regardless of pass or fail, so a failed assertion mid-test can never leave stack residue for the next test to inherit. Test scratch variables live at the very start of `APPVARS` - safe only because tests run strictly before `COLD` initializes `VARHERE` to that same address, which correctly and immediately re-purposes that space afterward; this is a real constraint worth remembering if the call site ever moves. Reporting is shared across every test via `TSTREPORT` rather than duplicated per test: prints the test's name (a counted string, via `COUNT + TYPE` - the same mechanism `BADWORD` itself uses to print a failing word), then "OK" or "FAIL", then a `CR` . `TSTDUP` verifies both the stack's contents after `DUP` (the duplicate and the original both equal a deliberately non- trivial pushed test value - not `0` , `1` , or `-1` , so a test that only appears to pass due to a special-cased value would be caught - and a sentinel pushed beneath is confirmed undisturbed) and the data stack pointer's movement (exactly one cell, 2 bytes - `DUP` 's own net effect, isolated from the

two setup pushes by capturing `U` immediately before and after the `JSR DUP` specifically, not around the whole test). `TSTRUNNER` and `TSTSTACK` are structured as extensible dispatchers (`JSR` a list of test groups, `JSR` a list of tests within each group) so more tests and more groups can be added without restructuring what's here. Verified across the full 2x2 matrix of `SERIALPOLL` / `UNITTESTS` combinations: zero duplicate symbols, dictionary chain intact (224 entries) in all four; byte-exact split-file reassembly. **MAME-CONFIRMED:** activated by the user, `TSTDUP` runs successfully and reports the correct message to the terminal - both the framework itself and `DUP`'s own test are now verified on real hardware, not just structurally. Whatever earlier concern prompted excluding this block was resolved by other, unrelated fixes made since it was first written.

- **Three real bugs found by inspection (not MAME) and fixed: `DEFER`, `2CONSTANT`, and `MARKER` all had the identical self-referential PFA bug already confirmed and fixed in `CONSTANT` via the MAME debugger.** Flagged by a parallel 68000 port of this codebase, which spotted the same construction pattern (`LDD CODEHERE / PSHU D / JSR CODECOMMA`, writing the PFA field's own current address into itself before the real value is appended two bytes later) recurring in these three defining words, unfixed. Verified each by inspection rather than deferred to hardware testing, since this is a deterministic compile-time address computation, not a timing/register- interaction issue - the same reasoning already used to confirm `VALUEW` was *not* affected when `CONSTANT` was first fixed. Traced both the compile-time construction and the runtime consumer for each: `DODEFER`'s `LDD ,X` would have jumped to the PFA's own address on execution of any `DEFER`'d word (crash); `DOMARKER`'s four fixed-offset reads (`,X / 2,X / 4,X / 6,X`) would have restored garbage `DPHERE / CODEHERE / VARHERE / LATEST` values - the most severe of the three, since using a `MARKER`-created word would corrupt live dictionary state, not just misbehave locally. Also checked, and confirmed genuinely unaffected rather than assumed safe by association: `DEFER@ / DEFER!` (go through `TOBODY`, a runtime offset computation from an existing `xt` - structurally immune, never compiles a self- referential field), `IS / ACTION-OF` (construct no PFA at all, just locate a word and call `DEFERSTORE / DEFERFETCH`), and `BUFFER:` (`BUFFERCOLON`, checked because it sits directly next to `2CONSTANT` in the source - uses the `VALUEW`-style pattern, PFA into `VARHERE` rather than `CODEHERE`, and `VALLOT` only advances `VARHERE` without ever writing through the PFA, so no self-reference exists there either). Fixed all three with the same `ADDD #2` pattern already established for `CONSTANT`. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME - the mechanism is identical to the already- hardware-confirmed `CONSTANT` bug, but these three fixes themselves are still unverified against real execution.

- Real bug found by inspection (not MAME) and fixed:** `2@ ( DFETCH )` read the two cells in the opposite order `2! ( DSTORE )` actually writes them, so a `2! 2@` round trip swapped the values. Flagged by the parallel 68000 port. Traced both routines precisely: `DSTORE` writes `x1` to the low address and `x2` to the high address; `DFETCH` was reading the high address first (pushing it deep) and the low address second (pushing it on top) - the reverse order. Confirmed via a concrete round-trip trace (`[x1, x2, a-addr]` before `2!` came back as `[x2, x1]` after `2@`), which is deterministic and fully verifiable from source - no MAME needed to establish it, same reasoning as the `MARKER / 2CONSTANT / DEFER` PFA bugs. `2!` itself was already correct and internally consistent; only `2@` 's read order was wrong. Fixed by swapping `DFETCH` 's two `LDD / PSBU` pairs to read low-then-high instead of high-then-low. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.
- Simplification (not a bug fix): the `PULS X / PSBS X` pair left over in `LEAVE` from the earlier `DOTEST` -class fix was a genuine no-op and has been removed.** Flagged by the parallel 68000 port, which dropped it rather than port a verified no-op forward. Confirmed by inspection: nothing runs between the pop and the push (no call, no other touch of `s` or `x`), so `x` held exactly what it held before the pop and `PSBS X` restored `s` to exactly where it already was - `RTS` would find the same return address there either way. Unlike `DOTEST` 's own deferred `PULS X` (genuinely load-bearing - its value feeds `LDD ,X / LEAX D,X` afterward), `LEAVE` never uses `x` 's value for anything; the pair was vestige of the pre-fix structure (where `PULS X` ran first and `PSBS X` was needed to restore `s` afterward), not something the fix itself required. `LEAVE` is now just `LDD #TRUEV / STD 6,S / RTS` . Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.
- Simplification (not a bug fix): `U. ( UDOT )` and `U.R ( UDOTR )` each had a literal pop-then-immediate-push-back on the value being formatted - a genuine no-op, removed.** Flagged by the parallel 68000 port, same class as the `LEAVE` finding immediately above. Confirmed by inspection: in both routines, `PULU D` immediately followed by `PSBU D` with nothing in between leaves `U` exactly as it was: the following `LDD #0 / PSBU D` (building the double-number `NUMGT` needs) pushes `0` on top of the value either way, whether or not it was ever popped and pushed back first. In `UDOTR` specifically, only the second pop/push pair (the value) was the no-op - the first (`PULU D / STD DRWIDTH` , consuming the width argument) is genuinely functional and was left untouched. `UDOT` now starts directly with `LDD #0` ; `UDOTR` now goes straight from the width pop to `LDD #0` . Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file

reassembly.

- **Real bug found by inspection (not MAME) and fixed: UNESCAPEW double-decremented UESRCLEN for a lone trailing backslash, underflowing the counter so the loop would run past the end of the intended input.** Flagged by the parallel 68000 port. Traced precisely: on encountering `\`, the escape-handling path decrements `UESRCLEN` once (correctly, accounting for the backslash itself) and checks `BEQ` - if the backslash was the *last* character in the input, `UESRCLEN` is now correctly `0`, but the branch fell into the shared `UEPLAIN` (plain-character) path, which decrements `UESRCLEN` a second time for a character that doesn't exist. `0 - 1` underflows to `$FFFF` (nonzero), so the next `UELOOP` pass's `BEQ UEDONE` never fires and the loop reads memory past the intended input. Confirmed the normal paths are unaffected: an ordinary plain character decrements exactly once (`BNE UEPLAIN` from the top); a real two-character escape like `\n` decrements once for the backslash and once more via the shared `UEEMIT` path for the second character - correctly 2 total for 2 characters consumed. The bug is specific to the lone-trailing-backslash case, where the escape branch's own decrement already accounted for the backslash before falling into a path that decrements again. Fixed by giving that one case its own landing point, `UELASTBS` - emits the backslash literally (`A` still holds it from the earlier `LDA ,X+`, so nothing needs reloading) without re-decrementing `UESRCLEN`, since it's already correctly `0`. `UEPLAIN` itself, still used by the normal plain-character path, was left completely untouched. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.
- **Design change (not a bug fix): UNLOOP is now a true no-op, resolving the UNLOOP / EXIT return-stack corruption from the prior turn.** Reasoning: traced `EXIT` (via `EXITUNLOOP`) across every scenario - no enclosing `DO` (compiled count `0`, the unwind loop never executes, falls straight through to a plain return), one enclosing `DO` with no prior manual `UNLOOP` (discards the correct, still-intact 8-byte frame). Both correct. The only broken combination was `UNLOOP` immediately before `EXIT`: `UNLOOP`'s own discard (6 of 8 bytes) left the frame partially gone, then `EXITUNLOOP` unconditionally discarded a full 8 more, overshooting into whatever sat beneath - typically the enclosing word's own caller's return address - branching into random memory on return, requiring a soft reboot to recover. Given `UNLOOP` has no legitimate use other than immediately preceding an exit from the definition, and `EXIT` already handles that correctly and automatically on its own, the conflict is resolved by removing the second, redundant discard mechanism rather than by making `EXIT`'s runtime detect how much of the frame remains (a more complex, riskier change). Checked for a new failure mode before applying: `UNLOOP` called without an immediate exit, then falling

through to `LOOP / +LOOP` again - already broken before this change too (`DOTEST / DOPLUSTEST` already expect the frame to still be there), so this introduces no new risk, and actually removes one instance of it. `UNLOOP` is now just `RTS`. Documentation updated to match: `UNLOOP` 's glossary entry rewritten to describe the no-op and why, kept in the dictionary for source compatibility with code written for other Forth systems where `EXIT` doesn't auto-unwind; `EXIT` 's own entry given a cross-reference noting `UNLOOP` isn't required beforehand. The Assembler Source appendix (Section 8.13, Control Flow) was updated to match the real source rather than left stale. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly; both new documentation passages confirmed present in the rendered PDF. Not yet confirmed via MAME.

- Real bug found by inspection (not MAME) and fixed: `CASE` pushed a 0 onto the compile-time `⋮` stack that nothing downstream ever consumed, causing every `CASE...ENDCASE` construct - even the simplest, with no `OF` clauses at all - to throw -22 at the closing `;`.** Reported via `: Nname CASE 0 OF ENDOF ENDCASE ;`; traced with the debugger to a `CSP` mismatch of exactly one stray cell. Verified the full compile-time lifecycle: `CASE` pushes `[0, TAGCASE]`; `OF` pushes `[placeholder-addr, TAGOF]`; `ENDOF` pops exactly those two and pushes its own `[NEWFLD, TAGENDOF]`; `ENDCASE` loops popping `TAGENDOF / NEWFLD` pairs until it sees `TAGCASE`, then stops. None of `OF / ENDOF / ENDCASE` ever read or pop past `TAGCASE` - the 0 `CASE` pushed beneath it is structurally unreachable by every consumer, leaving it permanently stranded one cell below where `CSP` (set by `:` before `CASE` ever runs) expects the depth to return to. `;` compares against `CSP` and correctly detects the mismatch every time. This implementation's `ENDCASE` uses a `TAGCASE`-scan approach entirely, with no counter involved anywhere in its actual logic - the stray 0 looks like a genuine leftover from an earlier, counter-based design that was never fully removed when the scan-based approach replaced it. Fixed by removing `CASEW` 's `LDD #0 / PSHU D` - now just `LDD #TAGCASE / PSHU D / RTS`. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.
- Correction to an earlier fix, found via a real MAME test (`MULT-TABLE`) and confirmed by inspection: `JWORD` 's offset, "fixed" to 10,s a few turns ago, was itself wrong.** `J` returned a fixed value (the outer loop's limit) on every iteration instead of the progressing index - `11 1 DO 11 1 DO I J * . LOOP CR LOOP` printed the same row ten times instead of a real multiplication table. The earlier fix incorrectly assumed `DOSETUP` 's own `JSR` -pushed return address persists on `s` after each nesting level - it doesn't, since `DOSETUP` 's own `RTS` pops and consumes it to jump into the loop body, so

it never actually occupies a slot for `J` to count past. Re-traced by counting every real push and pop across a nested `DO` rather than assuming a fixed frame size per level: after the inner `DOSETUP` finishes, `S` is `[inner-index@0][inner-limit@2][inner-leave@4][outer-index@6][outer-limit@8][outer-leave@10]`; `J`'s own `JSR` pushes one more cell on top, landing outer-index at `8`. Offset `10` (the earlier fix) lands on outer-limit instead - a value that never changes across outer iterations, exactly matching the observed symptom. `IWORD`'s own offset (`2`) was re-checked against this same corrected reasoning and confirmed still correct - it only ever sits one frame deep, not two, so the earlier error (specific to counting across a second, nested `DOSETUP` call) doesn't apply there. `JWORD` is now `LDD 8,S`. This is being logged transparently as a correction to a prior fix, not presented as if it were right the first time - worth being honest that even a change I was confident in and verified structurally at the time turned out to need real hardware testing to actually confirm. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- Real bug found via a real MAME test (`CREATE DOES>`) and fixed: `CREATE` had the same self-referential PFA-pointer bug already confirmed and fixed in `CONSTANT`, `DEFER`, `2CONSTANT`, and `MARKER` - but was never on the previously-flagged list, so it went unfixed until now.** Reported via `: ENUM CREATE , DOES> @ ; - DOES>` returned the address one cell before the value, actually stored, instead of the value's own address. Traced precisely: `CREATE` compiles `[JSR DODOES][FDB DOESRT0][FDB <CODEHERE's current value>]` - the same `LDD CODEHERE` - without- `+2` pattern as the other four, so the PFA-pointer field pointed at itself instead of two bytes further on, where, appends the real value next. Confirmed the runtime consequence through `DODOES` itself: it reads the *value stored at* the PFA-pointer field and pushes that (by design - this is how `CONSTANT`'s indirection works correctly, `DODOES` following a pointer to the real data rather than holding the data directly). With the field self-referencing, `DODOES` pushed the field's own address instead of following through to where the real value lives - exactly "the address before the 16-bit constant" instead of "the address 1 cell further on," matching the report precisely. Checked `VARIABLE` (the structurally similar routine immediately below `CREATE`, sharing the same `DOESRT0` trampoline setup) for the same bug rather than assume safety from resemblance alone - confirmed genuinely unaffected, since it uses the `VALUEW`-style pattern (`VARHERE`, a separate region) rather than `CODEHERE` self-reference. Also confirmed `SETDOES` needs no change - it patches a different field entirely (the behavior field, 2 bytes past `JSR DODOES`), fully independent of the PFA-pointer field this bug affects. Fixed with the same `ADDD #2` pattern already established for the other four. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- False alarm, resolved: a reported 2CONSTANT / 2@ discrepancy (666 777 2CONSTANT cc0 returning 666 \$700C instead of 666 777 ; a separate 2@ test on a 2VARIABLE appearing to read back in reverse order) turned out to be testing against a stale binary, not a live bug.** Traced the reported symptom carefully against the current source before this was resolved: `CREATE 's +2 fix` (this session, correcting the same self-referential PFA-pointer class as `CONSTANT / DEFER / 2CONSTANT / MARKER`) and `DFETCH 's` low-then-high read order (fixed to match `DSTORE` several turns earlier, flagged by the parallel 68000 port) both checked out correctly on paper against the reported symptom - neither reproduced the described behavior in static trace, which was the first signal something didn't add up. Confirmed only one `DFETCH` definition exists in the file (no stale duplicate). Resolved once the user rebuilt against the current `forth6809.asm` /split files and re-tested: the updated binary shows neither problem. Both are confirmed fixed by the prior `CREATE` and `DFETCH` changes already logged above - no additional code change was needed. Recorded here explicitly so the history is clear: this was a real, reasonable bug report at the time, not a mistake in reporting it - the fix had simply already landed in a turn the tested binary predated.
- Real, significant bug found via MAME debugger and fixed: QLOOP unconditionally reset STATE to interpret mode at the start of every line, silently breaking any colon definition split across more than one line of input.** Reported via `: test0 TRUE IF ." Hy" ELSE ." Hee" THEN ;` compiling correctly on one line but failing with a -22 (CSP mismatch) when split across several. `COLON` sets `STATE=-1` and `SEMI` sets it back to 0 (after checking `CSP`) - the correct, sole places `STATE` should change during normal operation. `QLOOP` 's own `LDD #0 / STD STATE` , run at the top of the per-line loop, was a third, redundant reset that fired every time `QUERY` read a new line - including lines in the middle of an still-open colon definition, regardless of whether `;` had actually been reached. `TRUE` (and everything after it on a continuation line) got interpreted and, for anything with a stack effect, executed instead of compiled - `TRUE` pushing `TRUEV` onto the data stack instead of being compiled, leaving a stray cell `CSP` correctly caught as a mismatch at `;` . Confirmed the fix's placement carefully rather than just deleting the reset outright: `QUIT` is only re-entered on cold boot or an uncaught error routing back through `ABORT` - confirmed ordinary successful lines loop back to `QLOOP` directly, never `QUIT` - so moving the reset to run once at `QUIT` (rather than removing it entirely) still correctly forces interpret state exactly when ANS's own `QUIT` semantics call for it (including recovering from an error that aborts an unfinished colon definition, which deleting the reset outright would have left permanently stuck in compile mode), just not on every single line of an otherwise-uninterrupted session. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- Real bug found via MAME testing and fixed, surfaced directly by the prior QLOOP/STATE fix:** `QOK` skipped the `CR` echo together with the "ok" message whenever `STATE` was nonzero, silently dropping the line ending for every line of a multi-line colon definition after the first. Once the QLOOP fix let multi-line colon definitions actually compile successfully, the echoed source no longer resembled what was typed - every continuation line ran together onto one visual line, since the terminal never saw the `CR` the user had genuinely pressed. Traced to `QOK: LDD STATE / BNE QLOOP` running *before* `JSR CRW` - while compiling, the branch away skipped both the `CR` and the "ok" message together, when only the "ok" message is actually supposed to be conditional (correctly not shown mid-definition). The error path just above `QOK` already got this right (unconditional `JSR CRW` before printing the error message) - only the success path had the bug. Fixed by moving `JSR CRW` to run unconditionally, first, with the `STATE` check (and the "ok" message it guards) coming after. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.
- Real bug found via MAME debugger and fixed:** `FILL` corrupted its own remaining-byte count on the very first iteration, running far past the requested length. Same register-clobber class as `CCALL`'s original bug: `FILLOOP` loaded `D` with the count, then `LDA FILLCHR` (needed for the fill byte) overwrote `A - D`'s high byte - before `SUBD #1` decremented what it believed was still the true count. The high byte took on the fill character's value instead, so the count almost never reached zero at the intended point. Fixed by keeping the count in `Y` instead of round-tripping it through `D` /memory each iteration - `Y` is untouched by loading the fill character into `A`, so the conflict is structurally impossible now, not just avoided by careful ordering. Also addressed the redundant scratch-register usage flagged alongside the bug report: the target address is now kept directly in `X` (via `TFR D,X`) rather than round-tripping through `FILLADDR` first, matching the same treatment given to the count. `FILLCNT` / `FILLADDR` (aliases for the shared `MVCNT` / `MVDST` scratch cells) became genuinely unused once the loop no longer touches memory for them - confirmed via search before removing, and removed along with their comment-block entries, consistent with how `DICTTOP` was handled earlier this session. `MVCNT` / `MVDST` / `MVSRC` themselves are untouched, since `HSLEN` / `HSADDR` / `MRESULT` / `FILLCHR` still alias them for other routines. `ERASEW` (which `JMP`s into `FILLW`) confirmed unaffected, since only the internal loop changed, not the external calling convention. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.
- Real bug found via MAME debugger and fixed:** `HEXBYTE` (used by `DUMP`'s hex column) read from `MSCR+1` instead of `MSCR`, formatting whatever stale byte

happened to be at the never-written address instead of the byte actually passed in. `STB MSCR` writes one byte, at `MSCR` itself; both subsequent reads (`LDB MSCR+1` , once for the high nibble via four `LSRB` s, once for the low nibble via `ANDB #$0F`) used the wrong address - a cell `HEXBYTE` never writes at all, so both nibbles were derived from stale scratch content rather than the real value. Reported via `MOVE + DUMP` : a fill character of `$7B` showed as `$03` in the hex column, while the ASCII column (a separate code path, unaffected) correctly showed `}` . Confirmed `MSCR` is shared, general-purpose scratch used pervasively elsewhere via full 16-bit `STD / LDD - HEXBYTE` 's single-byte `STB / LDB` pattern was the outlier, and nothing else depends on `MSCR+1` holding anything meaningful at the point `HEXBYTE` runs. Fixed by changing both `LDB MSCR+1` occurrences to `LDB MSCR` , reading from where the byte was actually stored. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- Formatting change, not a bug fix: `DUMP` now emits a leading `CR` before its first line of output, matching the trailing `CR DULEND` already emits after every line (including the last).** Requested for consistent vertical alignment - without a leading `CR` , the first line started wherever the cursor already was (e.g. right after the echoed command itself), while every subsequent line correctly started at column 0 via the existing trailing `CR` . Added `JSR CRW` as the first thing `DUMPW` does, right after popping its two arguments and before the per-line loop begins. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.
- Real bug found via MAME debugger and fixed: `CMOVE` never terminated - same register-clobber class as `FILL` 's bug earlier this session.** `CMVLOOP` loaded `D` with the count, then `LDA ,X+` (needed to fetch the byte being copied) clobbered `A - D` 's high byte - before `SUBD #1` decremented what it believed was still the true count, so the terminating condition almost never fired. Checked the two adjacent, structurally similar routines rather than assume they shared the bug: `CMOVEGT` (the reverse-direction, overlap-safe variant) is genuinely unaffected, since its own copy (`LDA ,X`) happens *before* `LDD MVCNT` reloads the count fresh from memory each iteration, rather than after - the clobber happens, but the count is correctly reloaded afterward, overwriting it, not corrupted by it. `MOVEW` has no copy loop of its own at all - it's purely a dispatcher choosing `CMOVEW` or `CMOVEGT` based on overlap direction, so it's correct automatically once `CMOVEW` is. Fixed `CMOVEW` differently than `FILL` : unlike `FILL` (which only needed one address, leaving a register free for the count), `CMOVEW` already commits both `X` (source) and `Y` (destination) to the two addresses, leaving no spare 16-bit register to hold the count in the same way. Fixed instead by reordering - decrement and store the count while `D` still holds the true value, before the byte copy is free to clobber `A` .

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly.

- **Real design gap found via MAME testing and fixed: SUBSTITUTE never returned the ANS-required substitution count, and its core algorithm was a plain substring search-and-replace on the bare registered name rather than the ANS % -delimited scan.** Confirmed against the actual spec (forth-standard.org/standard/string/SUBSTITUTE and [.../REPLACES](http://forth-standard.org/standard/string/REPLACES), both fetched and read in full): SUBSTITUTE must scan for text between % (ASCII \$25) delimiter pairs specifically - %% collapses to a single % (count unchanged); a name matching the REPLACES registration has the *entire* %name% span, delimiters included, replaced by the substitution text (count incremented); a non-matching name is passed through unchanged, delimiters and all (count unchanged); a trailing % with no closing delimiter passes the residue through unchanged. Rewrote SUBSTITUTEW completely to implement this scan, no longer calling SEARCHW at all - confirmed SEARCHW still has its own dictionary header (H_SEARCHW) and remains a legitimate standalone word (SEARCH) independent of this change. Reused the existing COMPAREW for the name-matching step rather than writing a new comparison loop. GLOBALS is fully packed (256/256 bytes, confirmed) - the rewrite reuses MSCR / MSCR2 / MSCR3 / MSCR4 (confirmed untouched by SUBCOPY) for the new algorithm's scan state rather than adding dedicated cells. Also fixed the missing third return value - SUBSTITUTEW now pushes the substitution count (0 or positive on success) alongside the destination address and length, matching (addr1 len1 addr2 len2 -- addr2 len3 n) . The destination-overflow behavior (THROW -1) was left unchanged - the ANS spec explicitly permits throwing for negative- n cases in the standard THROW-code table.

```
**Separately identified, and deliberately NOT fixed on request:**
`S"`, when interpreted (not compiled - i.e. typed directly at
the prompt rather than inside a colon definition), always
writes into the same fixed, shared buffer (`SIBUF`), returning
that same address every time. `REPLACES` only stores the
addresses it's given, not copies of the text - so two `S"
calls used to register a name/value pair (and a third `S"
for `SUBSTITUTE`'s own source string) all silently overwrite the
same buffer, corrupting earlier registrations before
`SUBSTITUTE` ever runs. Traced precisely enough to reproduce a
reported garbled MAME output (`Dancinging, %girl%!`) by hand,
confirming the exact mechanism. User's explicit decision: don't
add dedicated copy-on-register storage to `REPLACES` - instead,
wrap each string in its own colon definition (`: name S"
text" ;`) for stable, independent storage, and document this
requirement in the glossary instead of changing the code. Done -
see the `REPLACES`/`SUBSTITUTE` glossary entries in
```

ClaudeForth.docx, now also corrected to match the real ANS stack effects and behavior rather than the previous, inaccurate descriptions.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly. Given the size of the rewrite, long branches (`LBRA`/`LBHS`/`LBNE`) were used liberally for loop-backs and larger jumps as a safety margin. ****MAME-CONFIRMED**** (single and double substitution both correct: `S" Dancing, %girl%!" poem 80 SUBSTITUTE` -> "Dancing, Agnetha!"; `S" Two dancing girls- %girl%! %girl%!" poem 80 SUBSTITUTE` -> "Two dancing girls- Agnetha! Agnetha!", with `.S` showing `2` as the substitution count) - real assembly and execution confirm the long-branch margin was sufficient and the `%`-delimiter algorithm itself is correct.

- **Verification task, requested and completed: confirmed PAD's current offset satisfies all three ANS transient-region minimums (forth-standard.org/standard/usage, section 3.3.3.6), added named equates documenting them, and redesigned S" to use PAD instead of a fixed, retired SIBUF.** Verification result: already satisfied, no numeric adjustment required. - PAD's own scratch region (≥ 84 chars): `PADW` computes `PAD = CODEHERE + 84` - meets the minimum exactly, with the usual ANS-expected transient/dynamic behavior (moves as `CODEHERE` advances). - `WORD`'s transient region (≥ 33 chars): `WORDBUF` spans `$01DA` to `$01FA` (up to where the now-retired `SIBUF` began) - computed precisely via address subtraction (`$01FB - $01DA = 33`), meeting the minimum exactly. Unrelated to the PAD offset - this system's `WORD` uses its own separate, fixed buffer rather than `HERE`-relative storage. - Pictured numeric output buffer ($\geq (2*16)+2 = 34$ chars): traced `LTNUM` ("`<#`") and `HOLD` precisely - `HLD` anchors to `PADW`'s own return value, and `HOLD` decrements *before* storing, so this buffer grows downward from PAD into the `CODEHERE`-to-PAD gap, not into PAD's own region above it. $34 \leq 84$, comfortably satisfied with 50 bytes of margin.

Added `PADMINSIZE`/`WORDMINSIZE`/`HOLDMINSIZE`/`PADOFFSET` as named, documented equates, replacing the bare `84` that used to sit directly inside `PADW`.

Redesigned `S``'s interpreted-mode path (`SQINTERP`) to write into PAD (computed fresh via `PADW` on every call) instead of the old, fixed `SIBUF` - giving interpreted `S`` access to PAD's full 84-character region instead of the previous fixed 32-character limit. Confirmed `SIBUF` was used exclusively by

`S`` (matching the user's own stated assumption, verified by search) before retiring its `EQU` definition; left its old 32-byte address range (\$01FB-\$021A) unclaimed rather than shifting `APPVARS` down to reclaim it, to avoid any risk to the rest of this carefully-verified memory map for the sake of 32 bytes on a system with headroom to spare.

****Difficulties check, as requested:**** confirmed no conflict with the pictured-numeric-output buffer (opposite growth directions from the same PAD anchor - traced precisely, they don't overlap). Confirmed a genuine, expected trade-off instead: per ANS's own allowance ("non-standard words... may use PAD, but such use shall be documented"), PAD is now shared between the user's own general-purpose use and interpreted S" - using one right after the other will overwrite one with the other, same as any other transient-region interaction the standard itself anticipates. Also identified and documented a subtler point: PAD redesign does NOT eliminate the need for the user's own colon-definition-wrapping workaround (established two turns ago) - PAD's address only changes when something is compiled or allocated between calls, so two interpreted S" calls in a row (e.g. REPLACES's two arguments) still collide, just via PAD instead of SIBUF. Compiled S" (inside a colon definition) never touches PAD at all, writing directly into the definition's own permanent storage instead - this remains the correct approach for anything needing more than one string to coexist. Confirmed `SIBUF` referenced nowhere else in the file before retiring it.

Glossary entries for `PAD`, `S``, and `REPLACES` updated to match - the `REPLACES` entry's prior `SIBUF`-specific wording (from the previous SUBSTITUTE fix) was stale and has been corrected to describe the current PAD-based mechanism instead.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; both updated Assembler Source appendix subsections (8.1 Memory Map and 8.20 String Words) and all three updated glossary entries confirmed present in the rendered PDF. ****MAME-CONFIRMED**** as part of the full substitution chain below - the PAD-based interpreted-mode S" storage this entry introduced is exercised correctly by every scenario in that confirmation.

- **Confirmed and fixed at the user's request: WORD had its own, separate, fixed 31-character-capped buffer (WORDBUF), entirely unrelated to CODEHERE or PAD - the PAD-based S" redesign two turns ago didn't help long S" strings because**

WORD truncates during parsing, an earlier stage than either S''s storage path.

Redesigned WORD to scan directly into the CODEHERE-to- PAD gap instead, matching the traditional fig-Forth layout (WORD's buffer at the CODEHERE end, the pictured numeric output buffer at the PAD end, growing toward each other) - added

WORDMAXCHARS as a named constant reserving HOLDMINSIZE bytes at the PAD end for that buffer, so the two don't collide even though ANS itself would permit the overlap (3.3.3.6). WORDBUF confirmed genuinely unused before retiring it, same treatment as SIBUF 's retirement two turns ago.

```
**Serious collision bug caught and fixed while verifying this
redesign, before delivery - not reported by the user, found by
tracing the change's own consequences.** `SQUOTE` (S"),
`DOTQUOTE` ("."), and `AQSTOK` (ABORT") all compile a 3-byte
runtime trampoline (`JSR CCALL`) directly at `CODEHERE` *after*
calling `WORD` but *before* copying the parsed text - with
`WORD` now writing that same text at `CODEHERE` too, the
trampoline would overwrite the first bytes of the very text
being staged, before the copy loop ever ran. Checked every
other caller of `WORD` in the file systematically (`TO`, `IS`,
`ACTION-OF`, the main interpreter loop, ```, `POSTPONE`,
`[']`, `CHAR`, `[CHAR]`, `( `) - confirmed none of them share
this pattern, since they either resolve to an xt via `FIND`
immediately (never needing the raw text again) or extract a
single character into a register before any compilation runs.
`HEADER` (used by every defining word) is also unaffected,
since it writes into `DPHERE`, a completely separate region
from `CODEHERE`. Fixed all three affected words identically:
reserve 3 bytes ahead of `CODEHERE` before calling `WORD`, so
the parsed text naturally lands exactly where it needs to end
up, with the trampoline safely compiled into the reserved gap
in front of it rather than on top of it - the existing copy
loop becomes a harmless no-op (reading and writing the same
address) rather than needing to move anything.
```

Caught a further, related overflow while finishing this: the 3-byte reservation itself pushes the compiled path's text 3 bytes further into the CODEHERE-to-PAD gap than plain WORD-parsing does (e.g. a dictionary name for FIND) - the original `WORDMAXCHARS=49` would have let the worst case (S"/./ABORT", at the cap) overflow 3 bytes into the pictured-numeric buffer's reserved space. Corrected to `WORDMAXCHARS=46`, sized uniformly for the worst case across all callers rather than tracking a different effective limit per caller - simpler and safer.

Also caught and fixed a transcription error in my own edit

before it was ever delivered: while inserting an explanatory comment into WORD's scan loop, the `BEQ ENDW` branch instruction immediately following the length-cap comparison was accidentally dropped, which would have made the length check compare but never actually branch - re-verified against the live file and corrected before proceeding further.

Glossary entries for `WORD`, `S``, `.``, and `ABORT`` updated to match - `S``'s entry from the previous turn specifically is corrected here, since its claimed "up to PADMINISIZE (84 characters)" limit was always inaccurate: WORD's own separate cap (31 at the time, now 46) was still the binding constraint, never 84.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; all four updated Assembler Source appendix subsections (8.1 Memory Map, 8.10 Outer Interpreter, 8.14 Compiling Words, 8.20 String Words) and all four updated glossary entries confirmed present in the rendered PDF.

****MAME-CONFIRMED across every scenario that mattered**, including the highest-risk one: `: expand S" Two dancing girls- %girl%! %girl%!" poem 80 SUBSTITUTE DROP SPACE TYPE ;` then `expand` -> "Two dancing girls- Agnetha! Agnetha!" - this specifically exercises the compiled-mode collision fix in `SQUOTE` (the reserve-3-bytes-before-calling-WORD pattern), confirming the runtime trampoline and the parsed text no longer overwrite each other. Also confirmed: the 34-character (and longer, 36-character) strings that originally exposed WORD's old 31-character cap now parse in full, in both interpreted (`S" ..." poem 80 SUBSTITUTE`) and colon-definition-wrapped (`: template S" ..." ;` then `template poem 80 SUBSTITUTE`) forms, with correct output in every case.**

`DOTQUOTE` (`.```) separately MAME-confirmed too - its own reserve-3-bytes fix, applied identically to SQUOTE's but as a distinct edit, exercised independently: `: .message ." <text> " ;` then `.message` types the text correctly, unmodified.

`AQSTOK` (`ABORT``) also separately MAME-confirmed, completing independent confirmation of all three fixed words: `: over18 18 < ABORT" Must be > 18" ;` - `3 over18` correctly types the message and throws -2 (the true-flag path); `18 over18` correctly falls through with no message and no throw (the false-flag path). Both branches of ABORT``'s own conditional

logic exercised, not just the message-compilation mechanics shared with SQUOTE/DOTQUOTE.

- **Real, significant bug found via MAME testing and fixed: UNESCAPE 's argument popping was completely wrong - only one of its three arguments was ever popped, and even that one went into the wrong variable.** ANS signature is (c-addr1 u1 c-addr2 -- c-addr2 u2) . The code did PULU D / STD UESRCLEN (storing the popped c-addr2 , the real destination address, into the *source-length* variable), then LDD ,U - a peek, not a pop - into *both* UEADDR and UEDST (misreading u1 , the real source length, as an address, and never popping it off the stack at all). c-addr1 , the actual source address, was never touched - left sitting on the stack the entire time. Net effect: UEADDR / UEDST both ended up holding a small length value misread as an address, the copy loop read and wrote through whatever garbage that pointed at (low memory, within GLOBALS itself, given typical string lengths), and the real source/dest addresses were either misfiled or left stranded on the stack - matching the reported symptom precisely (two addresses left on the stack, no real count, buffer untouched). Also fixed UEDONE 's return: it used to do LDD UEOUTLEN / STD ,U , overwriting whatever was left on top of the stack rather than pushing both required return values; now correctly pushes c-addr2 (destination address) then u2 (actual unescaped length), matching the ANS stack effect. Confirmed UNESCAPEW is referenced nowhere else in the file before applying the fix. The escape-processing loop itself (the \n / \t / \ / \" handling, including the lone-trailing-backslash fix from earlier this session) was untouched and unaffected - it operates correctly once its input variables are actually populated correctly. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.
- **Major finding, well beyond the prior turn's argument-order fix: UNESCAPE 's entire algorithm was wrong from the start, not just its argument popping.** After the previous fix corrected the broken argument order, MAME testing showed doubling a single % wasn't happening at all, and the returned count didn't grow to account for it. Investigated against the actual spec rather than assuming a smaller bug within the existing logic - fetched both forth-standard.org/standard/string/UNESCAPE and the complang.tuwien.ac.at ANS Forth reference text, which includes the canonical reference implementation. Confirmed: UNESCAPE has nothing to do with backslash escape sequences (\n , \t , \ / \") at all - the entire pre-existing implementation (predating this session) was decoding a plausible-sounding but entirely wrong operation. The real ANS UNESCAPE replaces every % character with two % characters and passes everything else through unchanged, specifically so that a literal % in text survives an eventual SUBSTITUTE pass unchanged ("If you pass a string through UNESCAPE and

then SUBSTITUTE, you get the original string"). Replaced the entire routine body accordingly - confirmed the old backslash-handling labels (UELASTBS / UEPLAIN / UECKT / UECKBS / UEEMIT) were internal to this one routine before removing them. New output length is computed directly as the final write pointer minus the starting destination address, correct regardless of how many % characters were doubled, rather than incrementally tracked per-character the way the old (wrong) loop did it. Manually traced the new algorithm against the ANS reference's own test case ("aaa%bbb" UNESCAPE should equal "aaa%bbb") before delivery - confirmed matching exactly, including the length (7 characters in, 8 out). No destination-buffer-overflow check added, matching the ANS reference implementation, which also has none (an ambiguous condition per the standard, not a required error case). Glossary entry corrected completely - the prior entry's stack effect, description, and even its "result length is always <= input length" side-effect claim were all wrong, describing the old, incorrect algorithm rather than the real one (output can now be longer than input, by design). Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLX UNITTESTS combinations; byte-exact split-file reassembly; updated Assembler Source appendix subsection (8.20 String Words) and the corrected glossary entry both confirmed present in the rendered PDF. **MAME-CONFIRMED:** correctly escapes % characters as required.

- **Feature gap confirmed and addressed: the text interpreter had no support for ANS's double-number input convention (trailing ') - "123." was silently treated identically to "123", leaving no way to get a genuine ud1 onto the stack for use with # and friends.** Confirmed against the precise spec text (forth-standard.org/standard/usage#usage:numbers, 3.4.1.3: "When the text interpreter processes a number... that is immediately followed by a decimal point... the text interpreter shall convert it to a double-cell number. For example, entering DECIMAL 1234 leaves the single-cell number 1234 on the stack, and entering DECIMAL 1234. leaves the double-cell number 1234 0 on the stack."). Traced NUMBERQ and found it already accumulated a full 32-bit value internally (UDHI : UDLO , via NUMLOOP) for every number parsed, single or not - but only ever returned the low cell, discarding UDHI unconditionally, and had no detection of a trailing ' at all. Also found the existing negation only negated the low 16 bits - harmless while only ever returning a single cell, but wrong once a genuine 32-bit value needed returning.

Redesigned `NUMBERQ`: detects a trailing '.' early (excluding it from the digit count before `NUMLOOP` runs), performs full 32-bit two's-complement negation (low cell negated first, any carry out of that addition propagated into the high cell, manually traced against both a normal case and the zero-wraps-to-zero edge case before delivery), and re-derives the double-

vs-single decision independently at the very end of the routine by re-reading from `CADDR` and the original count - both confirmed unchanged throughout the routine's entire execution, including through `NUMLOOP`/`UDMULADD` - rather than remembering the decision in a new persistent flag. This was a deliberate choice: `GLOBALS` is fully packed at 256/256 bytes (the 6809's own direct-page addressing limit, not an arbitrary number), so avoiding a new flag byte avoided touching that boundary at all. Return convention on success now distinguishes double (pushes low, high, then success code `1`) from single (pushes value, then the original `-1`, completely unchanged) - chosen so `TRYNUM` doesn't need a separate flag either, and so single-number behavior is provably identical to before by construction, not just by testing.

Updated `TRYNUM` to handle both codes: on double success while compiling, `UDHI` (on top of `U`) is stashed in `MSCR4` so `UDLO` can be compiled as the first of two `LIT` literals, then `UDHI` as the second - ensuring they execute in the order needed to land correctly on the stack at runtime (low deep, high on top, per 3.1.4.1's "the cell containing the most significant part... shall be above the cell containing the least significant part"). While interpreting, both cells are already correctly placed by `NUMBERQ` and need no further handling.

Confirmed `TONUMBER` (`>NUMBER`) is untouched - it calls the same shared `NUMLOOP`, but neither its own code nor `NUMLOOP` itself needed any change; only `NUMBERQ`'s own wrapper logic and `TRYNUM`'s consumption of it were modified.

Manually traced three cases by hand before delivery, beyond structural verification alone: "123." (interpreted) -> [123, 0] correctly, matching the ANS example exactly; "123" (no dot) -> [123, -1] internally, correctly falling through to the original, unchanged single-cell path, confirming backward compatibility; "-123." -> UDHI=\$FFFF/UDLO=\$FF85, confirmed by hand to equal \$FFFFFF85 (i.e. -123 as a 32-bit value).

Glossary entries for `<#` and `#` updated with a note on how to get a proper udl onto the stack from typed input, since that was the specific gap originally reported.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; updated Assembler Source appendix subsection (8.10 Outer Interpreter) and both updated glossary entries confirmed present in the

rendered PDF. Not yet confirmed via MAME - this is a substantial, multi-path change (single/double x interpret/compile x positive/negative all interact) and deserves thorough real-hardware testing across that full matrix before being considered settled.

- **Two real bugs found via a very thorough MAME test matrix (interpret and compile mode, both signs, single and double, cumulative .s traces across many lines) and fixed - both in the 32-bit negation logic delivered last turn for double-number input, and both only visible on negative doubles specifically.** Positive cases (123. - > 0 123, 123456. -> 1 -7616) were already confirmed correct by the trace, matching the design exactly - the negation path (which only positive numbers skip entirely) was where the actual problems were.

****Bug 1**:** the 6809's `COM` instruction (one's complement) unconditionally sets the carry flag to 1, regardless of its operand - a documented CPU quirk. The negation code computed the real carry from the low cell's `ADDD #1`, but then ran `COMA`/`COMB` on the high cell before checking that carry - and `COMB` silently overwrote it with its own, always-set value, so the intended "only propagate carry into the high cell when the low cell actually overflowed" check never worked. `-123.` returned high cell `0` instead of `-1`; `-123456.` returned `-1` instead of `-2` - both exactly matching "always add 1" regardless of the true carry, not the intended conditional behavior. Fixed by saving the true carry via `PSHS CC` immediately after the low-cell addition, before the high-cell `COM` can destroy it, then `PULS CC` to restore it right before the branch that depends on it.

****Bug 2,** only exposed by fixing Bug 1**:

with the carry check now actually working, the branch it guarded (`BCC`) turned out to jump all the way past `STD UDHI` itself when there's no carry - not just past the `+1`. The complemented high-cell value was computed in D but never actually stored back to UDHI in that case, leaving it at its stale, pre-negation value. This was completely masked by Bug 1 in practice - since carry was always corrupted to "set," the branch was never actually taken, so the store always ran (with the wrong value, per Bug 1, but it did run). Fixing Bug 1 alone would have made this second, previously-latent bug live and produced a different set of wrong answers. Restructured so the store always runs, with the conditional `+1` applied before it rather than gating whether it happens at all.

Manually re-traced both negative-double test cases against the fully corrected code before delivery: ``-123.`` -> carry clear on the low cell, branch skips only the ``+1``, raw complement ``$FFFF`` gets stored -> high cell ``-1``, correct. ``-123456.`` -> same pattern, raw complement ``$FFFE`` stored -> high cell ``-2``, correct (matches the true 32-bit value ``$FFFE1DC0`` computed by hand). Also re-traced the carry-set path (negating 0, as a sanity check) to confirm that branch still correctly falls through to the ``+1`` before storing.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four ``SERIALPOLL`x`UNITTESTS`` combinations; byte-exact split-file reassembly; updated Assembler Source appendix subsection (8.10 Outer Interpreter) confirmed present in the rendered PDF. Glossary entries (``<#``, ``#``) needed no further changes - this fix doesn't alter the stack effect or usage, only the underlying correctness.

****MAME-CONFIRMED****, both interpret and compile mode, across the full t0-t7 matrix: ``-123.`` -> high ``-1`` / low ``-123``; ``-123456.`` -> high ``-2`` / low ``7616`` - both matching the hand-traced prediction from the prior turn exactly (high ``$FFFF``/``$FFFE`` respectively, confirmed independently against ``$FFFFFFF85`` and ``$FFFE1DC0``). Positive cases (``123.`` -> ``0 123``, ``123456.`` -> ``1 -7616``) and single-cell cases (unaffected by either bug, per the negation code being skipped entirely) also confirmed correct. Both the carry-corruption bug and the masked missing-store bug are now verified fixed on real hardware, not just by static hand-tracing.

- **Two explicit, requested adjustments applied: `BASECODE` and `BASEDICT` origins shifted down \$40 (64 bytes) each, to make room for growth since they were last positioned; `UNITTESTS` set to 1 (excluded) since unit tests don't work yet and resolution has been postponed.** `BASECODE : $E02A -> $DFEA`. `BASEDICT : $D83F -> $D7FF`. Applied exactly as requested, without independently recomputing the resulting gaps/overlaps against `INITCODE` and each other - this region's comments have historically tracked those in precise byte counts, but doing so accurately requires real assembled sizes (`BASECODESIZE / BASEDICTSIZE` , computed via `*-start` at each section's end), which structural verification here cannot determine - it checks label integrity and dictionary-chain structure, not location-counter arithmetic. The comments at both `EQU` s now say this explicitly rather than presenting a static estimate as if it were a real assembled count - confirm the resulting gaps/overlaps on real assembly.

Verified: zero duplicate symbols, dictionary chain still 224

entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations (checked with both `UNITTESTS` values regardless of the file's own new default); byte-exact split-file reassembly.

- **Real bug found via MAME debugger and fixed: `S>D` performed no sign extension at all - the same register-flag class of bug already documented once this session in `TRYNUM`.** `PULU` does not affect condition codes on genuine 6809 - `STOD` did `PULU D` then immediately `BPL SDPOS`, testing whatever flags an unrelated, earlier instruction happened to leave set, not the sign of the value just popped. `-123 S>D` returned high cell 0 instead of -1. Fixed by adding an explicit `TSTA` (testing D's high byte, whose bit 7 reflects the full 16-bit value's sign) immediately before the branch that depends on it. Confirmed `STOD` is referenced only via its own dictionary header before applying the fix - nothing else calls it directly. Checked the neighboring routines (`DTOS`, `DMAW`) for the same pattern rather than assume they're clean by proximity alone - both confirmed unaffected: `DTOS` has no branch at all, and `DMAW` already uses an explicit `CMPD` before its own branch. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.
- **Real bug found via MAME testing and fixed: `SIGN` had the same `PULU` -doesn't-set-flags defect already found in `STOD` - plus a separate, genuine usage-pattern issue in the reported test definition, diagnosed but not a code bug.** Reported via : `format DUP ABS S>D <# #S SIGN #> TYPE` ; never printing a minus sign for either sign of input.

****Code bug**:** ``SIGN`` did ``PULU D`` then immediately ``BPL SIGNDONE`` - ``PULU`` doesn't affect condition codes on genuine 6809, so the branch tested whatever flags an unrelated, earlier instruction happened to leave set, not the sign of the value just popped. Fixed with the same ``TSTA`` pattern used for ``STOD``. Checked all four existing internal callers (``DOT`/`DOTR`/`DDOT`/`DDOTR``, i.e. ``.`/`.`.R`/`.`.D`/`.`.D.R``) before concluding anything about their own correctness - all four turned out to be accidentally unaffected, since each runs ``LDD SAVEN`` (a flag-setting load of the true signed value) immediately before ``PSHU D`/`JSR SIGN``, and ``PSHU`` doesn't disturb flags. This was fragile, caller-side luck rather than ``SIGN`` being correct on its own - confirmed precisely by the fact that a direct, standalone use (the reported test word) had no such protection and failed outright.

****Separate, non-code issue identified in the reported definition itself**:** even with ``SIGN`` fixed, ``DUP ABS S>D <#`

#S SIGN #>` would still never add a minus sign, traced precisely - `#S` fully consumes its `ud` down to `0 0`, so whatever sits on top of the stack immediately after `#S` is always `0`, genuinely not negative, regardless of the original number's sign. The true signed value (left underneath by `DUP ABS`) never reaches the top of the stack in that sequence at all. Standard idiom needs the signed value stashed on the return stack and restored right before `SIGN`: `DUP >R ABS S>D <# #S R> SIGN #>`. Not a defect - flagged so the fix to `SIGN` itself isn't mistaken for resolving this specific definition's own behavior too.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.

- Investigated, no code change made: an apparent CATCH success-path anomaly (0 process , i.e. no THROW induced, showing nothing on .s instead of the expected 0) turned out not to reproduce on retest, most likely a MAME single-stepping artifact rather than a genuine bug.** Traced CATCH 's success path precisely (the LEAS 2,S / PULS D / STD HANDLER / LDD #0 / PSHU D / RTS sequence after a normal, non-throwing return) and separately OBRANCH (ZBRANCH , the specific code path taken since 0<> of 0 is false in the reported test, meaning IF 's body is skipped) - both confirmed structurally balanced and correct via static tracing, with no bug found on paper. User's own leading hypothesis, given this: MAME's 6809 core "jumping into other routines on single-stepping instructions that were not branches/subroutines" - i.e. a debugger artifact from stepping through raw memory, not a defect in the actual compiled code. Consistent with the static trace finding nothing wrong and the issue not reproducing on a subsequent, real (non-single-stepped) run - a genuine code bug wouldn't self- resolve between otherwise-identical test runs. Confirmed working via a complete, realistic follow-up test (loge / process , chaining CATCH success and failure paths together with SIGN and HOLDS , both fixed earlier this session): 0 process -> OK. (success path correct); 5 process -> Err: -4000 (failure path correct, SIGN / HOLDS interaction correct end to end). No source change was needed or made.
- Real bug found via MAME testing and fixed: ((LPAREN) left a stray address on the stack after every comment.** WORD always pushes a c-addr (the parsed text's address) at the end of its own execution, per its established calling convention - LPAREN only ever needed WORD 's side effect of advancing >IN past the closing) , never the address itself, but never popped it either, going straight to RTS . Explained both reported symptoms precisely: a stray value left on the stack after a comment in interpret mode ((

Math for */) .S -> 28888 , a plausible CODEHERE value within the APPCODE region, matching what WORD 's redesign earlier this session would push), and a -22 CSP mismatch for any colon definition containing a comment, since the leftover throws off the compile-time stack-depth check between : and ; . Confirmed LPAREN is referenced only via its own dictionary header, and confirmed there's no . ((a related, commonly-paired comment word) in this system that might share the same pattern - (is the only comment word affected. Fixed with a single PULU D to discard the address, immediately after JSR WORD returns. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly. **MAME- CONFIRMED:** (now works correctly.

`\` (`BACKSLASH`) separately tested alongside `(` as part of this same comment-words sweep - confirmed already working correctly, no code change needed.

- **Verification, no code change: confirmed WORDLISTS (16.3.2, Search-Order word set) already, correctly returns false from ENVIRONMENT?** . Checked the full ENVTABLE (four entries: /COUNTED-STRING , MAX-N , MAX-U , ADDRESS-UNIT-BITS) - WORDLISTS is genuinely absent, not an oversight, since this system doesn't implement the Search-Order word set at all. Confirmed against ENVIRONMENT? 's own spec (6.1.1345): "If the system treats the attribute as unknown, the returned flag is false" - falling through to ENVNOTFOUND already produces exactly that, correctly and completely, with no dispatcher work needed. Clarified the pre-existing comment above ENVQUERY , which previously lumped WORDLISTS together with MAX-D / MAX-UD / FLOORED as all equally "need[ing] dispatcher extensions not yet built" - that's true for MAX-D / MAX-UD (genuinely incomplete: both are double-cell values, but ENVFOUND only ever pushes one cell before the TRUE flag) but not for WORDLISTS , whose correct answer is simply "unknown," which absence from the table already provides. FLOORED kept separately tracked as its own, still-open item - unlike WORDLISTS it's a core environmental query, not tied to an unimplemented optional word set, so this system likely has a genuine, definite division behavior it could report - "false by omission" isn't necessarily the right answer for it the way it is for WORDLISTS . Not investigated further here; left explicitly open rather than silently dropped from tracking. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly.
- **Two real bugs found via MAME testing and fixed in ENVIRONMENT? : x register clobbered across a COMPAREW call, and a wrong table value predating this - together explained the user's exact reported value precisely, not approximately.** Reported via S" /COUNTED-STRING" ENVIRONMENT? returning -1 \$4E4D instead of the

expected -1 255.

****Root cause**:** ``COMPAREW`` uses ``X`` as its own internal scratch register (``LDX CMPA1``, then ``LDA ,X+`` advancing through the byte-by-byte comparison) and never saves or restores it - any caller relying on ``X`` surviving the call gets it clobbered. ``ENVQUERY`` does exactly that: it needs ``X`` to still point at the current table entry after ``COMPAREW`` returns, both for ``ENVFOUND``'s own ``LDD 4,X`` (reading the value field) and for advancing to the next entry - but never saved it first. Traced precisely why this produces ``$4E4D`` specifically, not just "some garbage": after a successful match against ``/COUNTED-STRING`` (15 characters), ``COMPAREW`` leaves ``X`` sitting exactly at the start of the next table entry's own string, ``"MAX-N"`` - reading ``LDD 4,X`` from there reads two bytes 4 characters into that string: ``'N`` (``$4E``) and ``'M`` (``$4D``, the first character of the following entry, ``"MAX-U"``) - together, exactly ``$4E4D``. Checked the only other ``COMPAREW`` caller in the file (``SUBHASNAME``, from the ``SUBSTITUTE`` rewrite earlier this session) before concluding this was ``ENVQUERY``-specific - confirmed unaffected, since it does a fresh ``LDX`` immediately after its own ``COMPAREW`` call rather than relying on the prior value surviving. Fixed with ``PSHS X`/`PULS X`` bracketing the ``JSR COMPAREW`` call.

****Separate, pre-existing bug found and fixed alongside it**:** ``/COUNTED-STRING``'s table value itself was ``31``, not the ANS-standard ``255`` - a counted string's maximum size is bounded by its 1-byte count field (0-255), not 31, which looks like it was mistakenly copied from an unrelated constraint (``WORD``'s own, separate scan cap from an earlier version of this file, before this session's ``WORD`` redesign). This alone wouldn't have produced the user's exact reported value, but was independently wrong regardless and worth fixing at the same time. Corrected to ``255``.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four ``SERIALPOLL`x`UNITTESTS`` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.

- **Explicit, requested adjustment applied: `BASECODE` and `BASEDICT` origins shifted down a further \$10 (16 bytes) each - the second such adjustment this session, needed for assembly to succeed as the codebase has continued to grow.**

`BASECODE`: `$DFEA` -> `$DFDA` (originally `$E02A` before either shift). `BASEDICT`: `$D7FF` -> `$D7EF` (originally `$D83F`). Applied exactly as requested, same approach as the earlier

\$40 shift - not independently recomputing the resulting gaps/overlaps against `INITCODE` and each other, since that requires real assembled section sizes this environment can't determine; confirm on real assembly. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- **/HOLD added to `ENVTABLE` (previously genuinely absent, per the file's own pre-existing note), answering the design question of which of two candidate constants is correct: `HOLDMINSIZE` (34), not the larger `PADOFFSET` (84).** `HOLDMINSIZE` is specifically the portion of the `CODEHERE` -to- `PAD` gap reserved for the pictured numeric output buffer - the same amount `WORDMAXCHARS` deliberately holds back from `WORD` 's own use at the opposite end of that same shared gap. `PADOFFSET` is the gap's total width, most of which is actually earmarked for `WORD` 's parsing, not `HOLD` 's - reporting it would overstate what `HOLD` can safely use without risking collision with whatever `WORD` is doing in the same space. Confirmed `/HOLD` (the query string, with its leading slash) was genuinely absent from the file entirely before this - only the words `HOLD` / `HOLDS` themselves existed. The reported 255 result therefore couldn't have come from a legitimate table match; most likely a stale binary predating the `ENVQUERY X` -register fix from two turns ago, which would produce exactly this kind of unpredictable result for a query with no real match in the table. Updated both the top-of-file note and the `ENVQUERY` -local comment, which both referenced `/HOLD` as absent - now stale for `/HOLD` specifically; `/PAD` remains genuinely absent and separately tracked, not silently dropped. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly. **MAME-CONFIRMED:** `/HOLD` now correctly returns 34 (`HOLDMINSIZE`) - also confirms the stale-binary explanation for the earlier 255 result was correct, not a lingering bug in the fix.
- **/PAD added to `ENVTABLE` , completing the pair the top-of- file note originally flagged as absent (`/HOLD` and `/PAD`) - both now present.** Answers `PADMINSIZE` (84) directly, per ANS's own `/PAD` meaning (3.3.3.6: "the size of the scratch area whose address is returned by PAD") - PAD's own region, growing upward from PAD itself, conceptually distinct from `HOLDMINSIZE` (which answers for a different region entirely, the downward-growing pictured-numeric buffer in the same `CODEHERE` -to- `PAD` gap, added last turn). `PADOFFSET` happens to equal `PADMINSIZE` numerically in this implementation (`PADOFFSET EQU PADMINSIZE`), but `/PAD` is answered with the conceptually correct constant rather than the coincidentally- equal one, matching the same discipline applied to `/HOLD` . Updated both the top-of-file note and the `ENVQUERY` - local comment, which both still referenced `/PAD` as absent - caught and fixed a duplication error introduced while editing the top-of-file comment (a leftover fragment

line duplicated "open-items checklist") before it reached delivery. The DPHERE/CODEHERE/VARHERE boundary checks remain the one item still genuinely open in this area, kept explicitly tracked rather than dropped. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly. **MAME-CONFIRMED:** /PAD now correctly returns 84 (PADMINSIZE).

- **Explicit, requested adjustment applied: BASECODE and BASEDICT origins shifted down a further \$10 (16 bytes) each - the third such adjustment this session, needed for assembly to succeed as the codebase continues to grow.** BASECODE : \$DFDA -> \$DFCA (\$E02A originally). BASEDICT : \$D7EF -> \$D7DF (\$D83F originally). Applied exactly as requested, same approach as both earlier shifts this session - not independently recomputing the resulting gaps/overlaps against INITCODE and each other, since that requires real assembled section sizes this environment can't determine; confirm on real assembly. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly.
- **FLOORED investigated and added to ENVTABLE , closing the one item explicitly left open two turns ago pending its own investigation.** Unlike WORDLISTS (correctly false by omission, since Search-Order is genuinely unimplemented), FLOORED is a core query - this system definitely has some division behavior, so "false by omission" wasn't the right default. Investigated by tracing DIVCOMMON (shared by / , MOD , /MOD) against this system's own SM/REM and FM/MOD implementations directly, rather than inferring from DIVCOMMON alone: DIVCOMMON restores the remainder's sign from DNSIGN (the dividend's own original sign) after dividing absolute values - confirmed byte-for-byte identical in structure to SM/REM 's own logic. FM/MOD , by contrast, has an explicit flooring-adjustment step (correcting the quotient and remainder when the quotient would be negative with a nonzero remainder) that DIVCOMMON does not have. Conclusion: this system's primary division words use symmetric division, not floored. Added FLOORED to ENVTABLE with value 0 - a real, meaningful "recognized query, value false" result (two items: 0 then TRUE), distinct from the "unrecognized" single-item false fallthrough WORDLISTS correctly uses. Confirmed this doesn't interfere with the table's own terminator check, which reads each entry's address field (offset 0), not its value field - a zero value is safe regardless of position in the table. Updated the ENVQUERY -local comment, which previously described FLOORED as still open. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split- file reassembly. **MAME-CONFIRMED:** FLOORED now correctly returns 0 (recognized, value false). Closes out the full ENVIRONMENT? sweep - /COUNTED-STRING , /HOLD , /PAD , WORDLISTS , and

FLOORED are all confirmed correct on real hardware.

- **MAX-CHAR (3.2.6, maximum value of any character in the implementation's character set) added to ENVTABLE with value 255, matching this system's 8-bit character set (ADDRESS-UNIT-BITS, already in the table, confirms this).** Confirmed MAX-CHAR was genuinely absent from the file entirely before this - same pattern as /HOLD / /PAD / FLOORED before they were added. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly. **MAME-CONFIRMED:** MAX-CHAR now correctly returns 255 .
- **MAX-D / MAX-UD resolved via a genuine dispatcher extension, not just a table entry - the one item in this area explicitly flagged as needing real code work, not just data.** Both are double-cell values per their own ANS data type, but the original ENVFOUND path only ever pushed one cell before the TRUE flag - a table entry alone couldn't answer them correctly. Added a second table (ENVTABLE2, 8-byte entries: addr/len/low/high, vs ENVTABLE 's 6-byte addr/len/value) and a matching second loop (ENV2START / ENV2LOOP / ENV2FOUND), rather than reworking the existing, already-tested single-cell path to handle both entry shapes at once - lower risk. The original table's "not found" path now tries the new table before finally giving up, rather than going straight to ENVNOTFOUND . MAX-D (\$7FFFFFFF, low \$FFFF /high \$7FFF) and MAX-UD (\$FFFFFFFF, low \$FFFF /high \$FFFF) added to the new table - low cell pushed first/deep, high cell second/ top, matching standard double-cell stack order (3.1.4.1). Updated the now-stale comment from two turns ago that described this as still needing dispatcher work. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split- file reassembly. **MAME-CONFIRMED:** MAX-D correctly returns -1 32767 -1 (low \$FFFF, high \$7FFF, then TRUE), and MAX-UD correctly returns -1 -1 -1 (low \$FFFF, high \$FFFF, then TRUE) - both independently confirmed on real hardware, not just structurally. The new double-cell dispatch path (ENVTABLE2 / ENV2LOOP / ENV2FOUND) works correctly for both entries it holds.
- **Explicit, requested adjustment applied: BASECODE and BASEDICT origins shifted down \$40 (64 bytes) each - the fourth such adjustment this session, and larger than the three prior \$10 shifts, matching the larger amount of new code added this turn (the MAX-D / MAX-UD double-cell dispatcher extension - a whole second table plus a matching lookup loop, not just a table row).** BASECODE : \$DFCA -> \$DF8A (\$E02A originally). BASEDICT : \$D7DF -> \$D79F (\$D83F originally). Applied exactly as requested, same approach as every prior shift this session - not independently recomputing the resulting gaps/overlaps against INITCODE and each other, since that requires real assembled section sizes this environment can't determine; confirm on real

assembly. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

- **Explicit, requested adjustment applied: `BASECODE` and `BASEDICT` origins shifted down a further \$20 (32 bytes) each - the fifth such adjustment this session. The prior \$40 shift wasn't enough to resolve the collision, confirmed by trial and error against the real assembler.** `BASECODE : $DF8A -> $DF6A ( $E02A originally)`. `BASEDICT : $D79F -> $D77F ( $D83F originally)`. Applied exactly as requested, same approach as every prior shift this session - not independently recomputing the resulting gaps/overlaps against `INITCODE` and each other, since that requires real assembled section sizes this environment can't determine; the user's own real-assembler trial-and-error remains the actual source of truth for whether this resolves the collision, not a static estimate here. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.
- **`BASECODE` and `BASEDICT` shifted down \$80 (128 bytes) each - the sixth such adjustment this session. Both the prior \$40 and \$20 shifts proved insufficient - deliberately chose a larger jump this time (rather than another small increment) since two consecutive small adjustments had already failed to resolve the collision.** `BASECODE : $DF6A -> $DEEA ( $E02A originally)`. `BASEDICT : $D77F -> $D6FF ( $D83F originally)`. Same approach as every prior shift this session - not independently recomputing the resulting gaps/overlaps against `INITCODE` and each other, since that requires real assembled section sizes this environment can't determine; the user's own real-assembler trial-and-error remains the actual source of truth for whether this resolves the collision. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte-exact split-file reassembly.

`**CONFIRMED RESOLVED**` on real assembly - this shift cleared the collision, with roughly 128 bytes of padding now free. Consistent with the two prior insufficient shifts (\$40 then \$20, summing to \$60/96 bytes) having undershot the true collision size, and this \$80/128-byte shift landing with the full amount to spare - the actual gap needed was evidently somewhere between 96 and 128 bytes. Closes out this round of memory-map adjustments; the next collision, whenever the codebase grows enough to reintroduce one, will need its own fresh trial and error rather than assuming this same margin holds indefinitely.

- **`RETURN-STACK-CELLS` and `STACK-CELLS` added to `ENVTABLE` , completing the full set of**

MAME-testable environmental queries the user identified. Unlike every other entry this session, both use computed expressions tied to this system's own real stack-boundary constants (`RSTACK` , `DSTACK` , `CODETOP`) rather than hardcoded numbers - deliberate, given the memory map has shifted repeatedly this session (`BASECODE` / `BASEDICT` alone moved six times) and a hardcoded number would silently go stale on the next shift with nothing to catch it. `RETURN-STACK-CELLS` = $(RSTACK - DSTACK) / 2 = 384$ currently (`RSTACK`'s own existing comment already confirms the occupied range is `$BD00 - RSTACK` , 768 bytes, and `DSTACK+1` is `$BD00` exactly per both regions' own comments confirming they're contiguous - the algebra simplifies to `RSTACK-DSTACK` directly). `STACK-CELLS` = $(DSTACK-CODETOP+1) / 2 = 512$ currently (`CODETOP` 's own comment: "code space ceiling (data stack begins here)"). Caught and simplified a redundant `-1+1` in the first draft of the `RETURN-STACK-CELLS` expression before delivery - algebraically identical but needlessly complex for an expression this environment can't test against a real assembler. Explicitly verified both expressions evaluate to exact integers (no truncation risk regardless of the assembler's own integer-division behavior) in addition to the standard structural checks, since FDB arithmetic isn't something the duplicate- symbol/dictionary-chain verification evaluates. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL X UNITTESTS` combinations; byte- exact split-file reassembly. **MAME-CONFIRMED:** both correctly return `384` and `512` respectively - confirms the computed `FDB` expressions assembled and evaluated correctly on the real toolchain, not just in this environment's own arithmetic check. Closes out the full `ENVIRONMENT?` sweep - every entry in both `ENVTABLE` and `ENVTABLE2` is now independently confirmed correct on real hardware.

- **Documentation updated: title page duration entry, full Memory Map table rebuild, and a new Transient Region section - grounded in `forth6809_1st.txt` , a real assembler listing of the current source, uploaded this turn.** Confirmed the listing genuinely matches the current source before treating any of its numbers as authoritative - `BASECODE=$DEEA` , `UNITTESTS=1` , and this session's additions (`MAX-CHAR` , `MAX-D` , `RETURN-STACK-CELLS`) are all present in it, not a stale snapshot.

```
**Title page**: added "Manual MAME testing and bug-fixing
duration (guesstimate): 80 hours", matching the formatting of
the existing duration entries.
```

```
**Memory Map table, fully rebuilt** with real, assembler-
measured values rather than the old estimates: `BASECODE` =
8304 bytes (`BASECODESIZE`, was a stale 8110 estimate);
`BASEDICT` = 2027 bytes (`BASEDICTSIZE`). Discovered and
documented two previously-unverifiable facts directly from the
```

listing: `BASEDICTEND` equals `BASECODE` exactly (`\$DEEA`) - zero gap, zero overlap - and `INITEND` equals `VECTORS` exactly (`\$FFF0`) - also zero gap. The one real gap (`BASECODEEND` to `INITCODE`, 79 bytes) is explicitly `\$FF`-filled by the assembler's own `FILL` directive, not an unaccounted collision. `APPCODE` row relabeled `APPCODE + Transient`, with a pointer to the new section, since the Transient region is a floating subregion within it rather than a fixed range of its own - chosen over a nested/indented row given the table's existing one-row-per-fixed-range structure. `APPDICT`'s upper bound corrected to `APPCODE-1` (20480 bytes; the prior `\$6EA4`/20133 figure was stale, left over from an earlier `APPCODE` position before it moved to its current, "back to its original address" value). Retired `SIBUF`/`WORDBUF` rows removed (confirmed genuinely retired - their `EQU`s are commented out in the listing) and replaced with a single `(unclaimed)` row for the 65-byte range they used to occupy. "ROM Size Required" narrative rewritten with the corrected totals: 10418 bytes real ROM content (16+71+8304+2027), 5710 bytes free (~35%) - both now real, assembler-measured figures rather than the prior manual-count estimate the old text explicitly flagged as uncertain.

****New "Transient Region" section****, inserted at the same heading level as "Dictionary Entry Layout" (matching its style: `Normal` paragraph, bold run, no explicit size), immediately before it. Describes the region's function (floating, `CODEHERE`-relative, shared by `WORD`'s parsing buffer and the `HOLD` pictured-numeric-output buffer), size (`PADOFFSET`/`PADMINSIZE` = 84 bytes, `PAD = CODEHERE + 84`), and location, then a new sub-table (Subregion / Size / Address relative to `CODEHERE` / Description) with exact byte-level addressing traced from `WORD`'s and `ENVQUERY`'s own source comments rather than reconstructed from memory: WORD parsing buffer (up to 47 bytes, `CODEHERE` or `CODEHERE+3` for `S"/\."/ABORT"`, up to `CODEHERE+49` worst case); HOLD buffer (34 bytes, `CODEHERE+50` to `CODEHERE+83`, i.e. `PAD-34` to `PAD-1`); PAD itself (at least 84 bytes from `CODEHERE+84` onward). Confirmed by direct arithmetic that the worst-case WORD buffer boundary (`CODEHERE+50`) exactly touches the HOLD buffer's start with no gap and no overlap, matching `WORDMAXCHARS`'s own defining formula (`PADOFFSET-HOLDMINSIZE-1-3`).

Verified visually via rendered PDF pages, not just text extraction - both the rebuilt Memory Map table and the new Transient Region table/section render cleanly across their page breaks. Confirmed the field-based Table of Contents correctly refreshed after the new content shifted page numbers (`3`.

Glossary` moved from page 9 to 10, and every entry after it shifted accordingly). Confirmed remaining `SIBUF`/`WORDBUF` mentions elsewhere in the document are expected and correct - the new section's own historical explanation, and the verbatim Assembler Source appendix (Section 8), which legitimately still shows this history in its own preserved inline comments - not leftover staleness. No source-code changes were made or needed this turn; this was a documentation-only update.

- **Documentation resync: Glossary and Assembler Source appendix both brought up to date with the entire MAME-testing phase of this session, after verification found both had stopped being updated right after the double-number-input work.**

Confirmed the exact boundary before resyncing: `NUMBERQ / TRYNUM`'s appendix section was current (last thing explicitly synced), but `SIGN`, `S>D`, `LPAREN`, `ENVQUERY`, and the full `ENVTABLE / ENVTABLE2` expansion were all missing - the appendix still showed pre-fix code (`SIGN / S>D` missing their `TSTA` checks, `LPAREN` missing its stray-address fix, `ENVQUERY` missing its `PSHS X / PULS X` fix, `ENVTABLE` showing only the original 4 entries with `/COUNTED-STRING` still at the wrong value 31). Replaced five appendix subsections wholesale with the current split-file source (8.1 Memory Map, 8.16 Arithmetic, 8.21 Numeric Output, 8.24 Comment Words, 8.25 Environmental Query) - same approach used successfully for the double-number sync earlier, chosen over surgical patching of individual routines within a single large paragraph.

Glossary: found the `ENVIRONMENT?` entry actively wrong, not just stale - it explicitly stated "Table is incomplete: missing /HOLD, /PAD, MAX-D, MAX-UD, WORDLISTS, FLOORED," all of which had since been added and MAME-confirmed. Rewrote it to list every supported query and value, and widened the stack-effect notation to `( c-addr u -- false | i\*x true )` to correctly reflect that some queries (`MAX-D`/`MAX-UD`) return two cells, not one. Checked `SIGN`'s, `S>D`'s, and `(`'s own glossary descriptions before assuming they needed similar rewrites - confirmed they didn't: each already described the *intended* behavior the bug fixes now correctly deliver, so the bugs never made the documentation wrong, only the implementation.

Verified: paragraph and table counts unchanged before/after (1444 paragraphs, 3 tables); explicit presence checks for every fix/addition across the rendered PDF; visually confirmed two representative pages (the updated `ENVIRONMENT?` glossary entry, and the appendix's `SIGN` routine with its full fix comment) - both render cleanly with no formatting corruption from the large content replacement.

- TSTDROP added, following TSTDUP 's pattern exactly - second test in what's intended to become a test for every stack manipulation word.** Verifies both the stack's contents after DROP (the sentinel guard pushed beneath the dropped value is confirmed undisturbed and is now the new top) and the data stack pointer's movement (exactly one cell, 2 bytes, freed - DROP 's own net effect: LEAU 2,U). Traced by hand before delivery, not just structurally verified: with TSTUB4 captured at U_start-4 (both TSTGUARD and TSTVAL1 pushed) and TSTUAF captured at U_start-2 (after DROP 's LEAU 2,U), TSTUB4-TSTUAF computes to exactly -2 , matching the check - the sign is deliberately opposite TSTDUP 's own +2 check, since DROP frees a cell while DUP adds one. Used distinct internal labels (DFFAIL / DPDONE) rather than reusing TSTDUP 's (TDFAIL / TDDONE), to avoid a collision now and to establish a naming pattern (short prefix per test) that won't collide as more tests are added. Wired into TSTSTACK alongside the existing JSR TSTDUP . UNITTESTS reactivated (1 -> 0), confirmed appropriate since the user activated it independently and confirmed TSTDUP runs successfully on real MAME hardware - whatever earlier concern prompted excluding it was resolved by other, unrelated fixes made since. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations (checked with UNITTESTS both on and off, not just the file's new default); byte-exact split-file reassembly. Not yet confirmed via MAME.
- 14 more stack-manipulation tests added (TSTSWAP , TSTOVER , TSTROT , TSTQDUPNZ , TSTQDUPZ , TSTDEPTH , TSTDDUP , TSTDDROP , TSTD SWAP , TSTD OVER , TSTNIP , TSTTUCK , TSTPICK , TSTROLL , TSTDROT) - completing every word in the Data Stack Manipulation glossary section.** Each follows TSTDUP / TSTDROP 's established pattern exactly (guard pushed beneath, U snapshotted before/after, contents verified via PULU , pointer movement verified independently, unconditional restore at the end). ?DUP specifically got two tests (TSTQDUPNZ , TSTQDUPZ), per the explicit request - nonzero and zero are genuinely different code paths (QDUP branches on the popped value), not just different inputs to the same path. Six new test-value constants added (TSTVAL2 - TSTVAL6 , plus TSTGUARD / TSTVAL1 already existing) - all distinct from each other and from 0 / 1 / -1 , so a test that only appears to pass due to a trivial or coincidentally-matching value would be caught. One new scratch cell (TSTSCR) added for TSTDEPTH , which independently computes its own expected value via the same (SP0-U)/2 formula DEPTH itself uses, rather than assuming a fixed starting depth - robust regardless of whatever's already on the stack when it runs. TSTPICK / TSTROLL both use u=2 as a concrete representative case (0 PICK is DUP , 1 PICK is OVER - 2 is the first case distinct from both, and using the same u for both makes the two tests directly comparable).

Every test's expected stack layout was traced by hand,

instruction by instruction, against each word's real implementation before being written - not assumed from the glossary's stack-effect notation alone. The five trickiest (`PICK`, `ROLL`, `2SWAP`, `2OVER`, `2ROT`) were also cross-checked with an independent Python simulation. That simulation caught a real bug in itself, not in the assembly design: an incorrectly-ordered `2OVER` model that inserted both pushed items in one step rather than modeling them as two sequential pushes - once corrected to push explicitly in order, it confirmed the original hand-trace was right all along. Worth recording as a reminder that a "second check" needs its own verification too, not blind trust.

Careful, collision-free label naming: each test uses its own short prefix for internal `FAIL`/`DONE` labels (`SWFAIL`/`SWDONE`, `OVFAIL`/`OVDONE`, etc.), with the double-cell variants using their single-cell counterpart's prefix plus a `2` (e.g. `SW2FAIL` for `2SWAP`, distinct from `SWAP`'s own `SWFAIL`) - confirmed zero collisions via the same structural verification used throughout this session, which would have caught any duplicate automatically. All 16 tests (including the two already-confirmed ones) wired into `TSTSTACK` in glossary order. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly. Not yet confirmed via MAME.

- **INITCODE shifted down 3 bytes (\$FFA9 -> \$FFA6), and the TSTRUNNER call site fixed to always emit exactly 3 bytes regardless of UNITTESTS - both per explicit request, closing a real risk: COLDSTRT 's own size previously depended on UNITTESTS (the JSR TSTRUNNER call existed only when UNITTESTS=0 , emitting 0 bytes when UNITTESTS=1), while INITCODE 's position was fixed regardless - meaning toggling UNITTESTS could silently overflow the code into VECTORS without any assembler error to catch it.** Fixed by adding an `ELSE` branch: three `NOP` s (byte-for-byte the same size as the `JSR` they replace) when `UNITTESTS=1` , so the block now contributes exactly 3 bytes to `COLDSTRT` either way - its total size no longer depends on `UNITTESTS` at all. `INITCODE` shifted down by the same 3 bytes to compensate. Traced the resulting arithmetic by hand rather than assuming it worked out: the previously-confirmed 71-byte real content figure was measured under the *old* structure (0 bytes for this block, since that measurement was taken with `UNITTESTS=1`) - under the new, always-3-bytes structure, real content is reasoned to be 74 bytes (71+3), not yet re-measured by a real assembler run. The 3-byte shift in `INITCODE` 's own start and the 3-byte growth in content offset exactly, so the reasoned end (`$FFEF`) is unchanged - still one byte below

VECTORS , matching the previously-confirmed real value, if the reasoning holds; flagged as reasoned-not-measured rather than presented as confirmed, and a stale comment referencing the old \$FFA9 position corrected in the same edit. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; explicitly confirmed the ELSE branch selects mutually-exclusively between the real JSR and the NOP placeholder depending on UNITTESTS , not just that the file parses; byte-exact split-file reassembly. Not yet confirmed via MAME - given this directly concerns VECTORS safety, this one specifically warrants confirming on real assembly before being considered settled, not just structural verification.

- **TSTSTACK given a header (CR , "Stack", CR), and a new subroutine TSTSARITH added - single-cell arithmetic tests covering every word in glossary section 3.4, with its own matching header ("SArithmetic"). 28 new tests total, wired into TSTRUNNER alongside TSTSTACK . Two new negative test constants added (TSTNEG1 =-12345, TSTNEG2 =-321) - TSTVAL1 - TSTVAL6 are all positive, which wouldn't exercise the sign- dependent branches several of these words genuinely have (ABS , NEGATE , MIN / MAX , signed division, 2/ 's sign- preserving shift).**

The two behaviors the request flagged as uncertain were both resolved by reading this system's own glossary and source directly, not assumed: overflow is explicitly ANS-defined as truncation, not an error (``'s own glossary entry: "Signed multiply, truncated to one cell") - tested with a case chosen specifically to overflow 16 bits (1000 * 1000), verifying the wrapped result, not an exception. Division by zero throws -10 - confirmed both in the glossary ("Throws -10 if n2 is zero") and by reading DIVCOMMON`/`STARSLASHCOMMON` directly, both of which explicitly THROW -10 on a zero divisor before any division is attempted.

Divide-by-zero tests (`/`, `MOD`, `/MOD`, `\*/`, `\*/MOD` - one each) use a different pattern than the rest, built around `CATCH`: push the operands and the target word's `xt`, call `CATCH`, verify the popped result is -10, and verify `CATCH`'s own documented depth-restoration contract held (net 0 change in `U` across the JSR CATCH, matching the ANS guarantee - `CATCH` only promises stack *depth* is restored, not the `i\*x` values, confirmed precisely earlier this session). Deliberately does not try to inspect or discard a specific count of leftover `i\*x` items - the test's own final, unconditional LDU TSTU0 (already established practice for every test in this framework) discards whatever's left regardless of the count, matching the ANS standard's own guidance to DROP rather than interpret them.

All 28 expected values computed programmatically (Python) before writing any assembly, rather than by hand, given the volume - reduces transcription-error risk across that many cases. Operand pushes and expected-value comparisons both use symbolic constant references (`#TSTVAL1` etc) wherever they correspond to an existing named constant, matching this framework's established style, rather than embedding literal hex duplicates - including, correctly, for a few cases where the *expected result* itself equals an operand's own value (e.g. `ABS` of an already-positive number returns itself unchanged, so the comparison target is `#TSTVAL1` too, not a separately-typed literal that happens to match).

Careful, collision-free label naming continued: each of the 28 new tests uses its own short prefix (`PL`, `MN`, `S1`/`S2`, `SL`/`SN`/`SZ`, `MD`/`MZ`, `SM`/`MX`, `NG`, `A1`/`A2`, `N1`/`N2`, `X1`/`X2`, `1P`/`1M`/`2P`/`2S`, `D1`/`D2`, `TS`/`TZ`, `TM`/`TX`) - confirmed zero collisions against both each other and all 17 existing stack-manipulation-test labels via the same structural verification used throughout this session.

Caught and fixed two real mistakes in my own generation process before delivery, not after: (1) a first draft of the `/` quotient-only test incorrectly used the two-result template (`/` only pushes the quotient, not a remainder - confirmed by re-reading `SLASH`'s actual code); (2) an early, mechanical symbolic-reference conversion pass left the expected-value formatting still using a helper that assumed a literal integer, which would have crashed generation the moment an expected value legitimately needed to be a symbolic reference too (the `ABS`/`MAX` cases above) - caught by actually running the generator and inspecting output, not just writing the code and assuming it worked.

Also worth recording: my own post-generation verification script produced two rounds of false "missing" alarms, both from bugs in the *checking* script itself (a text-boundary search that matched an early, incidental occurrence of a label name instead of its actual definition; and a check that only recognized `EQU`-style definitions, not label-style ones) - not from the generated source, which was correct both times. Resolved by direct inspection of the actual file content rather than trusting the automated check's first result, and by fixing the check itself before relying on it again - the same "verify your verifier" lesson already recorded elsewhere in this file.

Verified: zero duplicate symbols, dictionary chain still 224

entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; every symbolic constant/label reference in the new block resolves to a real definition somewhere in the file, checked explicitly rather than inferred from the duplicate-symbol check alone (which doesn't catch undefined references, only duplicate definitions - a gap in this session's usual verification, closed for this delivery). Not yet confirmed via MAME - given the volume (28 new tests in one delivery, several using a genuinely new pattern involving `CATCH`), this specifically deserves a real test before being considered settled, more than most single-word additions this session.

- RESOLVED - real assembly-blocking bug, not caught by this session's own structural verification: four internal label pairs in the arithmetic tests**
 (1PFAIL / 1PDONE , 1MFAIL / 1MDONE , 2PFAIL / 2PDONE , 2SFAIL / 2SDONE , for 1+ , 1- , 2+ , 2*) started with a digit, which the real assembler rejects (labels must start with a letter, underscore, or period) - the prefix naming scheme mechanically derived these from the Forth word names themselves, which happen to start with digits, without checking that the resulting label text was itself valid. Confirmed by the user's real assembler run, not caught here first. Renamed to OPFAIL / OPDONE (1+), OMFAIL / OMDONE (1-), TPFALL / TPDONE (2+), TWFAIL / TWDONE (2*) - chosen to avoid colliding with any of the other 40+ short prefixes already in use across both test files (TS was already taken by STARSLASH 's own tests, which is why 2* couldn't reuse it despite being an obvious first choice). Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four SERIALPOLL X UNITTESTS combinations; byte-exact split-file reassembly; an explicit sweep for any other label starting with a digit anywhere in the entire file, not just the four reported - found none. Worth noting as a real gap in this session's own structural verification: the duplicate-symbol and dictionary- chain checks used throughout never actually checked label *validity* (only uniqueness), so a mechanically-generated invalid label like this could pass every check this session ran and still fail on real assembly - closed for this fix by adding the explicit digit-start sweep, worth keeping as a standing check for any future generated labels.
- RESOLVED - two distinct real bugs in the arithmetic tests themselves, found by the user's actual MAME run (14 of 28 failed), not by this session's own verification, which had checked structural validity but never independently re-derived the expected values or checked the depth-check sign convention against each word's real arity.**

****Bug 1, the dominant one (13 tests): the depth-check sign was backwards for every 2-in/1-out and 3-in/1-out test.**** The

generator's `binop\_test`/`triop\_test` templates compared `TSTUB4-TSTUAF` against a **positive** `2` (or `4`), but a net-consuming operation (more cells popped internally than pushed back) **increases** `U`, which makes `TSTUB4-TSTUAF` **negative** - exactly the same sign convention already established and correctly used by `TSTDROP`/`TSTNIP`/`TSTDDROP` earlier in this same file. The template's own docstring even correctly said "-2 (one cell freed)" while the template body used `#2` - the comment and the code disagreed, and nothing checked that they matched. Affected: `TSTPLUS`, `TSTMINUS`, `TSTSTAR1`, `TSTSTAR2`, `TSTSLASH1`, `TSTSLASH2`, `TSTMODW`, `TSTMIN1`, `TSTMIN2`, `TSTMAX1`, `TSTMAX2` (all `2`->`-2`), `TSTSTSL` (`4`->`-4`). Every 1-in/1-out test (`NEGATE`, `ABS`, `1+`, `1-`, `2+`, `2\*`, `2/`) and every `CATCH`-based divide-by-zero test correctly used `0` and were unaffected - confirmed by an independent, arity-derived recomputation of what every one of the 28 tests' depth check **should** be, not just the ones already reported failing.

****Bug 2, a separate issue affecting `TSTSLMOD` and `TSTSTSM`:** the two-result comparisons were in the wrong order.** Both `SLASHMOD` and `STARSLASHMOD` push the remainder first, then the quotient last (confirmed by re-reading both routines directly: `LDD DIVREM / PSBU D` then `LDD DIVNUM`-or-`PRODLO` / `PSBU D`) - meaning the quotient is on top, popped first. The generated tests had the two expected values swapped (remainder compared first, quotient second). `TSTSTSM` had both bugs simultaneously - its depth check **and** its value order were both wrong, compounding into the same single `FAIL` report from MAME. `TSTSLMOD`'s depth check (`0`, correct for a 2-in/2-out operation) had been right all along; only the value order was wrong.

Every one of the 28 tests' depth-check sign and numeric expected values were independently re-verified against a fresh Python recomputation after the fix, not just the ones the user reported - to catch anything else wrong but not yet exercised by the specific values chosen. None found. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; no digit-leading labels remain (re-checked after this edit too, not just the prior one).

Root-cause note for future generated test batches: this session's structural verification (duplicate-symbol and dictionary-chain checks) cannot catch a **logically** wrong but **structurally** valid comparison value or sign - both of these

bugs assembled and ran without any assembler or structural-check complaint, and would only ever have been caught by either real execution or an independent recomputation of expected values, which is what closed this out. Worth treating "does the generator's own template body match its own docstring's stated formula" as an explicit check for any future generated test batch, not just for this one after the fact.

- **TSTDARITH added - mixed & double-precision arithmetic tests covering every word in glossary section 3.5 (14 words, 15 tests since DABS gets two - positive and negative cases - plus 3 divide-by-zero tests, 18 total), wired into TSTRUNNER alongside TSTSTACK / TSTSARITH . Header ("DArithmetic") follows the same CR /name/ CR pattern as the other two groups.** Every implementation traced by hand before any test was written, specifically to avoid repeating the previous batch's mistakes: confirmed double-cell values are pushed low- cell-first, high-cell-last (on top) consistently across every word here; confirmed which operand is d1 vs d2 (or dividend vs divisor) from the actual pop order, not assumed from the glossary stack notation; confirmed FM/MOD (floored) and SM/REM (symmetric) genuinely diverge on the same negative- dividend test case ($-70000 / 9320$: floored gives quotient -8 , remainder 4560 ; symmetric gives -7 , -4760 - both satisfy $\text{quot} * \text{divisor} + \text{rem} = \text{dividend}$, confirming the difference is real, not an arithmetic error) - Python's native `//` / `%` used directly for the floored case, matching ANS's definition exactly, rather than hand-deriving it.

Depth-check values are now computed automatically from an explicit `(cells_in, cells_out)` parameter passed to the generator, never a separately hand-typed signed literal - a direct structural fix for the root cause of the previous batch's dominant bug (a docstring correctly saying "-2" while the actual code said "2", with nothing checking they matched). Independently re-derived what every one of the 15 core tests' depth check **should** be from each word's real arity before trusting the generator's own output, rather than assuming the fix worked - all 15 matched on the first attempt.

Label collisions checked programmatically before insertion, not assumed safe by inspection - caught and fixed 6 real collisions this way (``DP``, ``MN``, ``MX``, ``MS``, ``DS``, ``DT`` all clashed with either an existing test's prefix or, in ``MSTAR``'s case, an **internal** label already used inside the real word's own implementation - ``MSDONE`` - which a manual prefix-list cross-check would likely have missed entirely, since it's not another test's label at all). Replaced with ``PD``, ``NM``, ``XM``,

`MC`, `BD`, `NS` respectively, each verified against every label in the entire file, not just other tests' prefixes.

All 15 numeric expected values independently re-verified against fresh Python computation after generation, with proper symbolic-constant resolution this time (the verification script itself needed a fix mid-process - it initially flagged several correct results as mismatches because it compared raw symbolic references like `TSTD3HI` against literal hex without resolving them to their real `EQU` value first; fixed by resolving symbols before comparing, the same "verify your verifier" pattern recorded from the section 3.4 work). All 18 tests confirmed correctly wired into `TSTDARITH`'s call list, and every symbolic reference confirmed to resolve to a real definition somewhere in the file.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly; explicit sweep confirms no digit-leading labels anywhere in the file.
MAME-CONFIRMED: the user reports all `TSTDARITH` tests now pass, including `TSTMSTAR` after its own fix (logged separately above).

- **RESOLVED - a real bug in `MSTAR` itself, not in `TSTMSTAR` or its push/pop order, found by the user's actual MAME run (reported as "the 2 returned values are swapped" - the real cause turned out to be a wrong computation, not a swap, though the symptom is an understandable way to describe it from the outside).** Same bug class already found and fixed multiple times this session (`PULU` doesn't set condition codes on genuine 6809), but in a new specific shape not seen before: `CLR MSIGN` runs immediately after the `PULU D` that loads `n1`, and `CLR` unconditionally sets `N=0` (since it clears its destination to 0) - so the following `BPL MSN1POS` was testing `CLR`'s own result, not `n1`'s actual sign, and always branched. `n1` was never negated even when genuinely negative - its raw bit pattern got used as an unsigned magnitude instead, and `MSIGN` never got toggled for it, so the final `MNEG32` never ran either (since the code believed both operands were positive). Traced and confirmed by hand: for this test's actual inputs (`-12345 * 9320`), the buggy path computes `53191 * 9320` (treating `-12345`'s raw 16-bit pattern as the unsigned value 53191) instead of first negating to `12345` - a real, large, specific wrong answer, not a simple swap.

Confirmed this bug was genuinely isolated to `MSTAR`, not a systemic pattern - both `FMSLASHMOD` and `SMSLASHREM` have the same general shape (`CLR` calls before a sign-check branch) but

already do this correctly, with an explicit `TST PRODHI` between their own `CLR`s and the branch, re-establishing the correct flags - `MSTAR` was simply missing the equivalent re-test that its siblings already had. A programmatic sweep of the entire file for the same unguarded shape (`PULU D` followed within a few lines by a `CLR` then a signed branch, with no intervening flag-setting instruction) found no other instances.

Fixed with `TSTA` (testing D's high byte, the same fix pattern already used for `TRYNUM`/`STOD`/`SIGN` earlier this session) inserted between `CLR MSIGN` and `BPL MSN1POS`. Verified the fix by hand-tracing the corrected path for this test's exact inputs: `n1` now correctly detected as negative, correctly negated to `12345` before the unsigned multiply, `MSIGN` correctly toggled, final `MNEG32` correctly applied - produces exactly `TSTMSTAR`'s original expected values (`\$F924`, `\$64D8`) with zero changes needed to the test itself. Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all four `SERIALPOLL`x`UNITTESTS` combinations; byte-exact split-file reassembly. ****MAME-CONFIRMED****: the user reports all `TSTDARITH` tests now pass, including `TSTMSTAR` - the fix holds on real hardware, not just by hand-trace.

- **RESOLVED - TSTSELECTOR transferred in from the user's own tested approach, addressing a real problem this session hadn't caught: assembling all three test groups (TSTSTACK , TSTSARITH , TSTDARITH) together exhausts the available unused ROM space.** New flag, `TSTSELECTOR EQU 2`, gates each group's own call-list *and* its own test-body definitions separately (two `IFEQ TSTSELECTOR-N / ENDC` pairs per group, `N = 0/1/2` for `TSTSTACK / TSTSARITH / TSTDARITH` respectively) - so only the currently-selected group's test code actually compiles into the ROM image, not all three at once. Each group's header message (`CR /name/ CR`) stays unconditional, so it prints regardless of which group is selected; only the actual test-invoking calls and their bodies are gated.

Received as an uploaded file rather than a description, so verified by diffing the user's file against this session's version rather than assuming the transfer was needed at all - confirmed the diff was genuinely just this mechanism (plus this session's own `MSTAR` bug fix, absent from the user's copy since it predates that fix) before applying anything. Applied by locating each of the six exact insertion points in this session's own file and inserting the identical text, then diffed the result against the user's file again to confirm an exact match (aside from trailing-whitespace-only differences on

a few blank lines, and the `MSTAR` fix, correctly present only in this session's copy).

Verified beyond the standard structural check: extended the duplicate-symbol/dictionary-chain verification to cover `TSTSELECTOR`'s three values crossed with both `SERIALPOLL` settings (8 relevant combinations, since `TSTSELECTOR` only matters when `UNITTESTS=0`) - zero duplicate symbols, dictionary chain intact (224 entries) in all eight. Also explicitly confirmed the mechanism achieves its actual goal, not just that it assembles cleanly: simulated each `TSTSELECTOR` value and confirmed only that one group's test bodies are present in the output (checked via a representative test name per group), with the other two groups' bodies genuinely absent, not just unreachable - and confirmed total compiled size per selector value is roughly a third of what all three groups combined would be. Byte-exact split-file reassembly confirmed.

****MAME-CONFIRMED****: the user reports the whole `TSTDARITH` group (`TSTSELECTOR=2`) now assembles and runs successfully, validating the mechanism end to end, not just structurally.

- **RESOLVED - UNITTESTS and TSTSELECTOR converted from plain EQU to an IFNDEF / SET / ENDC fallback-default pattern, per explicit request, supporting future override via lwasm's -D command-line option without editing the source.** -D pre-defines a symbol before the source is processed, so IFNDEF correctly detects and skips the fallback SET when a value was supplied that way, leaving the -D -provided value in place; when no -D was given, the symbol is genuinely undefined at that point, IFNDEF fires, and the fallback value applies. TSTSELECTOR's fallback preserves its current working value (2) - the request's own example showed 0, but that was illustrating the pattern, not asking to reset the actual selector.

****`UNITTESTS`'s flag meaning deliberately reversed at the same time, per explicit request****: both `IFEQ UNITTESTS` occurrences (the main test-framework body, and the `COLDSTRT` call site) changed to `IFNE UNITTESTS` - inverting the semantics from "0=included" to "0=excluded, nonzero=included". This makes 0 (or no -D at all) the safe, production default - test-framework code and its `TSTRUNNER` call site both genuinely absent from a default build - with testing now opt-in via `-DUNITTESTS=1` rather than opt-out via editing a plain `EQU`. Both `IFEQ`→`IFNE` changes applied together deliberately, not independently - leaving one as `IFEQ` while the other became `IFNE` would have left the test-framework body and its own call site testing opposite conditions, a real risk of building correctly but crashing at the call site (or vice versa) that

was checked for explicitly, not just assumed avoided by intent. Updated stale comments referencing the old "0=included,1=excluded" convention in both places, including the historical note on the ``COLDSTRT`` call site's own ``NOP``-padding fix from earlier in the session, which had specifically referenced ``UNITTESTS=1`` as the excluded case - no longer accurate under the reversed meaning.

This session's structural verification simulator extended to handle the three new constructs (``IFNDEF``, ``SET``, ``IFNE``), none previously present in this file - ``IFNDEF`` modeled via an explicit set of "defined" flags (mirroring what ``-D`` would predefine), ``SET`` treated identically to ``EQU`` for verification purposes, ``IFNE`` as the logical inverse of the existing ``IFEQ`` handling. Verified across scenarios a plain duplicate-symbol check wouldn't distinguish: (1) no ``-D`` at all - confirmed the entire test framework, including ``TSTRUNNER`` and every individual test body, is genuinely absent, not just unreachable, and the call site correctly falls through to its ``NOP`` placeholder; (2) ``-DUNITTESTS=1`` with each ``TSTSELECTOR`` value explicitly passed - confirmed correct group selection still works under the reversed ``UNITTESTS`` convention; (3) ``-DUNITTESTS=1`` alone, ``TSTSELECTOR`` *not* separately overridden - confirmed it correctly falls back to its own default (2) rather than picking up ``UNITTESTS``' value or failing to resolve. All scenarios: zero duplicate symbols, dictionary chain intact (224 entries). Byte-exact split-file reassembly confirmed. Not yet confirmed via real `lwasm` - simulated ``-D`` behavior here, not run against the actual toolchain's own ``-D`` implementation, which is worth checking given this is new assembler-feature territory for this file, not just new source structure.

- **RESOLVED - unit test framework block comment fixed, per explicit request:**
"UNITTESTS=1 removes it entirely" was stale after the flag-meaning reversal, and would have told a reader the exact opposite of the current behavior. Rewrote to describe the current (0/undefined=excluded, nonzero=included) convention, and added the requested example `lwasm` command line (`--define=UNITTESTS --define=TSTSELECTOR=2`) at the end of the block comment. Also proactively fixed three further references to the old convention in `INITCODE``'s own historical comment (describing an earlier fix made before the reversal existed) - left unfixed, these would have recreated the same confusion just flagged, one layer further into the file.
- **TSTLOGIC added - logic, shift, and address-arithmetic tests (glossary section 3.6, 12 words, 12 tests), wired into TSTRUNNER and gated as TSTSELECTOR's fourth group (-3).** Several of these words modify the top of stack in place (`LDD / op/ STD`, never

PULU / PSHU at all - INVERT , CELLS , CELL+ , CHARS , CHAR+ , ALIGNED) - tests verify what's observable via the stack regardless of implementation mechanism, reusing the same generic push/pop test template as every previous group without needing special-casing. Three words (CHARS , ALIGN , ALIGNED) are documented no-ops on this system; CHARS / ALIGNED got ordinary single-value identity tests, while ALIGN - which takes no stack arguments at all - got a dedicated test pushing two decoy values to confirm the *whole* stack is undisturbed, not just a single value's persistence (the same template handles this too, called with cells_in=0, cells_out=0 and no real "operand" semantics). RSHIFT tested against a negative input specifically, since it's documented logical (zero-fill) rather than arithmetic (sign-preserving) - the same reasoning already applied to 2/ and FM/MOD -vs- SM/REM in earlier groups, confirmed to give a genuinely different result than an arithmetic shift would.

Label collisions checked programmatically before insertion - caught 3 real ones this way (`LS` , `RS` , `HS` all collided with internal labels already inside `LSHIFT`'s, `RSHIFT`'s, and `HOLDS`'s own implementations - `LSDONE`/`RSDONE`/`HSDONE` - the same class of issue as `MSTAR`'s own `MSDONE` collision found earlier). Replaced with `L2` , `R2` , `C3` , each re-verified against every label in the file.

****A real mistake caught and fixed before delivery, not after**:**
the initial insertion left `TSTSELECTOR`'s second `IFEQ`/`ENDC` pair (wrapping `TSTLOGIC`'s own test-body definitions) with no closing `ENDC` at all - the wrapper and generated test bodies were concatenated without one, unlike the `TSTDARITH` transfer earlier, which needed and got an explicit closing `ENDC` inserted by hand. The open block fell through and silently consumed the pre-existing outer `ENDC` that closes the entire `UNITTESTS` block, leaving everything unbalanced by one. This didn't fail the routine duplicate-symbol check (which doesn't verify `IFEQ`/`ENDC` balance at all) - it was caught by extending verification to actually simulate the production-default build (no `-D` at all) and finding the dictionary chain walk returned **zero** entries instead of 224, a stark, unmissable signal once checked, but one that a narrower check would have missed entirely. Root cause confirmed precisely (an explicit open/close trace of every `IFEQ`/`IFNE`/`IFNDEF`/`ENDC` in the file, showing depth ending at 1 instead of 0) before fixing, not guessed at. Fixed by adding the missing `ENDC` ; `IFEQ`/`IFNE`/`IFNDEF` vs `ENDC` counts now balanced (15/15) file-wide.

Verified after the fix: every one of the 12 depth-check values independently re-derived from real arity and confirmed correct;

every numeric expected value independently re-verified against fresh computation with proper symbolic-constant resolution; every symbolic reference confirmed to resolve to a real definition; all 12 tests confirmed wired into `TSTLOGIC`. Full scenario matrix re-run after the `ENDC` fix specifically (not just re-trusted): production default (no `-D`) now correctly shows the full 224-entry chain again, `-DUNITTESTS=1` crossed with all four `TSTSELECTOR` values and both `SERIALPOLL` settings (10 scenarios total) all pass; explicitly confirmed `TSTSELECTOR=3` includes only `TSTLOGIC`'s own bodies, with the other three groups' bodies genuinely absent. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTLOGIC` tests pass, including `ALIGN`'s dedicated no-op test and `RSHIFT`'s negative-input logical-shift case - the `ENDC` fix holds on real hardware too, not just in simulation.

- **TSTCOMPARE added - comparison tests (glossary section 3.7, 14 words, 15 tests since WITHIN gets two cases), wired into TSTRUNNER and gated as TSTSELECTOR's fifth group (-4).** Signed-vs-unsigned pairs ($</U<$, $>/U>$) each tested with an operand pair that would give the *opposite* answer under the other convention (`TSTNEG1` 's raw bit pattern is a large unsigned magnitude despite being a negative signed value) - confirms genuine sign-awareness rather than an accidentally- shared implementation. `WITHIN` tested against a wraparound range specifically (`n2` near `$FFFF` , `n3` wrapped past `$0000`), both inside and outside cases - the documented special case ("handles wraparound ranges correctly," via unsigned-offset comparison) its own implementation exists to handle, not just an ordinary non-wrapping range; traced the real implementation by hand first to confirm it computes $(n1-n2) < (n3-n2)$ unsigned before designing the test, not assumed from the glossary's one-line description. `D< / DU<` both tested with equal high cells and different low cells - the documented tie-break case ("low cells unsigned only if [high cells] equal") a naive high-cell-only comparison would get wrong.

Label collisions checked programmatically before insertion - caught one real one this way (`DU`, intended for `DULESSW`, collided with `DUDONE` already inside `DUMP`'s own implementation, part of the unrelated Tools word set - the same class of issue as `MSTAR`/`LSHIFT`/`RSHIFT`/`HOLDS` found in earlier groups). Replaced with `DZ`, re-verified against every label in the file.

****Applied the lesson from the previous group's real mistake immediately this time****: checked the file's `IFEQ`/`IFNE`/`IFNDEF` vs `ENDC` balance right after insertion, before any other verification - confirmed balanced (17/17) on the first

check, meaning both `TSTSELECTOR-4` wrap points (call-list and test-body definitions) got their closing `ENDC` correctly this time, not caught after the fact via the production-build simulation like last time.

Verified: every one of the 15 depth-check values independently re-derived from real arity and confirmed correct; every numeric expected value independently re-verified against fresh computation with proper symbolic-constant resolution; every symbolic reference confirmed to resolve to a real definition; all 15 tests confirmed wired into `TSTCOMPARE`. Full scenario matrix re-run with `TSTSELECTOR`'s fifth value included (12 scenarios: production default x2, `-DUNITTESTS=1` crossed with all five `TSTSELECTOR` values x2 `SERIALPOLL` settings) - all pass; explicitly confirmed `TSTSELECTOR=4` includes only `TSTCOMPARE`'s own bodies, with the other four groups' bodies genuinely absent. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTCOMPARE` tests pass, including the `WITHIN` wraparound cases and the `D</DU<` tie-break cases - the `IFEQ`/`ENDC` balance fix applied at insertion time (learned from the previous group's mistake) held on real hardware, not just in simulation.

- **TSTCTRLFLOW added - control-flow tests (glossary section 3.8, 22 words, 20 tests), wired into TSTRUNNER and gated as TSTSELECTOR 's sixth group (-5). A fundamentally different testing problem from every prior group: these are compile-time, immediate, code-generating words, not runtime operations on operands.** Every implementation traced by hand first (IF / THEN / ELSE , BEGIN / UNTIL / AGAIN / WHILE / REPEAT , RECURSE , DO / ?DO / LOOP / +LOOP , I / J / LEAVE / UNLOOP , EXIT , CASE / OF / ENDOF / ENDCASE , plus PATCH / CCALL / CODECOMMA / ZBRANCH / BRANCH and the return-stack loop-frame layout) - confirmed the "control-flow stack" is just the ordinary data stack (IF pushes (patch-addr, TAGFWD) ; THEN pops and patches), and everything reads/writes through CODEHERE , a plain variable, not tied to a real dictionary entry.

****Test methodology****: each test redirects `CODEHERE` to a new scratch buffer (`TSTCBUF`, 80 bytes), calls the real compile-time words directly (`JSR IF`, `JSR THEN`, `JSR DO`, etc - the actual routines the compiler itself calls, not a hand-simulated imitation), restores `CODEHERE`, then executes the compiled snippet directly. Tests the real interaction between compile-time correctness (right bytes, right patched offsets) and runtime correctness (right control flow) in one coherent check, without needing the outer interpreter, `FIND`, or a dictionary

entry. Given the planned future ANS test suite for broader standards-compliance coverage, this deliberately focuses on the compile/patch mechanism itself, not full end-to-end parsing. ``UNLOOP`` is the one exception - a genuine runtime no-op (bare ``RTS``), tested directly like ``ALIGN``'s own test, needing none of this harness. ``RECURSE`` needed its own scratch (``TSTFHDR``, a fake dictionary header with a known CFA, ``LATEST`` temporarily redirected to it) since it reads ``LATEST`` directly. ``EXIT`` needed the real ``CSP`` variable set correctly (matching what ``:`` does at the start of a real definition), since its own compile-time frame-counting scan depends on it.

Every compiled snippet's exact byte layout and patched offset values were hand-traced before being trusted - not just "it assembled." Two genuinely difficult cases specifically: ``IF`/`ELSE`/`THEN`` (verified `IF`'s patched offset lands exactly at the `ELSE`-body's start, `ELSE`'s own patched offset lands exactly past it) and ``BEGIN`/`WHILE`/`REPEAT`` (verified `WHILE`'s forward offset lands exactly at the final ``RTS``, `REPEAT`'s back-edge lands exactly at ``BEGIN``).

****A real, substantive finding, not just a test-writing detail**:** traced ``DOTEST``'s actual termination logic (simulated the real increment-and-compare-equal check) and confirmed ``DO`` with ``limit=index`` takes a full 65536-iteration wraparound to naturally reconverge on equality - true to the letter of the glossary's "runs at least once" but not in any fast sense. Deliberately left untested rather than build something impractical for a boot-time check or guess at it. ``?DO``'s own equivalent case (``TSTQDOLPEQ``) IS tested - confirmed by the same kind of trace that its skip check happens before the loop is entered at all, fast and safe.

Distinguished ``LEAVE`` from ``EXIT`` precisely by hand-tracing when each actually takes effect: ``LEAVE`` only sets a flag, checked at the **next** ``LOOP``, so the current iteration's trailing code still runs (``TSTLEAVE``'s sum is $0+1+2+3=6$, including the iteration where ``I`=3` triggers ``LEAVE``); ``EXIT`` fires immediately where encountered, so the same-shaped test (``TSTEXIT``) gives  $0+1+2=3$ , not including ``I`=3`'s own iteration. Getting this backwards would have been a real, silent bug in the tests themselves, not caught by any structural check.

Error-path tests (``TSTTHENZ`/`TSTUNTILZ`/`TSTENDOFZ``, covering "throws -22 on tag mismatch") reuse the established ``CATCH`` pattern from earlier sections directly - confirmed by reading ``CFERR`` that the error path never touches ``CODEHERE`` at all, so

no redirect is needed for these three specifically, unlike every other test in this group.

****Two real mistakes caught and fixed during this batch, not after**:** (1) three depth-check sign errors (``TSTIFT2``, ``TSTRECUR``, ``TSTDOLP``) - the same class of mistake as earlier sections - caught by systematically re-deriving every depth value from actual cells-in/cells-out arithmetic rather than trusting the first pass, before any insertion happened. (2) a genuinely new mistake: concatenating the 14 draft files via a shell glob (``tstctrl_part*.txt``) sorted lexicographically, not numerically, scrambling the test bodies' order in the first insertion attempt (``part1`` followed by ``part10`-`part14``, then ``part2`-`part9``) - likely harmless for actual assembly (forward references are fine), but sloppy and not what was intended. Caught via a reference-verification check that found expected words apparently missing from the inserted block, traced precisely to the real cause (wrong extraction boundary revealing the reordering, not missing content) rather than assumed fixed, and the whole insertion redone with correct numeric ordering before delivering anything.

Label collisions checked programmatically before insertion - caught 6 real ones this way, all simple prefix reuse from this session's own earlier sections (``DZ``, ``UL``, ``PL``, ``DL``, ``TZ``, ``UZ``), replaced with ``EO``, ``UO``, ``PO``, ``DW``, ``T3``, ``U3``, each re-verified against every label in the file.

Verified: zero duplicate symbols, dictionary chain still 224 entries intact across all 14 relevant ``SERIALPOLL`x`UNITTESTS`x`TSTSELECTOR`` scenario combinations (production default, and ``-DUNITTESTS=1`` crossed with all six ``TSTSELECTOR`` values and both ``SERIALPOLL`` settings); every one of the 20 tests' depth-check values independently re-derived from reality and confirmed correct; every symbolic and word reference confirmed to resolve to a real definition (a small number of apparent false positives - bare ``F`/`O`` tokens - traced to the quoted character literals in ``TSTRECUR``'s fake header construction, not real gaps); all 20 tests confirmed wired into ``TSTCTRLFLOW``; explicitly confirmed ``TSTSELECTOR=5`` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all ``TSTCTRLFLOW`` tests pass - including the ``IF`/`ELSE`/`THEN`` dual-branch cases, ``BEGIN`/`WHILE`/`REPEAT``'s two-offset patch, ``RECURSE``'s fake-header mechanism, every ``DO``-family loop variant, the ``LEAVE`-vs-`EXIT`` timing distinction, ``CASE``'s match/no-match paths, and all three tag-mismatch error paths.

This was this session's most structurally complex test group - the redirected-`CODEHERE`, compile-then-execute harness holds up on real hardware, not just in simulation, validating the whole methodology for any future control-flow-adjacent work.

- **TSTDEFWORDS added - defining-words tests (glossary section 3.9, 17 words, 12 tests since VALUE / TO and IS / ACTION-OF are each combined into one test), wired into TSTRUNNER and gated as TSTSELECTOR 's seventh group (-6). A harder testing problem than section 3.8: these words don't just compile code, they parse a name from the input source and build real dictionary headers.** Every implementation traced by hand first (HEADER , shared by : / CREATE ; the CREATE / DOES> trampoline mechanism via DODOES / SETDOES ; VARIABLE / CONSTANT / 2VARIABLE / 2CONSTANT / BUFFER: / VALUE / TO / DEFER / DEFER@ / DEFER! / IS / ACTION-OF / MARKER) - confirmed HEADER calls WORD directly to parse the new word's name, something nothing in section 3.8 needed.

****Harness extended beyond section 3.8's `CODEHERE`-only redirect**:** every test now also redirects `DPHERE` (a new 40-byte scratch dictionary buffer, `TSTDBUF`), `VARHERE` (a new 20-byte scratch buffer, `TSTVBUF`), and `SRCADDR`/`SRCLEN`/`TOIN` (a fake name, "TESTWD", reused safely across every test here since each redirect/restore cycle is fully isolated) - and established a new convention: save `CODEHERE` immediately before calling the defining word (this becomes the newly-defined word's own CFA), so its compiled trampoline can be executed afterward to verify runtime behavior.

Worked out `SETDOES`'s "double return" mechanism precisely by hand before trusting it: confirmed a direct call into the compiled "JSR SETDOES" instruction itself (rather than building an intermediate wrapper word) naturally supplies exactly the two return-stack levels it expects - its own `JSR` call provides one, this test's own call provides the other - without extra scaffolding. Also traced why bare `CREATE` (never followed by `DOES>`) isn't tested standalone: its placeholder behavior field is a genuine compiled `JSR DOESRT0` instruction, while every other defining word compiles a raw address there instead (via `CODECOMMA`, not `CCALL`) - `DODOES` reads that field as data either way, which only produces a valid jump target for the raw-address form.

****A real, structurally significant bug found and fixed during this batch, not after**:** `HEADER` writes to the real `LATEST` variable unconditionally - unlike `CODEHERE`/`DPHERE`/`VARHERE`,

there's no redirect for it. The first 9 tests written (`TSTVAR` through `TSTDEFER2`) didn't save/restore it, which would have left the real dictionary's `LATEST` pointer corrupted, pointing at scratch memory, after every one of them ran - a genuine risk to the real dictionary, not just a test-isolation nicety. Caught by tracing `HEADER`'s own code directly rather than assuming symmetry with the other three pointers, and fixed across all 8 affected files before any insertion happened, re-verified afterward by direct inspection of each save/restore block's exact placement, not just re-trusted.

`MARKER`'s own test needed the opposite execution order from every other test in this group: `DOMARKER` (confirmed by reading it directly) also writes to the real `DPHERE`/`CODEHERE`/`VARHERE`/`LATEST` unconditionally, so that test executes the marker word **while** still redirected, restoring the real environment only afterward - reversing this order would have corrupted the real dictionary pointers with scratch addresses.

Label collisions checked programmatically before insertion - zero found this time (all chosen prefixes verified safe on the first attempt). Learned the lesson from section 3.8's ordering mistake and applied it from the start this time: used numeric (not lexicographic) file ordering throughout, and checked the `IFEQ`/`ENDC` balance immediately after insertion rather than discovering a problem later - both came back clean on the first attempt.

Verified: every one of the 12 depth-check values independently re-derived from real arity and confirmed correct (one sign error caught and fixed during writing, before insertion); every symbolic and word reference confirmed to resolve to a real definition (a handful of apparent false positives - bare character-literal tokens from the repeated "TESTWD" name construction - traced to their real cause, not just dismissed); all 12 tests confirmed wired into `TSTDEFWORDS`. Full scenario matrix re-run with `TSTSELECTOR`'s seventh value included (16 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=6` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTDEFWORDS` tests now pass, closing out three real, distinct issues found via real hardware over several rounds - a byte-overflow needing `LBNE` (a genuine `BNE`/`ISFAIL` short-branch overflow, unrelated to this section's own logic), a `WORD`-overwrites-`CODEHERE` corruption bug in `TSTVALTO`/`TSTISOF`'s second phases, a `SETDOES`/`LATEST`-ordering bug in `TSTCRDOES`,

and a wrong comparison target in `TSTMARKER`'s own verification logic (not a bug in `MARKER` itself). This section's own complexity - name-parsing, real header-building, the trampoline mechanism - proved out exactly the concern flagged when this group was first delivered.

- **RESOLVED - real assembly-blocking bug found by the user's actual lwasm run, not caught by this session's own verification: four `BNE` instructions in `TSTVALTO` and `TSTISOF` exceeded short-branch range (Byte overflow).** Same bug class as `DULINE`'s own already-documented fix elsewhere in this file - `TSTVALTO` and `TSTISOF` are this section's two combined, two-round tests (`VALUE + TO`, `IS + ACTION-OF`), and their earliest checks sit far enough before the shared `FAIL` label - past the entire second round's own code - to exceed a short branch's range. Fixed the exact four flagged lines (verified by line number, not by text pattern, since identical `BNE VTFAIL / BNE ISFAIL` text appears multiple times in each test) by converting them to `LBNE` - confirmed this is an established, already-proven mnemonic in this file rather than assumed. Trusted the assembler's own precise report rather than pre-emptively converting every `BNE VTFAIL / BNE ISFAIL` occurrence - 7 of the 11 total instances were not flagged and were left as-is, matching their genuinely shorter distance to the target. Confirmed by reasoning through it that the fix doesn't newly endanger those remaining short branches either: the added bytes from converting `BNE` to `LBNE` sit entirely before every one of them, which doesn't change their own distance to the shared target at all.

Verified: `IFEQ`/`ENDC` balance unaffected (still 21/21, re-checked after this edit rather than assumed unaffected by a pure branch-instruction change); zero duplicate symbols, dictionary chain still 224 entries intact across all 16 relevant `SERIALPOLL`x`UNITTESTS`x`TSTSELECTOR` scenario combinations; exact counts of `LBNE`/remaining-`BNE` confirmed (4 and 7 respectively, matching the fix precisely). Byte-exact split-file reassembly confirmed. This specific fix awaits its own confirmation via a real lwasm run - the user's own report is what caught it, this response addresses exactly what was reported.

- **RESOLVED - real crash found by the user's actual MAME run, not caught by this session's own verification: `TSTVALTO` jumped into invalid memory at `TSTCBUF`'s own address on its second `JSR ,x`.** Root cause confirmed by reading `WORD`'s own code directly rather than guessed at: `WORD` writes its parsed- token output directly at `CODEHERE`, by design ("matches the ANS transient-region contract, where `WORD`'s region is expected to be overwritten by whatever gets compiled/allocated next" - `CODEHERE` itself isn't advanced by this). `TSTVALTO` redirects `CODEHERE` to `TSTCBUF` a

second time for its own TO phase's name-parse, after TSTCBUF already held VALUE 's own compiled trampoline from the first phase - WORD 's own internal parse (called by TOW) silently overwrote it. Confirmed precisely against the user's own memory dump: a length byte (06) followed by "TESTWD" sitting exactly at TSTCBUF 's address, where the trampoline's real JSR DODOES opcode should have been - not a vague "probably right," the exact byte pattern WORD 's own documented behavior predicts.

Checked every other test in this section for the same pattern before assuming this was isolated - confirmed it's structurally possible only for this section's two "combined, two-phase" tests (`TSTVALTO`, `TSTISOF`), since every other test does a single name-parse with nothing to conflict with. Found `TSTISOF` had a related but more serious version of the same bug: its `IS` and `ACTION-OF` phases had *no* `CODEHERE` redirect at all (not even the wrong one) - meaning `WORD`'s own internal parse there would have written into the real, unredirected system `CODEHERE`, unsafe during boot-time testing before `COLD` has set it to anything meaningful. The `IS` phase specifically mattered more than `ACTION-OF`'s, since it runs *before* this test's own execution step - the same class of corruption that crashed `TSTVALTO`, just not yet triggered by MAME's own test run order.

Fixed with a new, dedicated scratch buffer (`TSTCBUF2`, 20 bytes) used for every "second phase" name-parse in both tests, keeping it structurally distinct from `TSTCBUF` so a later parse can never overwrite an earlier phase's still-needed compiled trampoline. Confirmed `TOW`/`ISW`/`ACTIONOF`'s own interpreting-mode paths don't touch `DPHERE`/`VARHERE` (traced directly - they only call `WORD`+`FIND`+`TOBODY` plus a direct memory store, none of which touches those two), so the fix only needed to cover `CODEHERE`.

Verified: `IFEQ`/`ENDC` balance unaffected (21/21, re-checked rather than assumed); both tests' own depth-check values re-confirmed unchanged by this purely corrective edit (`[2,2]` for `TSTVALTO`, `[2]` for `TSTISOF`, matching their state before this fix); `TSTCBUF2` confirmed correctly declared and referenced exactly where intended; zero duplicate symbols, dictionary chain still 224 entries intact across all 16 relevant `SERIALPOLL`x`UNITTESTS`x`TSTSELECTOR` scenario combinations. Byte-exact split-file reassembly confirmed. This specific fix awaits its own confirmation via MAME - the user's own crash report is what caught the root cause, this response addresses exactly what was found, including the related `TSTISOF` issue

the crash report itself didn't directly point to.

- **RESOLVED - real bug found by the user's actual MAME run, pinpointed precisely from a register/memory dump, not guessed at: `TSTCRDOES` failed because `SETDOES` patched the wrong word entirely.** Root cause: this test's own restore block set `LATEST` back to its real, original value *before* triggering `SETDOES` (via a direct call into the compiled "JSR SETDOES" instruction) - but `SETDOES` (confirmed by re-reading its own code) reads `LATEST` directly to find which header's behavior field to patch. With `LATEST` already restored, `SETDOES` patched whatever real word `LATEST` actually pointed to, not `TESTWD` - leaving `TESTWD` stuck on its original `DOESRT0` placeholder (push PFA, return - a plain `VARIABLE` -like behavior), never running @ 1+ at all.

Confirmed precisely against the user's own register/memory dump before fixing anything: the crashing ``CMPD #6`` compared against ``$022C``, exactly matching ``TESTWD``'s own computed PFA address (``TSTCBUF+7``, the self-referential-PFA pattern already documented elsewhere in this file) - the exact value ``DODOES`` pushes under ``DOESRT0``'s default behavior, not a coincidence.

Fixed by reordering: ``LATEST`` now stays pointed at ``TESTWD`` until **after** the ``SETDOES`` trigger completes, only restored to its real value afterward - ``CODEHERE`/`DPHERE`/`VARHERE`/`SRCADDR`/`SRCLEN`/`TOIN`` can still restore early, since ``SETDOES`` doesn't depend on any of those (confirmed by reading its code - it only touches ``LATEST`` and the computed behavior-field address).

Verified: ``IFEQ`/`ENDC`` balance unaffected (21/21); this test's own depth-check value unchanged (``2``); zero duplicate symbols, dictionary chain still 224 entries intact across all 16 relevant scenario combinations. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation.

- **RESOLVED - `TST2VAR` confirmed passing after the `TSTCRDOES` fix above (the user's own re-run confirmed it), without any change of its own needed.**
- **RESOLVED - `TSTMARKER`'s failure, precisely diagnosed from the user's own register/memory dump, turned out to be a bug in this test's own verification logic, not in `MARKER` / `DOMARKER` themselves - the real dictionary/compile mechanism was already correct.** The failing comparison was `TSTCSAV2` (`CODEHERE` right after executing the marker) against `TSTMKCOD` (a snapshot of `CODEHERE` captured right *after* `MARKERW` finished building its own header) - decoded precisely against the user's own

memory dump: TSTCSAV2 held \$0225 (TSTCBUF exactly), TSTMKCOD held \$0234 (TSTCBUF +15, the full compiled size of MARKERW 's own trampoline+snapshot structure - confirmed byte for byte by re-reading MARKERW 's complete implementation, not estimated).

Traced why these genuinely differ rather than assuming a simple off-by-some-amount error: `MARKERW`'s own internal snapshot (`MKDP`/`MKCODE`/`MKVAR`/`MKLATEST`) is taken at its very first four instructions, before `HEADER` or any of its own compiling runs at all - meaning `DOMARKER` correctly restores to the state *before* the marker word itself was created, not the state right after. This is exactly `MARKER`'s own documented behavior ("forgets itself too", not just what comes after it) - this test's own comparison target (`TSTMKCOD`, captured too late) was simply wrong from the start, not the mechanism it was testing.

Fixed by comparing directly against `TSTCBUF`/`TSTDBUF`/`TSTVBUF` (the real, known pre-creation values) instead, and removed the now-unnecessary `TSTMKCOD`/`TSTMKDP`/`TSTMKVAR` capture step and their now-unused scratch declarations entirely, rather than leaving dead, misleading allocations behind. Updated this test's own header comment, which had described the old (wrong) design, to accurately reflect the fix and why it's correct.

Verified: `IFEQ`/`ENDC` balance unaffected (21/21); this test's own depth-check value unchanged (`0`); confirmed no remaining code references to the removed constants (only an explanatory comment, which is harmless); zero duplicate symbols, dictionary chain still 224 entries intact across all 16 relevant scenario combinations. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation.

- **** TSTRETSTACK added - return-stack tests** (glossary section 3.3, 6 words, 4 tests since `>R / R>` and `2>R / 2R>` are each combined into one round-trip test), wired into `TSTRUNNER` and inserted in the source **in glossary order** (right after `TSTSTACK` 's own tests, before `TSTSARITH` 's), per explicit request - **not** at the end of the file like every prior group. Gated as `TSTSELECTOR` 's eighth value (`-7`), the next available number rather than a value matching this glossary-order position - per explicit request, `TSTSELECTOR` 's own numbering stays as-is for now, with a renumbering-to-match-glossary-order pass explicitly deferred to a future request, not attempted here.

A genuinely different testing problem from every prior group:

these words operate directly on the return stack (`S`) - the same stack holding real subroutine return addresses, including this test framework's own. Traced each word's `PULS`/`PSHS` juggling by hand first: `>R`/`R>`/`2>R`/`2R>` all temporarily lift the caller's own return address off `S`, do the actual move, then restore it on top - so a moved value ends up nested one level inside the *current* subroutine's own return-stack frame, safely retrievable by a later `R>`/`2R>` in the same body. Every `>R`/`2>R` in these tests is balanced by a matching `R>`/`2R>` before this group's own `RTS` - leaving one unbalanced would corrupt the path back to whatever called this test, a real risk specific to this section that no prior group had to guard against.

`TSTTOR`/`TSTTWOTOR` each include an intermediate depth check (right after the move-out, before the move-back) specifically because a pure round-trip check could pass even if both halves were broken no-ops - the value would simply never have left. `TSTRFETCH`/`TSTTWORFETCH` verify "non-destructive" for real: peek, then retrieve the original afterward and confirm it's still correct, not just that the peek itself returned the right value once.

Label collisions checked programmatically before insertion - caught one real one this way (`TW`, colliding with `TSTBWR`'s own prefix from section 3.6's `BEGIN`/`WHILE`/`REPEAT` test), replaced with `T2F`, re-verified against every label in the file. Caught and fixed the same `IFEQ`/`ENDC` mistake as section 3.8's own first attempt - forgot the second wrap's closing `ENDC` around the test bodies - but this time caught it immediately via the balance check run right after insertion, not discovered later via a broken production-build simulation.

Verified: every one of the 6 depth-check values (across 4 tests, two of which check both an intermediate and a final depth) independently re-derived from real arity and confirmed correct; every real word reference (`TOR`/`FROMR`/`RFETCH`/`TWOTOR`/`TWOFROMR`/`TWORFETCH`) confirmed to resolve, with the apparent "missing" list from the check itself traced to false positives (register names, a mnemonic omitted from the exclusion list, and this batch's own local `FAIL`/`DONE` labels) rather than dismissed without checking; all 4 tests confirmed wired into `TSTRETSTACK`. Full scenario matrix re-run with `TSTSELECTOR`'s eighth value included (18 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=7` includes only this group's own bodies, with `TSTSTACK`'s and `TSTCTRLFLOW`'s bodies genuinely absent. Byte-exact split-file reassembly

confirmed, including the correct glossary-order placement (confirmed via label position, not just presence).
MAME-CONFIRMED: the user reports all `TSTRETSTACK` tests pass - including both round-trip intermediate checks and both non-destructive-peek verifications - confirming every `>R`/`2>R` in this group stayed correctly balanced against a matching `R>`/`2R>` on real hardware, with no corruption of the return path back out of this test group.

- ****TSTSYSIO added - System/Console I/O tests (glossary section 3.1, 12 words, 7 tests), wired into TSTRUNNER and inserted **first** in the source, ahead of TSTSTACK, per explicit request - matching this session's established practice of placing new groups in glossary order in the source even while TSTSELECTOR's own numbering stays non-sequential for now (gated as its ninth value, -8, the next available number).**

Six of the section's 12 words are deliberately not tested, confirmed necessary by reading their own implementations directly, not assumed from the glossary descriptions alone:
`KEY` genuinely spins forever without real input (`BEQ KEY` branches back to itself); `ACCEPT` calls `KEY` directly in its own main loop, and `EXPECT`/`QUERY` both call `ACCEPT` - all four inherit the same block, and none can complete during automated, headless boot-time testing where no real input will ever arrive. `ABORT` falls straight through into `QUIT`, which resets the return stack (`S`) to `RP0` - discarding this test framework's *entire* call chain and hijacking the boot sequence into the interpreter's own top-level loop; calling either directly would break the whole boot, not just fail one test. This is a materially different situation from every prior section's own deliberate coverage gaps (`DO`'s `limit=index` case, bare `CREATE`) - those were single edge cases within an otherwise-testable word; here, half the section's words are fundamentally incompatible with this framework's own execution model, not just one case of each.

The other 6 (`KEY?`, `EMIT`, `TYPE`, `CR`, `SPACE`, `SPACES`, the last split into two tests for its `n<=0` no-op case) are confirmed safe - and already proven so, since this whole test framework has used `TYPE`/`CR` (via `TSTREPORT`) thousands of times already across every prior test group without incident. Each test here verifies the real output-queuing mechanism (`OUTHEAD`/`OUTBUF`) directly - reading the actual character(s) back out of the output ring buffer to confirm they were genuinely queued, not just that the call returned without crashing.

****Two `FCB`-length/string mismatches caught by a programmatic check during writing, not left to be found later**:** `TSTSPACESZ` and `TSTSPACES`'s own name-length bytes were miscounted by hand (11 vs the true 10, 10 vs the true 9) - found by systematically re-checking every declared length against the actual string for every test in this batch, not just the ones that happened to look suspicious.

Label collisions checked programmatically before insertion - caught 3 real ones this way (`KQFALSE`, `SPDONE`, `TYDONE`, all internal labels already inside `KEY?`'s, `SPACES`'s, and `TYPE`'s own real implementations - the same class of issue as `MSTAR`/`LSHIFT`/`RSHIFT`/`HOLDS`/`DUMP` found in earlier groups), replaced with `KY`, `SC`, `TE` respectively, each re-verified against every label in the file.

****The same missing-second-`ENDC` mistake as sections 3.3 and 3.8's own first attempts, and a separately forgotten `TSTRUNNER` wire-up**** - both caught and fixed before delivery: the `IFEQ`/`ENDC` balance check (run immediately after insertion, per this session's now-standard practice) caught the missing `ENDC` right away; a follow-on check specifically confirming `TSTSYSIO` appears in `TSTRUNNER`'s own call list (not just assumed done from memory) caught that it had been omitted entirely.

Verified: every one of the 7 depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state (not just the pre-insertion draft); every real word reference (`KEYQ`/`EMIT`/`CRW`/`SPACEW`/`SPACESW`/`TYPE`) confirmed to resolve, including `CRW` specifically, which an over-aggressive local-prefix filter in the verification script itself initially hid - confirmed directly via a literal `JSR CRW` count rather than trusting the flawed automated check; all 7 tests confirmed wired into `TSTSYSIO`, and `TSTSYSIO` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s ninth value included (20 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=8` includes only this group's own bodies. Byte-exact split-file reassembly confirmed, including the correct glossary-order placement ahead of `TSTSTACK` (confirmed via label position).

****MAME-CONFIRMED**:** the user reports all `TSTSYSIO` tests now pass under the real, active `SERIALPOLL=1` build - including the terminal output itself visibly showing each test's own emitted character(s) (`A` before `TSTEMIT`, a leading space

before `TSTSPACE`, three before `TSTSPACES`, `AB` before
`TSTTYPE`) intermixed with the pass/fail reporting exactly as
expected, confirming both the `EMITCH`-based and the
`SERIALPOLL`-conditional fixes hold correctly on real hardware,
not just in simulation.

- **RESOLVED - real bug found by the user's actual MAME run, not caught by this session's own verification: 5 of `TSTSYSIO`'s 7 tests failed, despite the characters genuinely appearing in the terminal output.** Root cause confirmed by directly re-checking `EMIT`'s own code, having *already found this exact pattern* once already for `KEY` earlier in the same investigation and failing to check whether it applied to `EMIT` too: `EMIT` has two entirely different implementations, gated by `SERIALPOLL` the same way `KEY`'s own two variants are - the interrupt-driven one (`SERIALPOLL=0`) queues into a ring buffer (`OUTHEAD / OUTBUF`), but the polling one (`SERIALPOLL=1` , confirmed the currently active build) writes directly to `ACIADR` and never touches `OUTHEAD / OUTBUF` at all. Every one of these 5 tests assumed the interrupt-driven mechanism unconditionally - so the checks always failed under the actual active build, despite `EMIT` genuinely working (the character really was transmitted, exactly as the user's own terminal output showed).

Fixed with a mix of two approaches, chosen per test based on what's genuinely verifiable in each mode: ``TSTEMIT`/`TSTSPACE`` (single-`EMIT`-call` tests) now check `EMITCH`` universally, with no `SERIALPOLL`` conditional needed at all - confirmed by reading both `EMIT`` variants that `EMITCH`` is set unconditionally, as the very first step, by both, before either mode-specific branch even begins. ``TSTCR`/`TSTSPACES`/`TSTTYPE`` (multiple-`EMIT`-call` tests, where `EMITCH`` only reflects the *\*last\** call) now use an explicit `IFEQ SERIALPOLL`/`ELSE`/`ENDC`` split: the full ``OUTHEAD`/`OUTBUF`` check for interrupt-driven mode, and an honest, narrower check (last character via `EMITCH``, plus the pre-existing stack-depth check) for polling mode - a genuine, accepted limitation of a direct hardware write with no buffering to inspect, not something worth contorting around.

****A second, structural mistake caught and fixed during this same round, before delivery**:** the first attempt put `IFEQ SERIALPOLL`` directly on the same source line as  ``TSTCR:`/`TSTSPACES:``'s own label ( ``TSTCR: IFEQ SERIALPOLL``) - a pattern that appears nowhere else in this entire file, where every other `IFEQ`` sits on its own dedicated line. This confused this session's own ``IFEQ`/`ENDC`` balance-checking script (whose regex requires whitespace, not a label, before `IFEQ``), which is itself just a verification-script limitation - but rather than

trust that the real assembler would handle the untested pattern correctly, restructured both to put the label and `IFEQ` on separate lines, matching the proven convention used everywhere else in this file. Re-verified via an explicit, full-file, line-by-line nesting trace (not just a raw open/close count match, which can't by itself distinguish correct nesting from a coincidentally-equal but wrongly-ordered mismatch) that the result is genuinely, correctly balanced end to end, with zero negative-depth excursions anywhere in the file.

Verified: all 7 depth-check values re-confirmed unchanged by these purely corrective edits, using a more robust per-test isolation method after the original 200-character-window heuristic in this session's own check turned out too narrow once the new `SERIALPOLL` blocks pushed some tests' own depth captures further from their label; explicitly confirmed via simulation that `TSTCR`'s two branches genuinely diverge under each `SERIALPOLL` setting (`OUTHEAD` check present only when `SERIALPOLL=0`, `EMITCH` check present only when `SERIALPOLL=1`, mutually exclusive as intended) rather than assumed correct from the source alone. Full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value (20 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own report and register/memory-level detail is what caught the root cause; this response addresses exactly what was found.

- **RESOLVED - TSTSELECTOR 's comparison values renumbered sequentially starting at 0, matching glossary/source order, per explicit request.** The prior numbering had grown ad hoc across nine separate additions (each new group taking "next available" rather than matching its glossary-order position, per earlier explicit direction deferring exactly this renumbering to a later request). New mapping, all confirmed matching each group's own source position (which already matched glossary order): TSTSYSIO (3.1) 8 -> 0, TSTSTACK (3.2) 0 -> 1, TSTRETSTACK (3.3) 7 -> 2, TSTSARITH (3.4) 1 -> 3, TSTDARITH (3.5) 2 -> 4, TSTLOGIC (3.6) 3 -> 5, TSTCOMPARE (3.7) 4 -> 6, TSTCTRLFLOW (3.8) 5 -> 7, TSTDEFWORDS (3.9) 6 -> 8.

Applied via a position-bounded replacement, not a blanket find-replace across the whole file - each group's own two `IFEQ` TSTSELECTOR-N` occurrences were located and replaced only within that specific group's own source region (bounded by its own label through the next group's), avoiding any risk of a sequential renumber colliding mid-process (e.g. an already-renumbered `-1` accidentally being renumbered again when a

different group's original ``-1`` gets processed later). All 18 replacements (2 per group x 9 groups) confirmed exact.

Checked the rest of the file for stale comment references to specific ``TSTSELECTOR`` values before considering this complete - found none needing correction: the one literal ``TSTSELECTOR=2`` appearing in the source is a generic syntax example in the unit-test framework's own header comment, not tied to a specific group's identity, and left as-is. The checklist's own historical entries documenting each group's original, non-sequential value at the time it was added are left untouched as well, being a chronological record of what was actually done, not a live, always-current description of the file's present state.

Verified: ``IFEQ`/`ENDC`` balance unaffected (31/31, unchanged from before - this edit only changed values within existing pairs, not the pairs themselves) and correct nesting re-confirmed via a full-file line-by-line trace; full scenario matrix re-run crossing both real ``SERIALPOLL`` settings with ``TSTSELECTOR`` 0 through 8 (20 scenarios total) - all pass; explicitly confirmed via simulation, not just inferred from the source text, that every ``TSTSELECTOR`` value 0-8 now compiles exactly the test group matching its own glossary-section ordinal (a representative test spot-checked per group). Byte-exact split-file reassembly confirmed. Not yet confirmed via MAME, though given this touched only which literal number selects which already-working group - not any group's own logic - the risk profile is low relative to this session's typical finds.

- **** TSTCOMPWORDS added - compiling-words tests** (glossary section 3.10, 13 words, 17 tests since several get separate cases: ``` found/not-found, `[ ]` compiling/interpreting state, `POSTPONE` normal/immediate word, `SLITERAL` compiling/ interpreting state, `ABORT` false/true flag), wired into `TSTRUNNER` and inserted last in the source (after `TSTDEFWORDS`), matching glossary order - the first new group whose `TSTSELECTOR` value (`-9`) matches its glossary-order position from the moment it was added, since the prior renumbering (above) already made every earlier group's value sequential.

****`POSTPONE` confirmed this section's hardest case, as anticipated****: a genuine two-level compile-time mechanism. Traced its own code first rather than assumed from the glossary's one-line description - for a normal (non-immediate) word, it compiles ``[LIT xt][JSR COMPILE,]`` into the current definition, not a direct call - standard ANS semantics (``POSTPONE`` appends a word's own `*compilation*` semantics to the

current definition; for an ordinary word that semantics IS "compile a call to it", so appending it only makes sense if the current definition, when it later runs, itself performs that compiling action against whatever is being compiled at that later point - not immediately). For an immediate word, by contrast, it compiles a direct call to the immediate word's own xt - simpler, confirmed by reading `FIND`'s own code that it pushes `1` for immediate words, `-1` otherwise, and `POSTPONE` branches on exactly that. Given the genuine complexity of fully executing the normal-word case's second level (which would need a further redirected `CODEHERE` and a further round of execute-and-verify), both tests instead verify the compiled byte sequence directly - still a meaningful, unambiguous check of the actual mechanism, not a workaround.

`ABORT`` confirmed safely testable for both flag cases, not just assumed from its relationship to `ABORT`: traced `DOABORTQUOTE`'s own runtime code directly and confirmed the true-flag path uses `THROW -2` internally, not the raw, never-returns `ABORT` mechanism plain `ABORT` itself uses (still deliberately untested, per section 3.1's own reasoning - `ABORT` resets the return stack and would hijack this whole test framework's boot sequence). Since `ABORT``'s own exception is genuinely catchable, both its false-flag (resumes normally, confirmed by checking execution lands exactly past the inline message) and true-flag (`CATCH`-wrapped, confirms `-2`) cases are tested directly.

`LITERAL` and `>BODY` had already been used extensively as internal helpers in sections 3.8/3.9's own tests, but got dedicated, direct tests here too, per this section's own coverage - each confirms genuine end-to-end behavior (executing a compiled snippet, or computing a real offset), not just reconfirming what their existing indirect use already implied.

Three `FCB`-length/string-length mismatches caught by the now-standard programmatic check during writing, not left to be found later: `TSTCOMPCOMMA` (declared 13, actually 12), `TSTPOSTPONE1` (declared 13, actually 12), `TSTBRACKTICK1`/`TSTBRACKTICK2` (each declared 14, actually 13), `TSTSLITERAL1`/`TSTSLITERAL2` (each declared 13, actually 12) - checked every single test's own name length individually and immediately after writing it, rather than batching the check to the end, given how frequently this specific mistake had already recurred across earlier sections.

Label collisions checked programmatically before insertion -

caught 2 real ones this way (`LTFAIL`/`LTDONE`, colliding with section 3.7's own `TSTLT`; `EXFAIL`/`EXDONE`, colliding with section 3.8's own `TSTEXIT`), replaced with `LI` and `XQ` respectively, each re-verified against every label in the file. Checked the `IFEQ`/`ENDC` balance immediately after insertion, per this session's now-standard practice following the repeated missing-second-`ENDC` mistake in earlier sections - balanced (33/33) and correctly nested end to end on the first attempt this time.

Verified: every one of the 17 depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with the apparent "missing" list (a handful of character-literal tokens, a mnemonic omitted from the exclusion list, and this batch's own local labels) traced to its real, harmless cause rather than dismissed without checking; all 17 tests confirmed wired into `TSTCOMPWORDS`, and `TSTCOMPWORDS` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s tenth value included (22 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=9` includes only this group's own bodies, with `TSTDEFWORDS`'s and `TSTSYSIO`'s bodies genuinely absent. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTCOMPWORDS` tests now pass, including `TSTTOBODY` after its own fix (logged separately above) - the `POSTPONE` normal/immediate-word cases, both `ABORT` flag cases, and every other test in this section all confirmed correct on real hardware, not just in simulation.

- **RESOLVED - real bug found by the user's actual MAME run, pinpointed precisely from a register/memory dump: `TSTTOBODY` failed because the test's own assumption about `TOBODY`'s return value was wrong - `TOBODY` itself is correct.** Root cause confirmed by re-reading `TOBODY`'s own code directly rather than trusting the earlier trace: it computes `xt+5`, then *dereferences* it (`LDD ,X`), returning the value stored there - not `xt+5` itself. This test's own comment had described it as "a pure, fixed-offset computation... not dependent on what's actually at that offset" - simply wrong, and the test never initialized `TSTCBUF+5` to anything, so it compared the real returned value against whatever garbage happened to already be sitting there (`$BDEB` in the user's own register dump, exactly matching the two bytes shown at that address in the accompanying memory dump - not a coincidence, the precise mechanism the trace predicted).

Reconciled this directly against why section 3.9's `DEFER@`/

`DEFER!`/`TO`/`IS` tests, which also rely on this exact same `TOBODY` and already passed on real hardware, never surfaced this: those tests set up real trampolines (via `VALUEW`/`DEFERW`), whose own compile-time code stores a genuinely meaningful value at `xt+5` (for `DEFER`, the current target xt; for `VALUE`, the stored value itself) - for this system's own trampoline design, "the body" *is* the value stored at `xt+5`, not `xt+5`'s own address. `TOBODY` was never broken; this one test's own understanding of what it does was.

Fixed by setting a known value at `TSTCBUF+5` before calling `TOBODY`, and verifying it returns exactly that value - not the address. Also corrected the test's own header comment, which had baked in the same wrong assumption.

Verified: `IFEQ`/`ENDC` balance unaffected (33/33); this test's own depth-check value unchanged (`0`, since the fix only adds a memory write before the stack-tracking baseline, touching nothing on `U`); full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value 0-9 (22 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own register/memory-dump-level detail is what let this be pinpointed precisely rather than guessed at.

- **** TSTMEMORY added - memory tests (glossary section 3.11, 22 words, 16 tests since several are combined into round-trip tests that can only be meaningfully verified together - @ / ! , C@ / C! , 2@ / 2! , and the CODEHERE -region words (, / C, / ALLOT / HERE) and VARHERE -region words (V, / VC, / VALLOT / VHERE) each combined into one sequential walk-through per region), wired into TSTRUNNER and inserted last in the source (after TSTCOMPWORDS), matching glossary order, gated as TSTSELECTOR 's eleventh value (-10) - continuing the now- sequential numbering established by the earlier renumbering.**

Mostly simpler, direct memory-access words than the last few sections (no compile-time behavior, no name-parsing), but `,`/`C`,`/`ALLOT`/`HERE`/`PAD` touch `CODEHERE` and `V`,`/`VC`,`/`VALLOT`/`VHERE` touch `VARHERE`, so those still needed the same redirect-based safety established for sections 3.8/3.9; `TSTCBUF`/`TSTVBUF` reused as general scratch for the simpler fetch/store tests too, safe since each `TSTSELECTOR` group runs exclusively of every other.

****`MOVE` confirmed via its own code to genuinely choose its copy direction by comparing addresses**, delegating to `CMOVE`**

(low-to-high) when `dst <= src` and `CMOVE>` (high-to-low) otherwise, specifically to stay overlap-safe. Rather than test only a trivial non-overlapping case, wrote two dedicated overlapping-region tests (`TSTMOVE2`/`TSTMOVE3`, one for each direction) - each expected result hand-derived byte by byte and then independently re-derived via a small Python simulation before writing the actual test, not guessed at or assumed correct from the reasoning alone. `CMOVE`/`CMOVE>` themselves are tested with non-overlapping regions only, given they are explicitly documented as not overlap-safe in one direction each - `MOVE`'s own existence is precisely the mechanism for handling the overlapping case correctly, not something to duplicate in `CMOVE`'s own tests.

`TSTCFETCHSTORE` (`C@`/`C!`) specifically pre-fills the cell with a distinctive nonzero high byte before testing, to verify both that `C!` genuinely only touches the low byte (checked via a direct memory read, not just inferred from the later `C@`) and that `C@` genuinely zero-extends (not just returns whatever was already there) - a plain round trip alone could not have distinguished either from a byte-store/byte-fetch pair that happened to leave the high byte untouched by coincidence.

Zero `FCB`-length/string-length mismatches survived to the collision-check stage this time - all nine caught (`TSTFETCHSTORE`, `TSTPLUSSTORE`, `TSTPAD`, `TSTUNUSED`, `TSTVUNUSED`, `TSTCMOVE`, `TSTCMOVEGT`, and two others) individually and immediately after writing each test, per this session's now-standard practice. Label collisions checked programmatically before insertion - caught 4 real ones this way (`PDFAIL`/`PDDONE`, colliding with section 3.5's own `TSTDPLUS`; `CMDONE`/`CGDONE`, colliding with `CMOVEV`'s/`CMOVEGT`'s own internal completion labels - the same class of issue as `MSTAR`/`LSHIFT`/`RSHIFT`/`HOLDS`/`DUMP`/`TYPE` found in earlier sections), replaced with `PA`/`CV`/`CX` respectively, each re-verified against every label in the file, including this batch's own internal loop labels.

Checked the `IFEQ`/`ENDC` balance immediately after insertion, per this session's now-standard practice - balanced (35/35) and correctly nested end to end on the first attempt.

Verified: all 16 depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with the apparent "missing" list (character literals, hex literals, a mnemonic, and this batch's own local loop labels) traced to its real, harmless cause rather than dismissed

without checking; all 16 tests confirmed wired into `TSTMEMORY`, and `TSTMEMORY` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s eleventh value included (24 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=10` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTMEMORY` tests now pass, including `TSTCFETCHSTORE` after its own fix (logged separately above) - both overlapping-direction `MOVE` cases, the combined `CODEHERE`/`VARHERE`-region walk-throughs, and every other test in this section all confirmed correct on real hardware.

- **RESOLVED - real bug found by the user's actual MAME run: TSTCFETCHSTORE failed because the test's own expected value after c! was backwards - CSTOREW itself is correct.** Root cause confirmed directly against the user's own register and memory dump: this test expected \$FF34 after storing \$34 into a cell pre-filled with \$FFFF, assuming the untouched byte would end up as the word's high byte - but CSTOREW writes directly at the exact address given (TSTCBUF itself, the *lower* address), leaving TSTCBUF+1 (the higher address) untouched at \$FF. Since LDD on this big-endian 6809 reads the word's high byte from the lower address first, the correct result is \$34FF, not \$FF34 - exactly what the user's own dump showed (D=\$34FF at the failing CMPD, with the memory dump confirming 34 FF sitting at TSTCBUF / TSTCBUF+1). The single-byte write and the zero-extending fetch this test also checks were both already correct; only this one comparison's own arithmetic about which byte ends up where was wrong.

Fixed by correcting the expected value to `\$34FF` and rewriting the test's own header comment, which had described the pre-fill byte's position ambiguously enough to have contributed to the original mistake - now states precisely that `C!` touches the single byte at the exact address given, not "the other half" of a 2-byte cell in some assumed sense.

Verified: `IFEQ`/`ENDC` balance unaffected (35/35); this test's own depth-check value unchanged (`-2`, since the fix only changes a comparison constant); full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value 0-10 (24 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own register/memory-dump-level detail is what let this be pinpointed precisely rather than guessed at.

- **** TSTSTRPARSE added - strings & parsing tests (glossary section 3.12, 16 words, 19**

tests since several get separate cases: `[CHAR]` compiling/interpreting state, `S` compiling/ interpreting state (this word genuinely has both), `SEARCH` found/not-found, `REPLACES / SUBSTITUTE` combined across two tests, `SNAMES` found/not-found), wired into `TSTRUNNER` and inserted last in the source (after `TSTMEMORY`), matching glossary order, gated as `TSTSELECTOR` 's twelfth value (`-11`).

****A real, significant naming collision caught by the standard collision check, not just an internal-label clash like earlier sections' finds**:** this section's own ``COMPARE`` (string comparison) word's test was originally named ``TSTCOMPARE`` - which turned out to already be the group-level wrapper name for section 3.7's entire comparison-operators test group (``=`, `<`, `<>`, etc. - a completely different, pre-existing thing, not an internal label inside some word's own implementation`). Renamed to ``TSTSCOMPARE`` throughout, including its own display name (the string ``TSTREPORT`` prints), confirmed via simulation that the real section 3.7 ``TSTCOMPARE`` wrapper is completely unaffected and still compiles correctly, and confirmed ``TSTSCOMPARE`` itself doesn't collide with anything either. Three more internal-label collisions caught the same way as earlier sections (``PNDONE`` inside ``PARSENAME``'s own real implementation, ``UEDONE`` inside ``UNESCAPEW``'s own real implementation, ``DQFAIL`/`DQDONE`` colliding with an existing section 3.7 test's own labels), replaced with ``PZ`/`UX`/`DX`` respectively.

****`.`** confirmed to have no ``STATE`` check at all in its own code** - matches "no interpretation semantics, per ANS" without a specific ``-14`` throw (unlike ``[CHAR]`/`S``, which do throw), so only its compiling case is tested, matching how it's actually meant to be used. Its own runtime (``DOTSTR``) calls ``TYPE`` internally, which - per section 3.1's own already-debugged finding - has two entirely different implementations gated by ``SERIALPOLL``; this test applies that same lesson from the start rather than re-discovering it via a failed MAME run, using the same split established for ``TSTCR``: full ``OUTHEAD`/`OUTBUF`` verification under ``SERIALPOLL=0``, a narrower ``EMIT``-based check under ``SERIALPOLL=1``.

****`PARSE` and `PARSE-NAME` each get a test specifically designed to isolate their one documented behavioral difference**** - a fake source starting with the delimiter itself, confirming ``PARSE`` genuinely does not skip it (returns a zero-length token immediately) while ``PARSE-NAME`` genuinely does skip leading spaces - rather than two generic parse tests that could pass even if the two words were accidentally sharing one code path.

****`SUBSTITUTE`'s** own expected outputs were independently simulated in Python against the algorithm as traced from the real source**, not hand-derived and assumed correct - covering all 4 of its own documented `%`-delimiter cases across two tests: `TSTREPLSUBS1` (the `%`-collapses-to-`%` case, a registered-name substitution, and a non-registered-name pass-through, all in one template) and `TSTREPLSUBS2` (an unpaired trailing `%` with no closing delimiter, tested in isolation for clearer failure diagnosis). `UNESCAPE`'s and `REPLACES`'s own extensive prior-session bug-fix comments (doubling `%`, not backslash-escaping; single-slot registration) were confirmed fresh against the actual current code before designing tests around them, not trusted blindly.

Nine `FCB`-length/string-length mismatches caught individually and immediately after writing each test, per this session's now-standard practice (`TSTCHARW`, `TSTPARSE`, `TSTPARSENAME`, `TSTBRACKCHAR1`/`TSTBRACKCHAR2`, `TSTDASHTRAILING`, `TSTSLASHSTRING`, `TSTREPLSUBS1`, `TSTSCOMPARE` after its own rename left a stale length byte from the old, shorter name).

Checked the `IFEQ`/`ENDC` balance immediately after insertion, per this session's now-standard practice - balanced (39/39) and correctly nested end to end on the first attempt.

Verified: all 19 depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with the apparent "missing" list (character literals, mnemonics, directives, and this batch's own local labels) traced to its real, harmless cause rather than dismissed without checking; all 19 tests confirmed wired into `TSTSTRPARSE`, and `TSTSTRPARSE` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s twelfth value included (26 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=11` includes only this group's own bodies, and separately confirmed via simulation that the real section 3.7 `TSTCOMPARE` group's own body is still present and unaffected by this section's rename. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTSTRPARSE` tests pass, including `TSTDOTQUOTE` - the terminal output itself visibly showing "HI" (the test's own emitted string) ahead of "TSTDOTQUOTE OK", confirming the `SERIALPOLL`-conditional handling applied from the start (per section 3.1's own lesson) held correctly on real hardware on the first attempt, with no repeat of that earlier bug. Both `S` state cases, both `[CHAR]` state cases, `SUBSTITUTE`'s

independently-simulated expected outputs, and the renamed
`TSTSCOMPARE` (confirmed not colliding with the real section 3.7
group) all confirmed correct as well.

- **** TSTNUMOUT added - numeric output tests (glossary section 3.13, 14 words, 13 tests since #S and #> are combined into one test, matching how they're naturally used together as the tail of the standard pictured-output idiom <# ... #S #>), wired into TSTRUNNER and inserted last in the source (after TSTSTRPARSE), matching glossary order, gated as TSTSELECTOR 's thirteenth value (-12).**

****`SIGN` confirmed to carry real bug-fix history worth designing the test around, not just noting**:** a prior version used a bare `BPL` immediately after `PULU`, which doesn't affect condition codes on real 6809 hardware - the bug was masked for every internal caller (DOT/DOTR/DDOT/DDOTR) by "caller-side luck" (each happens to run a flag-setting `LDD` of the true value right before calling `SIGN`, and `PSHU` doesn't disturb flags). The current code is fixed (explicit `TSTA` immediately before the branch), but a naive test calling `SIGN` the same way those internal callers do would repeat the identical caller-side-luck pattern and fail to catch a future regression if `TSTA` were ever removed again. This section's own `TSTSIGN` instead deliberately poisons the CPU flags with an unrelated `CMPX` (setting them to reflect an unrelated comparison result, without touching the value about to be pushed) immediately before each direct `PSHU`/`JSR SIGN` call, for both a negative and a positive test value - genuinely exercising the "direct, standalone call" scenario the bug-fix comment says was previously broken, not a scenario likely to pass by the same kind of luck that hid the original bug.

****Every one of the 7 high-level printing words (`.`/`U.`/`.R`/`U.R`/`?`/`.D.`/`.D.R`) applies the `SERIALPOLL`-conditional split established in section 3.1 from the start** - full `OUTHEAD`/`OUTBUF` verification under `SERIALPOLL=0`, an `EMITCH`-based last-character check under `SERIALPOLL=1` - rather than discovering the need for it via a failed MAME run, as originally happened in section 3.1 itself. Explicitly confirmed via simulation (not just assumed from the source) that both branches genuinely compile and diverge as intended for a representative test. Accepted, documented limitation carried over honestly: for the three words with a fixed trailing space (`.`/`U.`/`.D.`), the polling-mode last-character check only confirms a space was emitted, not the actual printed digits - the same trade-off already established for `TSTCR`/`TSTSPACES`/`TSTTYPE`/**

`TSTDOTQUOTE` in earlier sections, not a new one invented here. ``.R`/`.U.R`/`.D.R`` specifically test the padding path (a width wider than the actual digit count - "42" with width 5 producing 3 leading spaces), not just a no-padding sanity case. ``.D.`/`.D.R`` use a genuine double-cell value (-100000: high cell ``. $FFFE``, low cell ``. $7960``) computed and independently verified in Python before writing the tests, specifically to exercise the 32-bit negation path (``.MNEG32``), not a value that happens to fit in a single cell.

****Caught and fixed a real mistake in this section's own first ``.TSTSIGN`` draft before it was ever inserted****: an intermediate check compared ``.HLD`` against ``.CODEHERE``'s raw redirected value instead of the actual computed ``.PAD`` address (``.CODEHERE`` plus ``.PADOFFSET``) - would have been silently wrong every time, not caught by any later verification step since it was fixed during drafting, before the collision/depth/reference checks ran.

A significant naming-collision risk from the prior section (``.TSTCOMPARE`` accidentally colliding with an existing group wrapper) prompted extra care here: this group's own wrapper name (``.TSTNUMOUT``) was explicitly checked against the whole file before use, not assumed safe by pattern alone. Five internal-label collisions still turned up via the standard check and were fixed the same way as every prior section's finds: ``.HSFAIL`/`.HSDONE`` (the latter colliding with ``.HOLDS``' own real internal completion label), ``.NSFAIL`/`.NSDONE``, ``.NGFAIL`/`.NGDONE``, ``.UDFAIL`/`.UDDONE`` (three colliding with existing, unrelated tests' own labels from earlier sections), replaced with ``.HO``, ``.NZ``, ``.NX``, ``.UF`` respectively.

Seven ``.FCB``-length/string-length mismatches caught individually and immediately after writing each test, per this session's now-standard practice.

Checked the ``.IFEQ`/`.ENDC`` balance immediately after insertion - balanced (55/55, an increase of exactly 16 pairs over the prior section's 39, matching the expected count precisely: 7 I/O tests x 2 ``.SERIALPOLL`` blocks each, plus 2 group-level pairs) and correctly nested end to end on the first attempt.

Verified: all 13 depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, including several (``.QMARK``, ``.DDOT``, ``.DDOTR``, ``.HOLD``, ``.HOLDS``) initially hidden from the automated check by this session's own recurring local-prefix-exclusion artifact, each

confirmed directly via a literal call-site count rather than trusted blindly; all 13 tests confirmed wired into `TSTNUMOUT`, and `TSTNUMOUT` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s thirteenth value included (28 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=12` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all `TSTNUMOUT` tests now pass, including `TSTNUMSIGN`/`TSTNUMSIGNSGT` after their own two fixes (the pre-`COLD` `BASE=0` infinite loop and the depth-check mistake, both logged separately above) - the deliberately-poisoned-flags `TSTSIGN` test, both padding-path tests (`R`/`U.R`), and the genuine double-cell negation in `D.`/`D.R` all confirmed correct on real hardware, with the terminal output itself visibly showing each printed value ("42", "42", " 42", "-7", "-100000", " -100000") exactly as designed.

- **RESOLVED - real bug found by the user's actual MAME run: an infinite loop in `TSTNUMSIGNSGT`, which turned out to affect every test in this section that does digit conversion, not just the one that happened to surface it.** Root cause confirmed by tracing the real boot sequence directly, not assumed: this entire test framework runs from `COLDSTRT`, which clears the whole `GLOBALS` region (including `BASE`, at offset `$02`) to zero and *then* calls `TSTRUNNER` - with `COLD`'s own `BASE=10` initialization not running until later, after the test framework has already finished. So `BASE` reads as genuine zero at test time, not 10. With `BASE=0`, `UDDIGIT`'s own restoring- division algorithm (traced precisely against the user's own register/memory dump, which showed `REM / UDHI / UDLO` all stuck at `$FFFF / $FF` after `DCNT` had reached zero) degrades into an unconditional shift - the "compare against `BASE`, then subtract" step can never actually subtract, since nothing can compare below zero - so the value being converted never genuinely decreases. `NUMSIGNS`' own loop-until-zero condition then never becomes true, and `HOLD` keeps decrementing without bound, exactly matching the user's own observation of memory filling with repeated '6' characters far beyond the pictured- output region's own bounds.

This was never a bug in `UDDIGIT`/`NUMSIGN`/`NUMSIGNS` themselves, which reasonably assume `BASE` already holds a valid, nonzero value by the time they're used - true for every normal, post-`COLD` use of these words, just not true during this framework's own pre-`COLD` execution window. Structurally the same class of gap as every other system-variable dependency already handled via save/redirect/restore throughout this whole session (`CODEHERE`, `STATE`, `SRCADDR` etc.) - `BASE` was simply the one variable in this dependency chain not yet

identified as needing the same treatment, since the words in this specific section were the first to genuinely depend on its value rather than just its address or general presence.

Fixed by adding an explicit save/set/restore of `BASE` (save the real value, set it to 10, restore afterward) to every one of the 9 affected tests: `TSTNUMSIGN`, `TSTNUMSIGNSGT`, `TSTDOT`, `TSTUDOT`, `TSTDOTR`, `TSTUDOTR`, `TSTQMARK`, `TSTD DOT`, `TSTD DOTR` - the same pattern already established for every other redirected system variable in this file, not a new mechanism invented for this fix. Added a new, dedicated scratch constant (`TSTBASAV`) for this, with a detailed comment explaining the root cause directly at its declaration, matching this session's established practice of documenting bug-fix history at the point of the fix, not just in this checklist. Applied the identical fix systematically across the 6 structurally-identical `SERIALPOLL`-conditional tests via a single Python transformation rather than six separate manual edits, then verified each one individually afterward (both the save/set point and the restore point) rather than trusting the bulk transformation blindly. Also cleaned up a cosmetic indentation inconsistency the bulk transformation left behind (a fixed 12-space indent on the inserted lines, not matching the surrounding code's own per-label indent width) before delivery, re-verifying balance and every depth check afterward to confirm the cleanup itself introduced no regressions.

Verified: `IFEQ`/`ENDC` balance unaffected (55/55, unchanged - this fix only added `LDD`/`STD` memory operations, no new conditional blocks); every one of the 13 depth-check values confirmed unchanged (the fix touches only `BASE`, never `U`); full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value 0-12 (28 scenarios total) - all pass, both before and after the indentation cleanup. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own register/memory-dump-level detail (specifically the `DCNT=0` reading, confirming the division's own internal 32-iteration loop had genuinely completed, combined with `REM`/`UDHI`/`UDLO` all reading `\$FFFF`/`\$FF`) is what allowed the root cause to be pinpointed to `BASE`'s own value precisely, rather than guessed at via the division algorithm's structure alone.

- **RESOLVED - real assembly error found by the user's actual lwasm run: the previous `BASE` fix's own cosmetic cleanup step introduced four column-1 mnemonics, each parsed by lwasm as a label rather than an instruction (producing**

both “Bad opcode” and, for the second occurrence at each pair, “Multiply defined symbol” - since `LDD / STD` as label names then collided across the two structurally-identical, previously-fixed tests). Root cause: the Python script used to normalize the earlier bulk fix's own inconsistent indentation computed each inserted line's indent from the immediately preceding line in the output - but at both of the user's own reported locations (`TSTNUMSIGN` 's and `TSTNUMSIGNSGT` 's own `BASE` restore points), that preceding line was blank, so the computed indent came out as zero instead of matching the surrounding code.

Given the risk that this indentation-computation bug could have produced other, unreported column-1 lines elsewhere in the file (not just the two the user's own `lwasm` run happened to reach before stopping), searched the entire file for every mnemonic appearing at column 1, not just the two locations named in the error report. Found two more, unrelated to this specific fix and pre-dating this session's own ``BASE`` work entirely (``LDD #10`/`PSHU D`` pairs within ``TSTCODEHERE``'s and ``TSTVARHERE``'s own ``ALLOT`/`VALLOT`` steps, both in section 3.11) - fixed all four rather than only the two reported, then re-scanned the entire file programmatically (checking every real 6809 mnemonic against every line's own leading whitespace) to confirm zero remaining column-1 occurrences anywhere, not just in the sections touched by this specific round of fixes.

Verified: ``IFEQ`/`ENDC`` balance unaffected (55/55); every one of the 13 depth-check values in this section confirmed unchanged (these were purely whitespace fixes, touching no logic); full scenario matrix re-run crossing both real ``SERIALPOLL`` settings with every ``TSTSELECTOR`` value 0-12 (28 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own successful `lwasm` assembly run to confirm - the user's own assembler error output is what caught this, not a runtime/MAME-level issue.

- **RESOLVED - real bug found by the user's actual MAME run: `TSTNUMSIGN` failed on its own final depth check - `NUMSIGN` itself is correct, the test's own expected value was wrong.** Root cause confirmed directly against the user's own register dump (`D=$0004` exactly at the failing `CMPD`): this test captures `TSTUAF` immediately after `NUMSIGN` returns, with `ud2` (the computed quotient) still sitting on the stack - 2 more cells than `TSTUB4` , which was captured before `NUMSIGN` even ran. The expected depth had been written as `0` , apparently by mistakenly copying the reasoning from the neighboring `TSTNUMSIGNSGT` (which calls `#>` to consume the accumulated result *before* capturing `TSTUAF` , genuinely landing back at net zero) without re-deriving it for this

test's own, structurally different capture point. Fixed to the correct value (4), with a comment at the fix explaining the mistake directly, matching this session's established practice.

Verified: `IFEQ`/`ENDC` balance unaffected (55/55, unchanged - this was purely a comparison-constant fix plus an explanatory comment); this test's own depth-check sequence re-confirmed correct (`0`/`2`/`4`) using a more careful check that strips comments first, after an initial naive regex check was fooled by the new multi-line comment sitting between the fixed `CMPD` and its own `BNE` (a verification-script limitation only - the comment itself is standard, valid assembly syntax with no bearing on real assembly or execution); confirmed zero column-1 mnemonics anywhere in the file, re-checked after this edit specifically given the immediately preceding round of fixes in this same file; full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value 0-12 (28 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own register-dump-level detail is what let this be pinpointed to the exact, precise expected value rather than re-derived from scratch or guessed at.

- **** TSTBASERADIX added - base/radix control tests (glossary section 3.14, 4 words: BASE / DECIMAL / HEX / BINARY), wired into TSTRUNNER and inserted last in the source (after TSTNUMOUT), matching glossary order, gated as TSTSELECTOR 's fourteenth value (-13). A small, simple section - all four words are pure BASE state-setting/reading with no I/O involved, so none of this section's own testing needed the SERIALPOLL - conditional complexity every printing-heavy section since 3.1 has required.**

Combined into a single test (`TSTBASE`) rather than four separate ones - each word's own effect is a single, trivially-verified `BASE` read or write, so four separate tests would just repeat the identical save/set/verify/restore shape four times over with nothing meaningfully different to isolate. Sets `BASE` via `HEX`/`BINARY`/`DECIMAL` in turn (verifying the correct value lands each time), then calls `BASE` itself and verifies both that it returns the variable's own address (not its value, per its own documented `( -- addr )`) and that reading through that returned address reflects `DECIMAL`'s own, most recent setting - confirming the address is genuinely live, not just structurally correct.

Reused `TSTBASAV` (the scratch constant added in section 3.13,

when the real, pre-`COLD` `BASE=0` bug was found and fixed) to save and restore the real `BASE` across this test too, per the same established reasoning - this section's own tests deliberately set `BASE` to specific values as their whole point, so leaving the real value disturbed afterward would be a regression of exactly the kind that bug-fix was meant to prevent going forward.

Checked the group wrapper name (`TSTBASERADIX`) against the whole file before use, per the practice established after the `TSTCOMPARE` collision in an earlier section - safe, no collision. Checked every column-1 mnemonic occurrence immediately after insertion too, per the practice established after the indentation mistake found via a failed lwasm run in section 3.13 - zero found.

Verified: `IFEQ`/`ENDC` balance immediately after insertion - balanced (57/57) and correctly nested end to end on the first attempt; the test's own depth-check value and all four intermediate value checks confirmed correct; real word reference confirmed to resolve for all four words; the test confirmed wired into `TSTBASERADIX`, and `TSTBASERADIX` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s fourteenth value included (30 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=13` includes only this group's own body. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports `TSTBASE` passes - `HEX`/`BINARY`/`DECIMAL` all confirmed setting the real `BASE` correctly in sequence, and `BASE` itself confirmed returning a genuinely live address, on real hardware.

- **** TSTEXCEPTION added - exception handling tests (glossary section 3.15, 2 words: CATCH / THROW , 4 tests since the two can only be meaningfully tested together - THROW 's own effect is only observable through a CATCH that traps it), wired into TSTRUNNER and inserted last in the source (after TSTBASERADIX), matching glossary order, gated as TSTSELECTOR 's fifteenth value (-14).**

****`CATCH`/`THROW`** had already been used extensively as this whole session's own error-testing mechanism throughout sections 3.8-3.14** - dozens of prior `-13`/`-14`/etc. throw-code checks across earlier sections already give substantial indirect evidence they work correctly - but this section still got its own, direct, dedicated tests, matching the established practice of every glossary section getting real coverage rather than

resting on incidental use elsewhere.

****Critical safety constraint confirmed by reading `THROW`'s own code directly, not assumed from the glossary alone****: with no active `CATCH`, `THROW` jumps straight to `ABORT` - already established in section 3.1 as unsafe to trigger directly, since it resets the return stack and would destroy this whole test framework's own call chain. Every `THROW` in this section's own tests is either wrapped in a genuine, active `CATCH` (via a dedicated internal helper, `TSTTHROWHLP`, defined specifically so it can only ever be invoked that way - never called directly) or is a `THROW(0)` call, confirmed via `THROW`'s own code to be a pure, unconditional no-op that never touches `HANDLER` or `ABORT` at all, making that one specific case safe to call directly without a wrapping `CATCH`.

`TSTCATCHTHROW` specifically has `TSTTHROWHLP` push two dummy values before throwing, then verifies they're genuinely gone afterward - confirming `CATCH`'s own documented "restores data stack depth... on either path" for real, not just that the final result happens to look right. `TSTHANDLERSAVE` verifies the `HANDLER` half of that same documented guarantee directly, reading `HANDLER` itself before and after `CATCH` calls on both the success path (wrapping `DUP`) and the throw path (wrapping `TSTTHROWHLP` again), rather than only inferring it from stack-level behavior the other tests already cover.

****Caught and fixed a real depth-check mistake in `TSTHANDLERSAVE` before it was ever inserted**** - the same class of mistake found via MAME in section 3.13's own `TSTNUMSIGN`: `TSTUAF` is captured immediately after the second `CATCH` call returns, with the thrown code still sitting on the stack alongside `GUARD` (2 cells), not after the subsequent pops that verify it (which would have left only `GUARD`, 1 cell) - the expected depth was originally written as `0`, re-traced carefully and corrected to the true value (`2`) during drafting, the same way the earlier instance of this exact mistake had to be found by the user via a live MAME run instead.

Added a new scratch constant (`TSTHANDSAV`) for saving/restoring the real `HANDLER` across `TSTHANDLERSAVE`'s own dual-path check.

Checked the group wrapper name (`TSTEXCEPTION`) against the whole file before use, per the practice established after the `TSTCOMPARE` collision in an earlier section - safe. Checked every column-1 mnemonic occurrence immediately after insertion,

per the practice established after the indentation mistake found via a failed lwasm run in section 3.13 - zero found. One real internal-label collision caught by the standard check (``TZFAIL`/`TZDONE``, colliding with an existing, unrelated test's own labels from an earlier section), replaced with ``TV``.

Verified: ``IFEQ`/`ENDC`` balance immediately after insertion - balanced (59/59) and correctly nested end to end on the first attempt; every one of the 4 tests' own depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve; all 4 tests confirmed wired into ``TSTEXCEPTION``, and ``TSTEXCEPTION`` itself confirmed wired into ``TSTRUNNER``. Full scenario matrix re-run with ``TSTSELECTOR``'s fifteenth value included (32 scenarios total) - all pass; explicitly confirmed ``TSTSELECTOR=14`` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all 4 tests pass, including ``TSTHANDLERSAVE`` after its own pre-insertion depth-check correction - both ``CATCH``'s success and exception paths, ``THROW(0)``'s no-op behavior, and ``HANDLER``'s correct restoration across both paths all confirmed correct on real hardware.

- **** TSTCOMMENTS added - comments tests (glossary section 3.16, 2 words: (and \), wired into TSTRUNNER and inserted last in the source (after TSTEXCEPTION), matching glossary order, gated as TSTSELECTOR 's sixteenth value (-15). Both words parse from source, so both redirect SRCADDR / SRCLEN / TOIN , matching the established pattern from sections 3.8-3.12.**

****`(`** confirmed to carry real bug-fix history worth verifying directly, not just noting****** - its own code comment documents a prior version that omitted consuming ``WORD``'s own returned c-addr, leaving a stray value on the data stack after every ``"(...)"`` comment, confirmed and fixed via MAME in an earlier session (symptomatic both as a visible stray value in interpret mode and as a ``-22`` CSP mismatch for any colon definition containing one). This section's own test specifically confirms the data stack is genuinely unchanged (net 0) after the call, not just that parsing itself advanced to the right place - directly re-verifying the fix holds, not merely trusting the comment.

``\`` specifically redirects ``SRCLEN`` (not just ``TOIN``) to a value distinct from the real terminal input buffer's own

length, confirming it genuinely reads `SRCLen` directly rather than some fixed buffer size - the actual mechanism behind its own documented "follows `SOURCE`... operates correctly inside `EVALUATE`", not just trusting the description.

****Caught and fixed a real bug in this section's own first `TSTLPAREN` draft before it was ever inserted****: referenced an undefined scratch constant (`TSTSASAV2`) to re-read `SRCADDR` after the call, when `(` in fact never touches `SRCADDR` itself (only `TOIN`) - simplified to check the fake source's own already-known address directly, removing the unnecessary and broken re-read entirely rather than defining a new constant to paper over it.

Checked the group wrapper name (`TSTCOMMENTS`) against the whole file before use, per the practice established after the `TSTCOMPARE` collision in an earlier section - safe. Checked every column-1 mnemonic occurrence immediately after insertion, per the practice established after the indentation mistake found via a failed lwasm run in section 3.13 - zero found. Zero internal-label collisions this time.

Verified: `IFEQ`/`ENDC` balance immediately after insertion - balanced (61/61) and correctly nested end to end on the first attempt; both tests' own depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with `LPAREN` itself initially hidden from the automated check by this session's own recurring local-prefix-exclusion artifact, confirmed directly via a literal call-site count rather than trusted blindly; both tests confirmed wired into `TSTCOMMENTS`, and `TSTCOMMENTS` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s sixteenth value included (34 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=15` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports both tests pass - `(`'s own stray-value bug fix confirmed holding on real hardware, and ``'s genuine `SRCLen`-following behavior confirmed correct.

- **** TSTENVSYS added - environmental & system queries tests (glossary section 3.17, 10 words, 8 tests since TIB / #TIB / >IN / SPAN / BL are combined into one test, matching the established approach from section 3.14's own TSTBASE, and ENVIRONMENT? gets three separate tests of its own - single- cell, double-cell, and unsupported-string cases), wired into TSTRUNNER and inserted last in the source (after TSTCOMMENTS),**

matching glossary order, gated as TSTSELECTOR 's seventeenth value (-16).

****Two safety findings confirmed directly from the actual code, not assumed from the glossary**:** `REFILL`'s terminal-source branch calls `QUERY`, which - per section 3.1's own already-established finding - calls `ACCEPT`, which calls `KEY`, which spins forever without real input; only its safely-testable path (source is a string, `SRCID<>0`, returning `FALSE` immediately) is tested, matching the same reasoning already applied to `KEY`/`ACCEPT`/`EXPECT`/`QUERY` themselves. `EVALUATE` genuinely executes its given string as real Forth source via `JSR INTERPRET` - confirmed safe by reading `INTERPRET`'s own code directly: it simply stops and returns when input is exhausted, with no automatic `REFILL` call, so no blocking risk - making a short, deliberately harmless string (`"1 2 +"`) safe to evaluate directly. `TSTEVALUATE` verifies both the arithmetic result and that the whole input-source state (`SRCADDR`/`SRCLEN`/`SRCID`/`TOIN`) is genuinely restored afterward, not left pointing at the evaluated string.

****`ENVIRONMENT?` confirmed to carry substantial, already-documented bug-fix history worth designing a test around, not just noting**:** its own code comment describes a prior `COMPAREW`-clobbers-`X` bug, confirmed via MAME with the exact `\$4E4D` symptom (reading two bytes of the *next* table entry's own name text as the current entry's value) - specifically when `/COUNTED-STRING` was queried, since its own match left `X` sitting exactly at the start of `MAX-N`'s own string. `TSTENVQUERY1` queries that exact adjacent pair again, `/COUNTED-STRING` immediately followed by `MAX-N`, so a regression of that specific mechanism would be caught again, not just assumed still fixed from the comment alone. `TSTENVQUERY2` separately exercises the double-cell table path (`MAX-D`, `\$7FFFFFFF`) that the single-cell path can't cover, and `TSTENVQUERY3` confirms an unsupported string correctly reports `false`.

Added a new scratch constant (`TSTSISAV`) for saving/restoring the real `SRCID` across `TSTSOURCEID`'s and `TSTREFILL`'s own tests, both of which redirect it directly.

Caught a real naming collision during the standard check before insertion: `TSTENVQUERY1`'s own display-name label (`TSTEQNAME`) collided with an *existing* test's own label from an earlier section (`TSTEQ`, the `=` comparison test) - not just an internal `FAIL`/`DONE` pair this time, but the name-

string label itself. Renamed the whole prefix (`EQ` to `EN1`) rather than patching around just the one colliding label.

Verified: `IFEQ`/`ENDC` balance immediately after insertion - balanced (63/63) and correctly nested end to end on the first attempt; zero column-1 mnemonic issues; every one of the 8 tests' own depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with `EVALUATEW`/`SOURCEW`/`SOURCEID` initially hidden from the automated check by this session's own recurring local-prefix-exclusion artifact, confirmed directly via literal call-site counts rather than trusted blindly; all 8 tests confirmed wired into `TSTENVSYS`, and `TSTENVSYS` itself confirmed wired into `TSTRUNNER`. Full scenario matrix re-run with `TSTSELECTOR`'s seventeenth value included (36 scenarios total) - all pass; explicitly confirmed `TSTSELECTOR=16` includes only this group's own bodies. Byte-exact split-file reassembly confirmed. ****MAME-CONFIRMED****: the user reports all 8 tests pass, including `TSTEVALUATE` after its own two-round fix (`LATEST` and `CODEHERE` both needing redirection, logged separately above) - both `ENVIRONMENT?` bug-fix-regression checks (`/COUNTED-STRING`+`MAX-N` and the double-cell `MAX-D` path), `REFILL`'s safely-testable string-source path, and every other test in this section all confirmed correct on real hardware.

- **RESOLVED (in two rounds - the first fix was real but incomplete) - real bug found by the user's actual MAME run: the full test sequence terminated partway through TSTEVALUATE , never reaching TSTENVQUERY1 / 2 / 3 .**

****Round 1**** identified a genuine issue - `LATEST`, like `BASE` (section 3.13's own earlier finding), is only set by `COLD` (`LDD #BASELATEST`/`STD LATEST`, right alongside `BASE=10` in the same init block), which hasn't run yet at this whole test framework's own pre-`COLD` execution point. With `LATEST=0`, `FIND` would report every word as not found. This fix (save/set/restore `LATEST` around the call) was applied and delivered - but the user's own follow-up MAME run showed the sequence still terminated, with the specific value left visible in the terminal output differing between runs (`2`, then `1`) - a strong sign of genuine memory corruption rather than a single, deterministic missing-value problem, prompting a full second round of investigation rather than assuming the first fix needed only a minor adjustment.

****Round 2 found the actual, complete root cause**:** ``CODEHERE`` had never been redirected by this test at all - unlike every other test in this entire session involving ``WORD`` (which ``EVALUATE``'s own internal ``JSR INTERPRET`` calls repeatedly, once per token). With ``CODEHERE=0`` (also only set by ``COLD``, same as ``BASE`/`LATEST``), ``WORD``'s own write of each parsed token's ``[len][text]`` landed straight at address ``$0000`` - which is ``STATE``'s own address (confirmed by re-reading its own ``RMB`` declaration directly: offset ``$00``, the very first bytes of ``GLOBALS``, which itself starts at ``$0000``). Even a single-character token like ``"1"`` overwrites ``STATE`` with a nonzero value, flipping ``INTERPRET`` into thinking it's compiling rather than interpreting mid-parse - fully explaining both the sequence's own termination and why the visible output varied between runs (the exact corruption pattern depends on residual memory state from whatever ran immediately before).

This was never a bug in ``EVALUATE`/`INTERPRET`/`FIND`/`WORD`` themselves, all of which reasonably assume the real dictionary chain and a valid ``CODEHERE`` are already established by the time they're used - true for any normal, post-``COLD`` use, just not true during this framework's own pre-``COLD`` execution window. ``TSTEVALUATE`` is the first test in this whole section to invoke the **real** ``INTERPRET`/`WORD`/`FIND`` path together against the real, unredirected global state, rather than calling a single lower-level primitive directly the way most other tests do.

Fixed by adding a second, separate redirect: ``CODEHERE`` saved, set to ``TSTCBUF2`` (the existing, second scratch buffer already established in section 3.9 for exactly this kind of conflict - not ``TSTCBUF``, which this test also uses to hold its own source string, ``"1 2 +"``; redirecting ``CODEHERE`` there too would have let ``WORD`` overwrite the very source text still being parsed), and restored afterward alongside the existing ``BASE`/`LATEST`` restores.

Verified: ``IFEQ`/`ENDC`` balance unaffected (63/63, unchanged - this fix only added ``LDD`/`STD`` memory operations, no new conditional blocks); this test's own depth-check values unchanged (the fix touches only ``CODEHERE``, never ``U``); zero column-1 mnemonic issues; full scenario matrix re-run crossing both real ``SERIALPOLL`` settings with every ``TSTSELECTOR`` value 0-16 (36 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation - the user's own observation across two separate runs (the sequence stopping in the same place both times, but

with a different value left visible) is what made the distinction between "wrong value" and "memory corruption" possible to draw, and what prompted going back for the actual, complete root cause rather than assuming the first, real but partial fix was sufficient.

- **** TSTTOOLS added - tools word set tests (glossary section 3.18, 3 words: .S / WORDS / DUMP), wired into TSTRUNNER and inserted last in the source (after TSTENVSYS), matching glossary order, gated as TSTSELECTOR 's eighteenth value (-17).**

****`WORDS` confirmed via its own code to depend directly on `LATEST`**** (walks from `LATEST` until the chain terminates at 0) - applying the lesson section 3.17's own `TSTEVALUATE` surfaced from the start this time: built a small, two-entry fake dictionary chain (`"CD"` pointing back to `"AB"`, terminated with 0) and redirected `LATEST` to it, rather than discovering the same real, pre-`COLD` dependency the hard way again.

****`DUMP` confirmed to carry real, already-documented bug-fix history worth designing a test around, not just noting****: a prior version's ASCII column read past the actual byte count on a partial final line, showing garbage for the unused portion - fixed via `DUVALID` tracking. This section's own test deliberately dumps a non-multiple-of-16 byte count (5), the exact scenario that bug affected, and specifically verifies the ASCII column's own blanking for the remaining 11 columns (via a loop, not 11 unrolled checks) - the bug-relevant part - rather than only checking the real hex data was correct.

****`.S` required a genuinely different testing approach from every other test in this whole session****: it prints the **entire** real data stack, including whatever the calling chain (all the way up through `TSTRUNNER`) already left there before this test even started - unlike every other test, there's no way to predict the complete output in advance. Rather than force an exact-output check the way every other printing word's test in this session has done, this test instead verifies the core, documented guarantee directly: the stack is genuinely unchanged afterward (confirmed by popping the same values back and comparing), plus, in interrupt mode only (where the very first characters queued are unambiguously this test's own top value, since `.S` starts from the current top of stack), that those first characters match - a deliberately narrower check than usual, given the genuine unpredictability of what precedes or follows in the real stack.

****Caught and fixed two real mistakes during drafting, before either was ever inserted**:** ``TSTDOTS``' own first draft put ``IFEQ`` on the same line as its own label (``TSTDOTS: IFEQ SERIALPOLL``) - the exact same mistake, and exact same real risk (an untested pattern `lwasm` might not parse as intended), that caused a genuine assembly error in section 3.13; caught by this session's own balance-check script failing to match the line (the same detection signal as before), not by a live `lwasm` run this time. ``TSTDUMP``'s own first draft had a genuine register conflict in its ASCII-column blank-verification loop - using ``D`/`B`` as both the loop counter and the ``OUTBUF`` index simultaneously, which would have corrupted the index on the very first iteration; caught by manually re-tracing the loop's own register usage rather than trusting it compiled cleanly, and fixed by moving the counter to ``Y``, leaving ``B`` free.

Label collisions checked programmatically before insertion - caught one real one this way (``DSFAIL`/`DSDONE``, colliding with ``DOTS``' own real internal completion label - the same class of issue as ``MSTAR`/`LSHIFT`/`RSHIFT`/`HOLDS`/`DUMP`/`TYPE`/`WORD`` found in earlier sections), replaced with ``DY``. Checked every column-1 mnemonic occurrence immediately after insertion too - zero found.

Verified: ``IFEQ`/`ENDC`` balance immediately after insertion - balanced (71/71) and correctly nested end to end; every one of the 3 tests' own depth-check values independently re-derived from real arity and confirmed correct against the final, fully-assembled file state; every real word reference confirmed to resolve, with ``WORDSW`/`DUMPW`` initially hidden from the automated check by this session's own recurring local-prefix-exclusion artifact, confirmed directly via literal call-site counts rather than trusted blindly; all 3 tests confirmed wired into ``TSTTOOLS``, and ``TSTTOOLS`` itself confirmed wired into ``TSTRUNNER``. Full scenario matrix re-run with ``TSTSELECTOR``'s eighteenth value included (38 scenarios total) - all pass; explicitly confirmed ``TSTSELECTOR=17`` includes only this group's own bodies, and explicitly confirmed via simulation that ``TSTDUMP``'s own ``SERIALPOLL`` branches genuinely diverge as intended. Byte-exact split-file reassembly confirmed.

****MAME-CONFIRMED**:** the user reports all 3 tests pass, including ``TSTDOTS`` after its own ``BASE=0`` fix (logged separately above) - the terminal output itself visibly confirming each test's own design: ``.S`` printing ``7`` then ``TSTVAL1``'s own decimal value (this test's own top-of-stack pushes) ahead of the rest of the real stack beneath them; ``WORDS`` printing ``CD AB`` in the

correct, most-recent-first chain order; and `DUMP` showing the correct hex bytes for "ABCDE" with the ASCII column's own blanking beyond the 5 real bytes, confirming the `DUVALID` partial-line fix holds on real hardware.

- **RESOLVED - real bug found by the user's actual MAME run: TSTDOTS entered an infinite loop when .s internally called DOT .** Root cause: the exact same `BASE=0` infinite loop already found and fixed in section 3.13 - `.s` calls `DOT` for each stack item, `DOT` calls `NUMSIGN / UDDIGIT`, and with `BASE=0` (only set by `COLD`, which hasn't run yet at this whole test framework's own pre- `COLD` execution point), that restoring-division mechanism degrades into an unconditional shift, so the value being converted never genuinely decreases - the same mechanism, the same class of gap, just not yet applied to this specific test when it was first written. Simply forgot to carry the section 3.13 lesson forward to `TSTDOTS` specifically.

Checked whether `WORDS` and `DUMP` needed the same fix before assuming only `TSTDOTS` was affected: confirmed via their own code, not assumed. `WORDS` does no numeric conversion at all (just `TYPE`'s name text via the dictionary chain walk), so it was never at risk. `DUMP`'s own hex conversion (`HEXDIGIT`/`HEXBYTE`) is a fixed, hardcoded base-16 converter - confirmed directly by re-reading its own code, which compares against the literal value `10` and adds `A-10` or `0` outright, never referencing `BASE` at all - genuinely independent of it, not just assumed safe by inference.

Fixed by adding the same save/set/restore of `BASE` already established in section 3.13, reusing the existing `TSTBASAV` scratch constant rather than adding a new one.

Verified: `IFEQ`/`ENDC` balance unaffected (71/71, unchanged - this fix only added `LDD`/`STD` memory operations, no new conditional blocks); this test's own depth-check values unchanged (the fix touches only `BASE`, never `U`); zero column-1 mnemonic issues; full scenario matrix re-run crossing both real `SERIALPOLL` settings with every `TSTSELECTOR` value 0-17 (38 scenarios total) - all pass. Byte-exact split-file reassembly confirmed. This specific fix awaits its own MAME confirmation.

- **Unit test framework extracted into a separate file (`unit_tests.asm`), replacing its previous location inline within `forth6809.asm` with a conditional `INCLUDE` .** The framework's own explanatory comment block (previously immediately after `ORG USROMSTRT`) and its entire body (every test group, `TSTRUNNER`, all scratch-constant `EQU` s

specific to testing) both moved into `unit_tests.asm` verbatim - confirmed via a full line-multiset comparison against the pre-extraction source (every original line accounted for, zero lost, only 2 cosmetic blank lines added for readability at the new transition point).

``forth6809.asm`` itself keeps the ``IFDEF UNITTESTS`/`IFDEF TSTSELECTOR`` default-value blocks in place (unchanged in substance - still the only place either symbol's fallback is set when no ``-D`` is passed), immediately followed by the existing, already-proven ``IFNE UNITTESTS`` wrapper - unchanged in kind, just narrowed to gate a single ``INCLUDE unit_tests.asm`` line instead of the framework's full body directly. Deliberately kept this session's own established ``IFNE`` convention (matching the existing, tested call site in ``COLDSTRT`` that decides whether to ``JSR TSTRUNNER``) rather than introducing a new, untested ``IFEQ``-based convention, even though an earlier illustrative example used ``IFEQ`` - the person who requested this explicitly allowed for "whatever the correct syntax is."

``unit_tests.asm`` itself does ****not**** carry its own outer ``IFNE UNITTESTS`/`ENDC`` pair - that wrapping now lives once, at the include site in ``forth6809.asm``, rather than being duplicated in both files. Confirmed via ``lwasm``'s own manual that ``INCLUDE`` is the correct directive name (there is no alternate spelling or pseudo-op required), and that ``-D SYM[=VAL]`` predefines a symbol exactly as though set via ``SET`` in-source - matching this file's own existing ``-D UNITTESTS=1`` command-line convention precisely, unaffected by this change.

Verified: ``IFEQ`/`ENDC`` balance checked three ways - the combined, simulated source (as ``lwasm`` would see it after processing the ``INCLUDE``) balances at 71/71; ``unit_tests.asm`` considered entirely on its own also self-balances (64/64, confirming no test group's own conditional block was left dangling by the extraction); and 7 (``forth6809.asm``'s own remaining conditionals) + 64 = 71 confirms the split accounts for every block precisely. Full 38-scenario matrix (both real ``SERIALPOLL`` settings crossed with every ``TSTSELECTOR`` value, plus the two production-default scenarios) re-run against the combined, simulated source - all pass, dictionary chain and label-collision checks included. Split reference files regenerated from the new, shorter ``forth6809.asm`` and confirmed to byte-exact-reassemble back to it. ****`lwasm`-CONFIRMED****: the user reports the ``INCLUDE unit_tests.asm`` extraction has been tested as working on a real ``lwasm`` run - the file split, the

narrowed `IFNE UNITTESTS` wrapper, and the removal of
`unit\_tests.asm`'s own duplicate outer conditional all assemble
correctly in practice, not just in this session's own
simulation.

- **RESOLVED - real bug found by the person's own automated test runner (the very first run of the new `run_all_tests.sh` script): `TSTSNAME1` (section 3.12, `SNAME`'s own found-case test) failed.** Root cause: the same `LATEST=0` pre- `COLD` dependency already found and fixed once in section 3.17's own `TSTEVALUATE` - confirmed directly by re-reading `SNAMEW`'s own code, which starts its dictionary-chain walk from `LATEST` (`LDD LATEST / STD SNXT`) exactly the way `FIND` does. With `LATEST=0` (only ever set by `COLD` , which hasn't run yet at this whole test framework's own pre- `COLD` execution point), `SNLOOP`'s own `BEQ SNNOTFOUND` fires immediately, before ever comparing against `DUP`'s own `CFA` - `TSTSNAME1` was written in section 3.12, before the `LATEST=0` finding existed at all, so it never received the fix section 3.17 later established for exactly this dependency.

****Also caught and fixed `TSTSNAME2` (the not-found case) for the same root cause, even though it was reported passing**** - its own passing result was masked, not genuine: with ``LATEST=0``, the chain is trivially empty, so ``SNAME`` reports not-found immediately regardless of the target address, meaning this test was only ever exercising the degenerate "empty chain" case, not the real one it claims to (walking a genuine, populated chain and finding no match for an address that isn't in it). Fixing ``TSTSNAME1`` alone would have left this one quietly untested in substance while still showing green.

Fixed both the same way section 3.17 fixed ``TSTEVALUATE``: added an explicit save/set/restore of ``LATEST`` (save the real value, set it to ``BASELATEST``, restore afterward) to each test, reusing the existing ``TSTLSAV`` scratch constant already established for this exact purpose, rather than adding a new one.

Verified: ``IFEQ`/`ENDC`` balance unaffected (64/64 within ``unit_tests.asm`` on its own, unchanged - this fix only added ``LDD`/`STD`` memory operations, no new conditional blocks); both tests' own depth-check values unchanged (the fix touches only ``LATEST``, never ``U``); zero column-1 mnemonic issues. Full 38-scenario matrix re-run against the combined, simulated source (``forth6809.asm`` with ``unit_tests.asm`` substituted in for its own ``INCLUDE``) - all pass. ****AUTOMATION-CONFIRMED****: the person's own full ``run_all_tests.sh`` run across all 18 glossary

sections reports every section passing, `3.12\_StrParse` included - the first time this whole test suite has been confirmed end-to-end via the automated runner rather than individual manual MAME sessions.

Structural duplication (identified, some resolved, some not)

- `:` / `CREATE` / `VARIABLE` 's header-building — resolved via `HEADER` (section 7).
- `COMMA` / `CODECOMMA` / `CCOMMA` / `VCOMMA` / `VCCOMMA` / `CCOMMA1` — resolved via `APPENDCELL` / `APPENDBYTE` (section 6).
- `<# / #>` recomputing PAD's address inline — resolved, both now call `PADW` (section 20 / section 6).
- **Identified via the real listing (`forth6809_1st.txt`), not yet acted on: `NUMBERQ` 's inline 32-bit negation (added this session for double-number text input) duplicates `MNEG32` , an existing, already-correct shared routine `DNEGATE` and `DABS` both already call.** Traced precisely why this duplication mattered beyond just code size: `MNEG32` 's instruction order (complement both cells and store both, *then* do the `+1` /carry-propagation step separately afterward) happens to avoid the exact 6809 quirk (`COM` unconditionally sets carry) that `NUMBERQ` 's own, separately hand-written version fell into - which is precisely why `NUMBERQ` 's negation needed two separate bug-fix turns this session while `MNEG32` itself never had the problem at all. Not refactored here - the user's own stated reason for deferring: automated testing should be in place before touching working, duplicated logic, given hand-verifying a refactor across every affected call site the way this session has been doing (MAME tracing turn by turn) doesn't scale well to consolidation work.
- **Identified via the real listing, not yet acted on: the self-referential PFA bug (`CONSTANT` / `DEFER` / `2CONSTANT` / `MARKER`) got four separate, duplicated fixes rather than one shared helper.** Confirmed directly in the listing: `DEFER` 's, `2CONSTANT` 's, and `MARKER` 's fixes are each explicitly commented "same self-referential PFA bug as [`CONSTANT`]," each repeating the identical `LDD CODEHERE` / `ADDD #2` /store sequence independently rather than calling a shared routine. Same deferral reasoning as above - not acted on pending automated testing.
- No further known duplication beyond the two items just above, but the codebase still hasn't been given a full, systematic pass specifically hunting for more - both findings above came from genuinely investigating specific leads, not an exhaustive scan.

Real, load-bearing gaps

- **Software (XON/XOFF) handshaking is still not implemented.** Now that hardware RTS/CTS handshaking is in place, this is the one remaining flow-control gap: this system still never transmits XON/XOFF and would not recognize either byte as a control signal if the remote device sent them — an incoming `$11 / $13` is just queued as an ordinary character. Only relevant for links with no RTS/CTS wiring (plain 3-wire serial); would need a byte- interception layer between the input ring and `KEY` 's caller.
- **No hard boundary checks** anywhere `DPHERE / CODEHERE / VARHERE` grow toward each other or toward the stacks. `UNUSED` and `VUNUSED` both correctly *report* the distance to their boundary now (`CODETOP` and the newly-added `APPVARS` respectively), but nothing enforces either — a runaway compile can silently corrupt an adjacent region.
- **`VUNUSEDW` (the `VUNUSED` word) computed a meaningless number, not “how much `APPVARS` space remains” - a real, distinct bug, not just the absence of a check.**
Fixed. `LDD #APPCODE / SUBD VARHERE` computed `APPCODE - VARHERE` (roughly `$7000` minus wherever `VARHERE` currently sits within `APPVARS`), which had nothing to do with `APPVARS` 's own boundary - looked like a copy-paste slip from `UNUSEDW` immediately above it (`LDD #CODETOP / SUBD CODEHERE` , correct for `CODEHERE` 's own boundary). Surfaced directly by a question about how `APPVARS` 's size is actually set: it wasn't - `APPVARS EQU $021B` defined only its start; there was no end/size constant anywhere. Fixed by adding `APPVARS` `EQU APPVARS+8000` (an exclusive upper bound, `$215B` , matching `CODETOP` 's own convention for `APPCODE`) and changing `VUNUSEDW` to `LDD #APPVARS` `/ SUBD VARHERE` . Verified: at cold start (`VARHERE = APPVARS`), this computes exactly 8000, matching `APPVARS` 's documented size precisely; zero duplicate symbols; dictionary chain unaffected. `APPVARS` initially fell within `APPDICT` 's current range - expected at the time, the same, already-tracked `APPDICT / APPVARS` overlap, not something this fix caused. Since resolved by redefining `APPVARS` as `APPDICT-1` - see above.
- **`ENVTABLE` is incomplete.** Missing entries: `/HOLD` , `/PAD` (this build never fixed a capacity for either — filling them in would mean deciding a real bound first, not just picking a number), `MAX-D` , `MAX-UD` (need the dispatcher extended to push *two* cells for double-cell answers — current `ENVQUERY` only handles one), `WORDLISTS / FLOORED` (need the dispatcher to be able to report a recognized-but-false answer, distinct from “unrecognized name” — no such path exists yet). **RESOLVED - all six entries added and MAME-confirmed this session; see the dedicated entry earlier in this file (search “ENVTABLE completeness”).** `ENVIRONMENT?` is complete.
- **`CATCH -wrapped QUIT / INTERPRET` rollback is scoped to one input line only.** A colon

definition spanning multiple lines that fails partway through a later line only rolls back that line's contribution, not the whole definition back to `:`. A complete fix needs the region-pointer snapshot taken once at `:` and held until `;` or an error, not refreshed every `QLOOP` pass.

- **RESOLVED - confirmed via the real listing (`forth6809_1st.txt`), not just the source text.** `ABORTHDR` 's `LINK` field is a placeholder `0` in the consolidated/split source — must be set to the real prior `LATEST` value once final ROM layout/assembly order is fixed. `ABORTHDR` doesn't exist as a symbol anymore - renamed to `H_ABORT` and moved into `BASEDICT` along with every other header, as part of the broader dictionary-consolidation work referenced above. `H_ABORT` 's actual `LINK` field is a real, resolved address, not a placeholder: `FDB H_FALSE`, with the source's own comment documenting its history (placeholder `0 -> H_DUMPW -> H_DOESGT -> H_DUMPW` again `-> H_FALSE`, tracking the chain's actual newest entry as the dictionary grew during the session). Independently confirmed by the 224-entry dictionary chain walk from `H_M2` to `0`, checked repeatedly throughout this session - that walk only succeeds if every `LINK` field in the chain, including this one, is a real, correct address; a placeholder `0` mid-chain would have broken it outright. No action needed on the source.

Never resolved after being explicitly raised

- **`J` doesn't generalize past one level of loop nesting.** No `J2` /deeper equivalent exists. A triple-nested loop's innermost body has no built-in way to reach the outermost index without manually replicating `J` 's offset arithmetic one frame further.
- **`LEAVE` 's correctness depends on always being textually inside the loop it affects** — never verified against a `LEAVE` called from a separately-defined word invoked from within a loop (which would read/set the wrong stack frame, or crash).
- **RESOLVED via direct code inspection of the real listing (`forth6809_1st.txt`), not just address arithmetic - the underlying inconsistency is now fully understood, though the answer is a genuine mix, not a simple "fixed" or "still open."** Traced `CHARW` (`CHAR`) and `BRACKCHAR` (`[CHAR]`) precisely: both simply push a space delimiter, `JSR WORD`, and read the first character of whatever comes back - no separate cap logic of their own at all. They automatically, correctly inherited `WORD` 's new 46-character cap (`WORDMAXCHARS`) with zero code changes needed - this part of the original item is stale. `FIND` also has no independent cap of its own - traced its comparison down to `HDRFLAGS ANDA #$1F`, masking the dictionary header's `LEN/FL` byte to its low 5 bits. This is the real finding: dictionary header names genuinely are still capped at 31 characters, but not because they "inherit `WORD` 's cap" as the original item claimed - the header format itself allocates only 5 bits (bits 4-0) for the name length, an absolute maximum of 31

regardless of anything `WORD` does or how it's redesigned. That's a separate, permanent architectural constraint (the header storage format), not a leftover inconsistency from `WORD` 's own redesign. The original item's framing was imprecise even when written - it grouped "things that call `WORD` " (`CHAR` / `[CHAR]` , no independent cap) together with "things bounded by the header format" (header names/ `FIND`) as if they were the same constraint, when they're two different ones. Original text: **`WORD` 's 31-character cap is now inconsistent within the system.** `PARSE` / `PARSE-NAME` were deliberately rewritten to not share it, but `WORD` itself — and everything still built on it (`CHAR` , `[CHAR]` , header names, `FIND`) — still has it.

- **The optional Search-Order word set is not implemented.** Every word resolves through one unified dictionary chain rooted at `LATEST` ; there is no multi-wordlist support (`WORDLIST` , `GET-ORDER` / `SET-ORDER` , `ALSO` / `ONLY` , etc.). This was a foundational decision fixed since `COLDSTRT` 's first version, not an oversight — now documented explicitly (ClaudeForth documentation, Section 4.5) along with what adding it later would actually require: restructuring `FIND` into a loop over an active search order, changing `HEADER` to link into a selectable "current" wordlist instead of unconditionally into `LATEST` , and fixing the few words (`RECURSE` , `;` 's unsmudge step, `WORDS`) that read `LATEST` directly today.

Open design questions (not defects — deliberate unresolved choices)

- Whether `ALLOT` / `VALLOT` / `PICK` / `ROLL` / `BASE` should range-check their inputs against actual stack/region bounds. Currently none of them do, consistently, by choice rather than oversight.
- Whether any additional distinction like `TIB` / `SOURCE` (fixed terminal buffer vs. current input source) is needed anywhere else in the system where `EVALUATE` 's redirection might still cause a mismatch.
- `SWIH` 's throw code (`-99` in the consolidated source) was never assigned a real, deliberate value — it's a placeholder for "some hardware trap occurred," not a considered ANS-style code.
- `REPLACES` / `SUBSTITUTE` remain single-slot (one registered name/value pair, overwritten by the next `REPLACES` call) rather than a true multi-entry table. A real table needs its own storage layout and lookup structure — a deliberate scope decision made when `SUBSTITUTE` was completed, not an oversight.

What's solid (for reference, not action)

Stack manipulation (Core + Core Ext + return-stack transfer), all arithmetic (single + double + mixed precision), logic, comparison (single + double), full control flow including `CASE` , all defining words including `DEFER` / `MARKER` / `VALUE` / `TO` (`VALUE` now correctly targeting mutable space), memory and string operations (including a completed `SUBSTITUTE`), `SOURCE` / `EVALUATE` / `REFILL` mechanics, and the Tools word set (`.S` / `WORDS` / `DUMP` , with `DUMP` 's edge-case bug fixed) are complete and were traced/verified against concrete cases during the build, not merely asserted correct.